SPECIFICATION DOCUMENT

Simula Training Data Framework

Reasoning-Driven Synthetic Data Generation for SAP AI Agents

Version	**1.2.0**
Reference	`SIM-TRAIN-2026-v1.2`
Date	April 2026
Author	SCB GCFO, AI Research & Development

Internal Use Only • Confidential

0 •

Contents

1 Executive Summary **5**

Frontmatter **7**

Executive Summary 7

Document Conventions 8

How to Read This Specification 9

Reading Guide 10

Version History 13

Document Scope Matrix 14

2 Overview and Scope **15**

2.1 The Problem 15

2.2 Goals 16

2.3 In-Scope Domains 17

2.4 System Context . . . 18

2.5 Research Paper Traceability . . . 18

2.6 Out of Scope . . . 19

3 Business Requirements Specification **21**

Integration with existing chapters . . . 21

Scope of this chapter . . . 22

3.1 Requirements Overview . . . 22

3.1.1 Requirement Identification Convention . . . 22

3.1.2 Priority Classification . . . 23

3.2 BR-1: Strategic and Preparation Phase . . . 23

3.2.1 BR-1.1: Use Case Definition . . . 24

3.2.2 BR-1.2: Value-Risk Assessment . . . 25

3.2.3 BR-1.3: Governance Framework Integration . . . 26

3.2.4 BR-1.4: Cross-Functional Team Formation . . . 27

3.2.5 BR-1.5: ROI Baseline Establishment . . . 28

3.3 BR-2: Generation Phase . . . 28

3.3.1 BR-2.1: Domain Taxonomy Construction and Validation . . . 29

3.3.2 BR-2.2: Controllable Parameter Configuration . . . 30

3.3.3 BR-2.3: Cost and Resource Budget Control . . . 31

3.3.4 BR-2.4: Metadata and Provenance Recording . . . 32

3.3.5 BR-2.5: Incremental Generation and Dataset Evolution . . . 33

3.4 BR-3: Quality Assurance and Validation . . . 33

3.4.1 BR-3.1: Intrinsic Quality Assessment . . . 34

3.4.2 BR-3.2: Downstream Task Validation . . . 35
3.4.3 BR-3.3: Privacy Risk Assessment . . . 36
3.4.4 BR-3.4: Bias Detection and Mitigation . . . 37
3.4.5 BR-3.5: Continuous Monitoring and Drift Detection . . . 38
3.5 BR-4: Governance and Compliance . . . 38
3.5.1 BR-4.1: Access Control and Permission Management . . . 39
3.5.2 BR-4.2: Regulatory Framework Alignment . . . 40
3.5.3 BR-4.3: Audit Logging and Traceability . . . 41
3.5.4 BR-4.4: Transparency and Explainability . . . 42
3.5.5 BR-4.5: Incident Response and Remediation . . . 43
3.6 BR-5: Integration and Infrastructure . . . 43
3.6.1 BR-5.1: ML Pipeline Integration . . . 44
3.6.2 BR-5.2: Data Storage and Version Management . . . 45
3.6.3 BR-5.3: Elastic Compute Resource Scheduling . . . 46
3.6.4 BR-5.4: Multi-Environment Deployment Support . . . 47
3.7 Requirements Traceability Matrix . . . 47

4 Canonical Data Schema **49**

4.1 Layout of `src/training/pipeline/` . . . 50
4.2 Training Example Record . . . 50
4.3 Schema Registry . . . 51
4.4 Taxonomy Structure . . . 51
4.5 Enumerations . . . 52
4.6 Unified Schema Registry . . . 52

5 HANA Schema Extraction Pipeline 55

5.1 Pipeline at a Glance . . . 55

5.2 Stage 1 — MCP Discovery . . . 56

5.3 Stage 2 — Table Enumeration . . . 56

5.4 Stage 3 — Column Extraction . . . 57

5.5 Stage 4 — Registry Population . . . 57

5.6 Stage 5 — JSON Export . . . 58

6 Taxonomy Generation Engine 59

6.1 Algorithm 1: Reasoning-Driven Taxonomy Generation . . . 59

6.2 Step 1: Factor Identification . . . 60

6.3 Step 2: Breadth-First Expansion . . . 61

6.4 Step 3: Best-of-N Proposal . . . 61

6.5 Step 4: Critic Refinement . . . 62

6.6 Step 5: Level Planning . . . 63

6.7 Taxonomy Depth and Dataset Size Relationship . . . 63

6.7.1 Node Count Growth . . . 63

6.7.2 Target Examples vs Taxonomy Depth . . . 64

6.7.3 Configuration Guidelines . . . 65

6.7.4 Cost Estimation . . . 65

6.8 Worked Example: Trial Balance Taxonomy . . . 66

7 Training Data Generation Engine **67**

7.1 Algorithm 2: Agentic Data Synthesis . . . 67

7.2 Step 1: Mix Sampling (Global Diversity) . . . 68

7.3 Step 2: Meta-Prompt Generation (Local Diversity) . . . 69

7.4 Step 3: Complexification . . . 69

7.5 Step 4: Example Generation . . . 70

7.6 Step 5: Double-Critic Filtering . . . 70

7.7 Schema Context Pruning . . . 71

7.7.1 Token Budget . . . 72

7.7.2 Pruning Algorithm . . . 72

7.7.3 Relevance Scoring . . . 73

7.7.4 Configuration . . . 74

7.7.5 Quality Impact . . . 74

7.8 Output Format . . . 74

8 Environment and CLI **77**

8.1 Configuration Table . . . 77

8.2 CLI Usage . . . 78

8.3 Output Files . . . 79

9 Evaluation Framework **81**

9.1 Calibrated Complexity Scoring (Algorithm 3) . . . 81

9.1.1 Elo Rating Algorithm . . . 82

9.1.2 Usage . . . 82

9.2 Taxonomic Coverage Evaluation . . . 83

9.3 Embedding-Based Diversity Metrics . . . 84

9.4 Critic Calibration Validation . . . 85

9.5 Student-Teacher Gap Analysis . . . 86

9.6 Adaptive Features (v1.2) . . . 86

9.6.1 Adaptive Complexity Ratio (Figure 7) . . . 86

9.6.2 Adaptive Critic Thresholding (Figure 3b) . . . 87

9.6.3 Automatic Coverage Gap Filling . . . 87

9.6.4 Diversity Optimizer . . . 88

9.6.5 Taxonomy Quality Gate (Section B.1, Table 2) . . . 89

9.7 Evaluation Environment Variables . . . 90

9.8 CI Integration of Evaluation Tolerances . . . 91

10 Worked Example: Trial Balance Domain **93**

10.1 Step 1: Schema Extraction . . . 93

10.2 Step 2: Factor Identification . . . 95

10.3 Step 3: Taxonomy Generation . . . 95

10.4 Step 4: Mix Sampling and Meta-Prompt Generation . . . 96

10.5 Step 5: Example Generation . . . 96

10.6 Step 6: Evaluation Metrics . . . 97

10.6.1 Reproducibility . . . 97

10.7 Step 7: Output Files . . . 98

11 Testing Specification **99**

11.1 Test coverage requirements . . . 99

11.2 Test fixtures . . . 99

11.2.1 Taxonomy fixtures . . . 100

11.2.2 Training example fixtures . . . 100

11.3 Unit test cases . . . 101

11.3.1 Taxonomy traversal . . . 101

11.3.2 Complexity scoring . . . 101

11.3.3 Critic evaluation . . . 102

11.3.4 Diversity metrics . . . 102

11.3.5 Adaptive features . . . 103

11.3.6 HANA type handling . . . 103

11.4 Integration test suite . . . 105

11.4.1 End-to-end scenarios . . . 105

11.5 Performance benchmarks . . . 106

11.5.1 Reference hardware profile . . . 106

11.6 Quality metrics tests . . . 107

11.7 CI/CD integration . . . 107

11.8 Acceptance criteria . . . 108

12 Test Fixture Specification **109**

12.1 Overview . . . 110

12.2 Fixture Count Calculation Methodology . . . 110

12.2.1 Coverage-Based Calculation . . . 110
12.2.2 Fixture Count Derivation by Domain . . . 111
12.2.3 Edge Case Multiplier . . . 111
12.2.4 Extending Fixture Counts . . . 111
12.3 Fixture organisation . . . 112
12.4 Arabic AP fixtures . . . 113
12.4.1 Invoice fixtures . . . 113
12.4.2 OCR fixtures . . . 114
12.5 Trial Balance fixtures . . . 114
12.5.1 TB extract fixtures . . . 114
12.5.2 Variance fixtures . . . 114
12.6 Simula fixtures . . . 115
12.6.1 Taxonomy fixtures . . . 115
12.6.2 Training example fixtures . . . 116
12.7 Common fixtures . . . 116
12.7.1 Mock LLM responses . . . 116
12.7.2 Schema copies . . . 117
12.8 Data sanitisation requirements . . . 117
12.9 Fixture validation . . . 117
12.9.1 Schema validation . . . 118
12.9.2 Coverage verification . . . 118
12.10 Implementation checklist . . . 118

13 Agent Implementation Instructions **121**

13.1 Purpose and audience . . . 121

13.2 Rules of Engagement for AI Agents . . . 122

13.3 Agent Rule-Pack Architecture (.clinerules) . . . 123

13.3.1 Inheritance Hierarchy . . . 123

13.3.2 Safety-Critical Flags . . . 123

13.4 Platform Overview — One Codebase, One Instance . . . 124

13.4.1 Integration Stance . . . 124

13.4.2 Platform Integration Map . . . 124

13.5 Transactional Sequence Diagrams . . . 125

13.5.1 Universal Governance Handshake . . . 125

13.5.2 Arabic AP Golden Path . . . 125

13.6 Compute and Resource Alignment . . . 126

13.7 Human-in-the-Loop (HITL) Critical Gates . . . 126

13.8 Registry Integrity & Live Status . . . 127

13.9 Master Implementation Manifest (Union) . . . 127

13.10 Governance Obligations (MAS MGF) . . . 128

13.11 FMEA: Failure Mode and Effects Analysis for AI . . . 129

13.12 This domain in context . . . 129

13.13 Schema contribution . . . 130

13.14 Persona and MCP allow-list . . . 130

13.15 Agent Rule-Pack Binding . . . 131

13.16 Current Implementation Gap Analysis . . . 132
13.16.1 Key Findings Summary . . . 132
13.16.2 Gap Analysis Validity . . . 133
13.17 Implementation tasks for this domain . . . 133
13.18 Governance obligations applied to this domain . . . 136
13.19 Cross-spec integration — full detail . . . 137
13.19.1 With Arabic AP . . . 137
13.19.2 With Trial Balance . . . 137
13.19.3 With Regulations . . . 137
13.19.4 With Simula . . . 138
13.20 Wave assignment and sequencing . . . 138
13.21 Acceptance criteria (platform-wide) . . . 139
13.22 References . . . 140

14 Cross-Specification Dependencies 141

14.1 Architecture Position . . . 142
14.2 Schema Dependencies . . . 142
14.2.1 Schemas Owned by Simula . . . 142
14.2.2 Schemas Consumed from Other Domains . . . 143
14.3 Regulations Integration . . . 143
14.3.1 Pillar 1: Training Data Bounds . . . 143
14.3.2 Pillar 2: Human Accountability . . . 144
14.3.3 Pillar 3: Technical Controls . . . 144
14.3.4 Pillar 4: Data Provenance . . . 144

14.4 Arabic AP Integration . . . 144

14.5 Trial Balance Integration . . . 145

14.6 Cross-Reference Summary . . . 145

15 Schema Versioning and Interface Contracts 147

15.1 Versioning Policy . . . 147

15.1.1 Semantic Versioning for Schemas . . . 148

15.1.2 Current Schema Versions . . . 148

15.2 Schema Evolution Rules . . . 148

15.2.1 Backward Compatibility Requirements . . . 148

15.2.2 Deprecation Process . . . 149

15.3 Cross-Domain Interface Contracts . . . 149

15.3.1 Schema-In, TrainingExample-Out Contract . . . 149

15.3.2 Interface Versioning . . . 150

15.4 Schema Validation . . . 150

15.4.1 Validation Points . . . 150

15.4.2 Validation Implementation . . . 151

15.5 Changelog Management . . . 151

15.5.1 Schema Changelog Location . . . 151

15.5.2 Changelog Entry Format . . . 152

15.6 Version Negotiation . . . 152

15.6.1 Runtime Version Checking . . . 152

15.6.2 Graceful Degradation . . . 153

15.7 Implementation Checklist . . . 153

15.8 Governance . . . 154

15.8.1 Schema Change Approval . . . 154

15.8.2 Change Ownership . . . 154

16 Project Roadmap **155**

Scope of this chapter . . . 155

16.1 MoSCoW classification . . . 156

16.2 Dependency graph . . . 156

16.3 Project milestones . . . 156

16.4 Launch tiers and exit gates . . . 159

16.5 Kill-switch conditions . . . 159

16.6 Shared-platform dependencies . . . 162

16.7 v1.3 expansion trigger . . . 162

16.8 Governance and business-case cadence . . . 163

17 Implementation Roadmap **165**

Integration with existing chapters . . . 165

Scope of this chapter . . . 166

Relationship to Project Milestones . . . 166

17.1 Roadmap Overview . . . 167

17.1.1 Phase Dependencies . . . 168

17.2 Phase 1: Foundation Capability Building . . . 168

17.2.1 Phase 1 Objectives . . . 168
17.2.2 Phase 1 Deliverables . . . 169
17.2.3 Phase 1 Exit Criteria . . . 174
17.3 Phase 2: Scaled Deployment . . . 175
17.3.1 Phase 2 Objectives . . . 175
17.3.2 Phase 2 Deliverables . . . 175
17.3.3 Phase 2 Exit Criteria . . . 181
17.4 Phase 3: Mature Operations and Continuous Optimization . . . 182
17.4.1 Phase 3 Objectives . . . 182
17.4.2 Phase 3 Deliverables . . . 182
17.4.3 Phase 3 Exit Criteria . . . 188
17.5 Key Success Factors . . . 189
17.5.1 Critical Success Factors . . . 189
17.5.2 Risk Mitigation Summary . . . 189
17.6 Requirements Coverage by Phase . . . 190
17.7 SAP Technology Stack Summary . . . 190

18 Future Directions (v1.3+) 193

Integration with existing chapters . . . 194
Scope of this chapter . . . 195
18.1 Enhancement Overview . . . 195
18.2 Enhancement 1: Multi-Modal Extension . . . 196
18.2.1 Current Limitation . . . 196
18.2.2 Cross-Modal Reasoning Scaffolds . . . 196

18.2.3 Modality-Aware Complexification . . . 197
18.2.4 SAP Integration: Document Information Extraction . . . 198
18.3 Enhancement 2: Domain Adaptation Calibration . . . 199
18.3.1 Current Limitation . . . 199
18.3.2 Human-in-Loop Knowledge Anchors . . . 199
18.3.3 Vertical-Specific Complexity Metrics . . . 201
18.3.4 SAP Integration: S/4HANA Domain Schemas . . . 201
18.4 Enhancement 3: Closed-Loop Validation . . . 202
18.4.1 Current Limitation . . . 202
18.4.2 Training-Validation Feedback Loop . . . 203
18.4.3 Adversarial Verification Layer . . . 205
18.4.4 SAP Integration: AI Core MLOps . . . 206
18.5 Enhancement 4: Programmable Workflow Engine . . . 207
18.5.1 Current Limitation . . . 207
18.5.2 Declarative Dataset Specification Language . . . 207
18.5.3 Version Control and Reproducibility . . . 209
18.5.4 SAP Integration: BTP and HANA Cloud . . . 210
18.6 Regulatory Framework Alignment . . . 211
18.6.1 UK FCA Synthetic Data Expert Group . . . 211
18.6.2 IEEE Synthetic Data White Paper . . . 212
18.6.3 World Economic Forum Guidelines . . . 212
18.7 Implementation Checklist . . . 213

A Implementation Gap Review and Agent Rule Pack **215**

A.1 Review scope and constraints . . . 216
A.2 Executive assessment . . . 217
A.3 Validation performed . . . 217
A.4 Gap-to-Milestone Mapping Summary . . . 218
A.5 Findings in `src/training` . . . 218
A.6 Findings in `src/generativeUI/training-webcomponents-ngx` . . . 225
A.7 Proposed implementation-agent rule file . . . 229
A.8 Proposed runtime-monitor agent rule file . . . 230
A.9 Sequenced remediation order . . . 231

B SAP Integration Guide **233**

B.1 SAP HANA Cloud Integration . . . 233
B.1.1 Connection Configuration . . . 233
B.1.2 Required Database Permissions . . . 234
B.1.3 HANA-Specific Type Mapping . . . 235
B.1.4 Spatial Type Handling . . . 236
B.2 SAP AI Core Integration . . . 237
B.2.1 Deployment Configuration . . . 237
B.2.2 Training Pipeline Integration . . . 238
B.3 SAP BTP Services Integration . . . 239
B.3.1 Credential Store . . . 239
B.3.2 Audit Log Service . . . 240
B.3.3 Alert Notification Service . . . 241

B.4 SAP Analytics Cloud Integration . . . 242
B.4.1 Cost Tracking Dashboard . . . 242
B.4.2 Recommended Dashboard Widgets . . . 243
B.5 SAP GRC Integration . . . 243
B.5.1 Policy Management . . . 243
B.5.2 Compliance Evidence Collection . . . 244
B.6 SAP Build Apps Integration . . . 245
B.6.1 API Endpoints . . . 245
B.6.2 Build Apps Data Entity . . . 245
B.7 Quick Reference: Environment Variables . . . 247

C References **249**
C.1 Source Documents . . . 249
C.2 Implementation Files . . . 250
C.3 Tools and Dependencies . . . 251
C.4 Academic References . . . 251
C.5 Related Specifications . . . 252
C.6 Schema Registry . . . 252
C.7 External Standards . . . 253

D Global Wave-Dependency Manifest **255**
D.1 Critical Path Visualization (Waves 1–3) . . . 255
D.2 Inter-Domain Prerequisite Matrix . . . 256

Chapter 1

Executive Summary

Frontmatter

Executive Summary

This specification defines the **Simula Training Data Framework**—a comprehensive system for generating high-quality synthetic training data to enhance SAP AI agents' reasoning capabilities. The framework transforms production schemas and business logic into structured training examples suitable for supervised fine-tuning (SFT) and reinforcement learning (GRPO).

◎ Primary Objectives

1. **Schema-Driven Generation**: Extract taxonomy from HANA Cloud schemas via MCP
2. **Multi-Turn Reasoning**: Generate progressive reasoning chains with verification
3. **Quality Assurance**: Implement double-critic filtering and diversity scoring
4. **SAP Integration**: Native compatibility with AI Core, BTP, and GRC platforms

Key Deliverables

- **TrainingExample Schema**: Canonical data model for all generated examples
- **Taxonomy Builder**: Algorithm for extracting domain concepts from HANA schemas
- **Generation Engine**: Multi-turn reasoning generator with critic validation
- **Evaluation Framework**: Elo-based quality scoring and coverage metrics
- **CLI Tools**: Reproducible pipeline for dataset generation

Document Conventions

Requirement Classification

All requirements use MoSCoW prioritization:

Priority	Symbol	Definition
MUST	❗	Non-negotiable for v1.0 release
SHOULD	✅	Important but can slip to v1.1
COULD	❓	Nice to have if time permits
WON'T	❌	Explicitly deferred (this phase)

Table 1.1: MoSCoW Priority Definitions

Colored Boxes

This specification uses colored boxes to highlight different types of content:

> *"Source quotations from research papers or reference documents appear in green boxes with quotation marks."*

Requirement: D

erived requirements appear in blue boxes and include formal requirement IDs.

Implementation Note

Implementation notes provide technical guidance for developers in purple boxes.

Warning

Warnings highlight critical constraints or potential pitfalls in red boxes.

Future Work

Future enhancements (v1.3+) are marked in gold boxes to distinguish them from current scope.

Requirement ID Conventions

How to Read This Specification

This specification is organized into five parts plus appendices:

1. **Part I: Introduction and Context** – Establishes the problem domain, business requirements, and stakeholder needs

Prefix	Category	Example
SIM-DATA	Data Schema	SIM-DATA-001: TrainingExample format
SIM-TAX	Taxonomy Engine	SIM-TAX-001: Concept extraction
SIM-GEN	Generation Engine	SIM-GEN-001: Multi-turn generation
SIM-EVAL	Evaluation	SIM-EVAL-001: Elo scoring
SIM-CLI	CLI/Pipeline	SIM-CLI-001: Build command
SIM-INT	Integration	SIM-INT-001: HANA connectivity
SIM-HANA	HANA-specific	SIM-HANA-001: Schema extraction

Table 1.2: Requirement ID Prefixes

2. **Part II: Architecture and Core Design** – Defines data schemas, extraction pipelines, and generation algorithms
3. **Part III: Quality Assurance and Evaluation** – Covers evaluation metrics, testing, and worked examples
4. **Part IV: Implementation and Governance** – Provides implementation guidance and cross-specification dependencies
5. **Part V: Planning and Future Evolution** – Presents roadmaps and future enhancement plans
6. **Appendices** – Contains the implementation gap review, SAP integration guide, and references

Reading Guide

Chapter Relationship Flowchart

The following flowchart shows how the planning and future directions chapters (15–17) relate to other parts of the specification. This centralized view reduces the need to repeat integration context within each chapter.

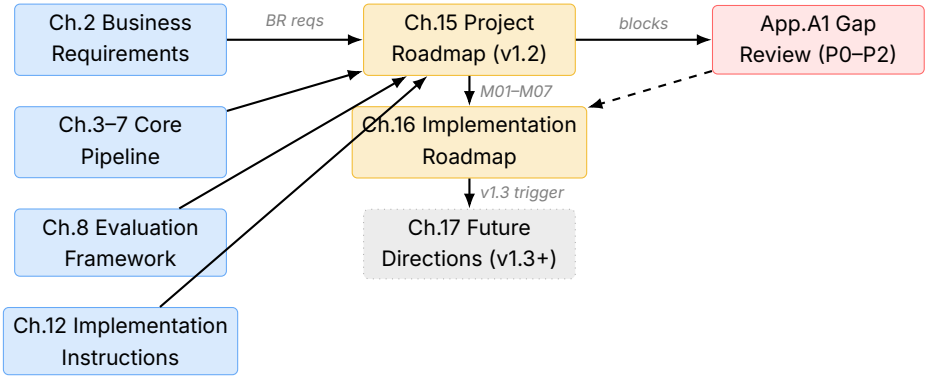

Figure 1.1: Chapter relationship flowchart. Blue: core specification. Gold: planning chapters (v1.2 scope). Dotted gray: future scope (v1.3+). Red: gap review. Solid arrows: hard prerequisites. Dashed arrows: remediation dependencies.

Key Reading Paths:

- **Understanding v1.2 scope**: Ch.2 → Ch.15 → App.A1 → Ch.16
- **Gap remediation**: App.A1 (P0 gaps) → Ch.15 (milestones blocked) → Ch.16 (enterprise phases)
- **Future planning**: Ch.16 (Phase 3 complete) → Ch.17 (v1.3+ enhancements)

Audience-Specific Paths

For Project Managers / Business Owners

1. Frontmatter (this section)
2. Chapter 1: Overview and Scope
3. Chapter 2: Business Requirements
4. Chapter 15: Project Roadmap
5. Chapter 16: Implementation Roadmap

For Implementers / Developers

1. Chapter 1: Overview (full)
2. Chapter 3: Canonical Data Schema
3. Chapters 4–7: Core Pipeline (Extraction, Taxonomy, Generation, CLI)
4. Chapter 12: Implementation Instructions
5. Chapter 10: Testing Specification
6. Appendix B: SAP Integration Guide

For Quality / Evaluation Engineers

1. Chapter 8: Evaluation Framework
2. Chapter 9: Worked Example
3. Chapter 11: Test Fixtures
4. Chapter 10: Testing Specification
5. Appendix A: Implementation Gap Review

For Auditors / Compliance Officers

1. Chapter 1: Overview (Sections 1.3–1.5 for research traceability)
2. Chapter 2: Business Requirements
3. Chapter 14: Schema Versioning
4. Chapter 13: Cross-Spec Dependencies
5. Appendix C: References

Version History

Version	Date	Changes
1.0.0	2026-04-01	Initial specification based on Simula research paper
1.1.0	2026-04-10	Added evaluation framework and CLI specification
1.2.0	2026-04-21	Added SAP integration guide, business requirements, and implementation roadmap
1.2.1	2026-04-21	Reorganized chapter structure for improved readability; moved business requirements early; added Part divisions; moved gap review to appendix
1.2.2	2026-04-21	Added Chapter Relationship Flowchart in Frontmatter (Figure 1.1); added Gap-to-Milestone mapping table in Appendix A1 (Table A.2)

Document Scope Matrix

Topic	In Scope	Chapter(s)	Notes
TrainingExample schema	✔	3	Canonical data model
Taxonomy extraction	✔	4, 5	Algorithm 1
Multi-turn generation	✔	6	Algorithm 2
Double-critic filtering	✔	6	Quality gate
Elo-based evaluation	✔	8	Quality metrics
SFT training	✔	3	Output format
GRPO training	✔	3	Reward format
HANA Cloud integration	✔	4, App. B	MCP-based extraction
Multi-modal input	✖	17	Future (v1.3+)
Real-time streaming	✖	–	Out of scope
Model training execution	✖	–	Downstream concern

Table 1.4: Document Scope Matrix

Chapter 2

Overview and Scope

2.1 The Problem

The SAP OSS training pipeline requires high-quality Text-to-SQL training data that reflects the actual schema structure of SAP HANA Cloud tables. The current implementation in `src/training/pipeline/` uses metadata from **Excel files** (`DATA_DICTIONARY.xlsx`, `ESG_DATA_DICTIONARY.xlsx`) that were originally exported from HANA or defined manually.

While the `HanaClient` in `src/training/pipeline/hana_client.py` has the capability to query `SYS.TABLE_COLUMNS` for live schema metadata, this HANA-direct extraction is **not currently integrated** into the main training data generation pipeline.

Table 2.1: Current vs. Target Training Data Pipeline

Dimension	Current State	Target State
Schema Source	Static Excel/CSV files	Live HANA Cloud metadata
Extraction Method	Manual export	MCP-based automated query
Training Data	Template-based expansion	Reasoning-driven synthesis
Diversity Control	None	Taxonomic coverage
Quality Assurance	None	Double-critic filtering
Complexity Control	None	Calibrated Elo scoring

2.2 Goals

The specification targets three outcomes:

1. **HANA-Direct Extraction.** Every training dataset is generated from live schema metadata queried via the AI Core PAL MCP service.
2. **Taxonomic Coverage.** The conceptual space of the target domain is mapped using hierarchical taxonomies, ensuring global diversity.
3. **Reasoning-Driven Synthesis.** Each training example is generated from first principles using agentic refinement, ensuring quality, local diversity, and controlled complexity.

2.3 In-Scope Domains

Table 2.2: Application Domains

Domain	HANA Schema	Training Task
Arabic Document Processing	`AP_INVOICES`, `VAT_CHECKLISTS`	Invoice field extraction, VAT compliance
Trial Balance Review	`TB_VARIANCES`, `GL_BALANCES`	Variance analysis, commentary generation
ESG Analytics	`ESG_METRICS`, `FINANCED_EMISSIONS`	Scope classification, net-zero alignment

2.4 System Context

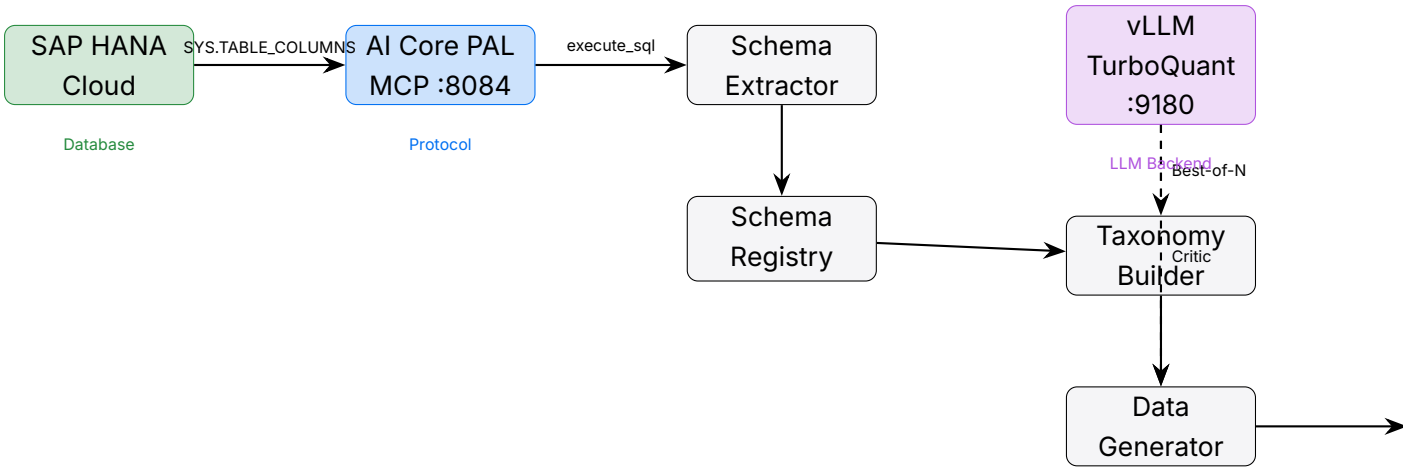

Figure 2.1: System context. HANA Cloud (green) provides live schema metadata via the MCP protocol (blue). The vLLM backend (purple) powers taxonomy expansion and data generation. Output is JSONL training data (gold).

2.5 Research Paper Traceability

This specification implements the Simula framework from Davidson et al. (2026), "Reasoning-Driven Synthetic Data Generation and Evaluation," published in Transactions on Machine Learning Research (TMLR). The following table maps paper sections to specification sections:

Table 2.3: Traceability: Research Paper to Implementation

Paper Reference	Paper Finding	Spec Implementation
Section 2.1, Figure 1	Taxonomy-based coverage	§6, Algorithm 1
Section 2.2, Algorithm 2	Agentic data synthesis	§7, Algorithm 2
Section 2.3, Algorithm 3	Calibrated complexity scoring	§9.1, Elo algorithm
Section 3.1, Figure 3	Double-critic validation	§7.5, §9.4
Section 3.3, Figure 4	Additive diversity	§9.3, Diversity analyzer
Section 4.3	83% saturation ratio	§9.5, Gap analysis
Figure 7	Complexity vs teacher strength	§9.6.1, Adaptive ratio
Table 2 (Appendix B)	Taxonomy quality metrics	§9.6.5, Quality gate
Appendix E.3	Elo algorithm details	§9.1.1, `_update_elo()`

2.6 Out of Scope

The following are **explicitly** out of scope for this document:

- Model fine-tuning and evaluation (covered by separate LoRA training specification).
- Downstream inference and routing (covered by `03-prompt-routing.tex`).
- Trial Balance business process (covered by `04-trial-balance-ai-specification.tex`).
- Arabic invoice pipeline (covered by the Arabic AP specification).

Chapter 3

Business Requirements Specification

Integration with existing chapters

This chapter provides the business governance layer around the technical specifications in earlier chapters:

- **Chapter 16** (Project Roadmap): BR-1.x strategic requirements align with the MoSCoW classification in Table 16.1. High-priority BRs map to Must milestones

(M01–M07); Medium-priority BRs map to Should (S01–S02); Low-priority BRs map to Could (C01).

- **Chapter 13** (Agent Instructions): BR-2.x generation requirements provide acceptance criteria for Tasks S1–S4 (schema extraction, taxonomy building, generation, evaluation).
- **Chapter ??** (Enterprise Enhancements): BR-3.x quality requirements are implemented by Enhancement 3 (Closed-Loop Validation); BR-4.x governance requirements by Enhancement 4 (Workflow Engine).
- **Table 13.3** (FMEA): BR-3.3 (Privacy) and BR-3.4 (Bias) address the "Critic Collusion" and "Schema Contraction" failure modes.

The traceability matrix at the end of this chapter maps each BR to the enhancement and regulatory framework that governs it.

Scope of this chapter

This chapter defines the complete business requirements for enterprise deployment of the Simula synthetic data framework. Requirements are organized across five categories: strategic preparation, generation phase, quality assurance, governance and compliance, and integration infrastructure. Each requirement includes regulatory citations where applicable.

3.1 Requirements Overview

3.1.1 Requirement Identification Convention

Requirements follow the identification pattern `BR-<category>.<sequence>`:

Table 3.1: Business requirement categories.

Category	Prefix	Scope	Count
Strategic & Preparation	BR-1.x	Use case definition, governance integration	5
Generation Phase	BR-2.x	Taxonomy, configuration, metadata	5
Quality Assurance	BR-3.x	Validation, privacy, bias detection	5
Governance & Compliance	BR-4.x	Access control, audit, transparency	5
Integration & Infrastructure	BR-5.x	ML pipelines, storage, deployment	5

3.1.2 Priority Classification

Requirements are classified using MoSCoW prioritization:

- **High**: Must have for initial deployment
- **Medium**: Should have for production maturity
- **Low**: Could have for optimization

3.2 BR-1: Strategic and Preparation Phase

3.2.1 BR-1.1: Use Case Definition

BR-1.1: Clear Use Case Definition

Priority: High

Description: Organizations must clearly define synthetic data use cases and objectives before initiating generation projects. The synthetic data market is projected to reach $3.4 billion by 2031[a], but success depends on starting with high-value, well-scoped use cases.

Typical Use Cases:

- Privacy-sensitive domain data substitution
- Long-tail scenario coverage augmentation
- Adversarial testing data generation
- Training data augmentation for rare events
- Model validation in regulated environments

Acceptance Criteria:

1. Documented use case with clear business objective
2. Identified target domain and data modality
3. Defined success metrics and measurement approach
4. Stakeholder sign-off on scope and expectations

SAP Integration: Align use cases with SAP AI Core deployment scenarios (Text-to-SQL, document extraction, financial analytics).

[a] Market Research Future, "Synthetic Data Generation Market Report," 2024

3.2.2 BR-1.2: Value-Risk Assessment

BR-1.2: Structured Value-Risk Assessment

Priority: High

Description: Before initiating synthetic data projects, organizations must conduct structured value-risk assessments. The UK FCA Synthetic Data Expert Group explicitly identifies clear objective definition and value-risk analysis as foundational to project success[a].

Assessment Dimensions:

- **Value**: Cost savings, time-to-market, compliance enablement
- **Risk**: Privacy leakage, bias amplification, quality degradation
- **Feasibility**: Technical capability, resource availability
- **Regulatory**: Jurisdiction-specific compliance requirements

Acceptance Criteria:

1. Completed value-risk assessment matrix
2. Identified risk mitigation strategies
3. Documented go/no-go decision criteria
4. Executive sponsor approval

[a] FCA, "Synthetic Data Expert Group Report," Section 3.2, 2024

3.2.3 BR-1.3: Governance Framework Integration

BR-1.3: Existing Governance Framework Integration

Priority: High

Description: In the absence of dedicated synthetic data governance frameworks, organizations should integrate synthetic data management into existing model risk management, data governance, and AI ethics frameworks rather than building from scratch.

Integration Points:

- Model Risk Management (MRM) — treat synthetic data as model input
- Data Governance — apply data quality and lineage standards
- AI Ethics Framework — extend bias and fairness requirements
- Information Security — apply data classification policies

Acceptance Criteria:

1. Synthetic data policies integrated into existing governance documentation
2. Roles and responsibilities defined within existing RACI matrices
3. Approval workflows configured in existing governance tools
4. Training materials updated for affected stakeholders

SAP Integration: Leverage SAP GRC (Governance, Risk, Compliance) for policy management and workflow orchestration.

3.2.4 BR-1.4: Cross-Functional Team Formation

BR-1.4: Cross-Functional Team Organization

Priority: Medium

Description: Synthetic data generation and usage requires multi-stakeholder collaboration. The World Economic Forum notes that synthetic data governance is typically constructed by multiple distinct teams working together[a].

Required Roles:

- **Data Science Team**: Technical implementation, model training
- **Domain Experts**: Taxonomy validation, complexity calibration
- **Legal/Compliance**: Privacy risk assessment, regulatory alignment
- **IT/Platform**: Infrastructure provisioning, security controls
- **Business Stakeholders**: Use case definition, acceptance criteria

Acceptance Criteria:

1. Documented team structure with named individuals
2. Defined communication and escalation paths
3. Regular sync cadence established
4. Clear decision-making authority matrix

[a] WEF, "Synthetic Data for Financial Services," p. 18, 2024

3.2.5 BR-1.5: ROI Baseline Establishment

BR-1.5: Investment Return Baseline

Priority: Medium

Description: Organizations should establish clear ROI measurement baselines. Industry data indicates typical benefits include: AI model development speed improvements up to 40%, compliance and data processing cost reductions of 50–60%[a].

Measurement Dimensions:

- Time-to-data: Days from request to available training data
- Data acquisition cost: Cost per training example
- Model performance: Accuracy, F1, or domain-specific metrics
- Compliance cost: Audit preparation, privacy review cycles

Acceptance Criteria:

1. Baseline metrics captured before synthetic data adoption
2. Target improvement thresholds defined
3. Measurement methodology documented
4. Reporting cadence established

[a] Gartner, "Synthetic Data Adoption Trends," 2025

3.3 BR-2: Generation Phase

3.3.1 BR-2.1: Domain Taxonomy Construction and Validation

BR-2.1: Taxonomy Construction and Validation

Priority: High
Description: Simula's global diversification step requires constructing deep hierarchical taxonomies for target domains. Business requirements include domain expert participation in taxonomy review and version management as domain knowledge evolves.
Requirements:

- Expert review interface for taxonomy inspection and correction
- Version control for taxonomy artifacts
- Coverage gap identification and reporting
- Domain vocabulary injection mechanism

Acceptance Criteria:

1. Taxonomy review workflow with expert approval gates
2. Taxonomy versioning with diff visualization
3. Coverage reports showing sampling breadth and depth
4. Domain-specific node validation rules

SAP Integration: Store taxonomy versions in SAP HANA Cloud with metadata tracking. Integrate with SAP Signavio for process-aware taxonomy construction.

3.3.2 BR-2.2: Controllable Parameter Configuration

BR-2.2: Parameter Configuration Interface

Priority: High
Description: Provide business user-friendly interfaces for independent control of generation parameters without requiring technical expertise.
Controllable Parameters:

- **Global Coverage**: Taxonomy sampling breadth and depth
- **Local Diversity**: Variants per concept node
- **Complexity Distribution**: Easy/medium/hard ratio curve
- **Quality Thresholds**: Critic pass rate, diversity minimums

Acceptance Criteria:

1. Visual configuration interface with parameter sliders
2. Preview mode showing expected distribution
3. Parameter validation with constraint warnings
4. Configuration templates for common scenarios

SAP Integration: Expose configuration via SAP Build Apps for low-code workflow construction. Store configurations in SAP HANA Cloud.

3.3.3 BR-2.3: Cost and Resource Budget Control

BR-2.3: Budget Control Mechanism

Priority: Medium
Description: Synthetic data generation involves inference model API costs. Organizations require budget control mechanisms to manage expenses.
Requirements:

- Per-project and per-department generation quotas
- Single-task cost ceiling configuration
- Resource consumption estimation before execution
- Low-cost "draft mode" for taxonomy design validation

Acceptance Criteria:

1. Budget allocation interface by cost center
2. Real-time cost tracking during generation
3. Alerts when approaching budget limits
4. Cost-per-example reporting

SAP Integration: Integrate with SAP AI Core resource quotas and SAP Analytics Cloud for cost visualization.

3.3.4 BR-2.4: Metadata and Provenance Recording

BR-2.4: Complete Metadata Capture

Priority: High
Description: Each generation task must automatically record complete metadata for audit compliance and reproducibility.
Required Metadata:

- Reasoning model version and configuration
- Taxonomy structure snapshot
- Sampling strategy parameters
- Generation timestamp and executor identity
- Random seeds for reproducibility

Acceptance Criteria:

1. Automatic metadata capture on every generation run
2. Metadata stored with tamper-evident integrity
3. Query interface for metadata retrieval
4. Export capability for audit packages

SAP Integration: Store provenance in SAP HANA Cloud with blockchain-style hash chaining. Integrate with SAP Data Intelligence for lineage visualization.

3.3.5 BR-2.5: Incremental Generation and Dataset Evolution

BR-2.5: Incremental Generation Support

Priority: Low
Description: Support incremental generation on existing datasets when taxonomies evolve, avoiding redundant regeneration.
Requirements:

- Detect new/modified taxonomy nodes
- Generate only for changed coverage regions
- Merge incremental examples with existing dataset
- Version-to-version diff analysis

Acceptance Criteria:

1. Taxonomy diff computation between versions
2. Targeted generation for changed nodes only
3. Dataset merge with deduplication
4. Generation efficiency metrics (examples saved)

3.4 BR-3: Quality Assurance and Validation

3.4.1 BR-3.1: Intrinsic Quality Assessment

BR-3.1: Standardized Quality Reporting

Priority: High
Description: Simula's built-in quality assessment capabilities must be productized into standardized reports for business consumption.
Required Reports:

- **Coverage Report**: Taxonomy branches sampled vs. omitted
- **Complexity Distribution**: Generated data complexity vs. target
- **Diversity Metrics**: Intra-concept sample similarity analysis
- **Critic Pass Rates**: Acceptance rates by taxonomy region

Acceptance Criteria:

1. Automated report generation on every dataset
2. Visual dashboards with drill-down capability
3. Threshold alerts for quality degradation
4. Historical trend analysis

SAP Integration: Publish reports to SAP Analytics Cloud. Configure alerts via SAP Alert Notification Service.

3.4.2 BR-3.2: Downstream Task Validation

BR-3.2: Downstream Model Validation

Priority: High
Description: Generated datasets must be validated on downstream models to verify practical utility. Recommended approach: "Train-Synthetic-Test-Real" (TSTR) methodology.
Validation Protocol:

1. Train model on synthetic data
2. Evaluate on held-out real data
3. Compute synthetic-to-real performance gap
4. Identify systematic failure patterns

Acceptance Criteria:

1. Automated TSTR pipeline integration
2. Performance gap quantification by taxonomy region
3. Failure pattern analysis reports
4. Threshold-based quality gates

SAP Integration: Execute validation via SAP AI Core training pipelines. Store results in SAP HANA Cloud for trend analysis.

3.4.3 BR-3.3: Privacy Risk Assessment

BR-3.3: Privacy Risk Evaluation

Priority: High
Description: Even in Simula's seedless mode, generated data may exhibit unintended similarity to real data seen during reasoning model training. Privacy risk assessment mechanisms are essential. IEEE emphasizes that privacy risk metrics and mitigation techniques are core governance components[a].
Assessment Mechanisms:

- Membership inference attack testing
- Similarity scanning against known real datasets
- Differential privacy budget tracking
- Re-identification risk scoring

Acceptance Criteria:

1. Automated privacy risk scoring on generated datasets
2. Configurable risk thresholds with blocking gates
3. Audit-ready privacy assessment reports
4. Remediation workflow for high-risk samples

SAP Integration: Integrate with SAP Data Privacy Integration service for sensitive data detection.

[a] IEEE, "Synthetic Data: Generation, Validation, and Governance," Section 4.3, 2025

3.4.4 BR-3.4: Bias Detection and Mitigation

BR-3.4: Systematic Bias Detection

Priority: Medium
Description: Establish systematic bias detection processes to identify and mitigate inherited biases from reasoning models.
Detection Approaches:

- Distribution analysis on sensitive attributes
- Reasoning model bias inheritance detection
- Over/under-sampling recommendations
- Fairness metric computation

Acceptance Criteria:

1. Automated bias scanning on generation
2. Fairness metric dashboards
3. Bias mitigation tool integration
4. Bias audit trail for compliance

3.4.5 BR-3.5: Continuous Monitoring and Drift Detection

BR-3.5: Quality Drift Monitoring

Priority: Medium
Description: Establish continuous monitoring to detect quality changes due to reasoning model upgrades or taxonomy drift. FCA includes "continuous monitoring and improvement" as one of nine governance principles[a].
Monitoring Capabilities:

- Periodic quality sampling and assessment
- Version-to-version quality drift detection
- Model upgrade impact analysis
- Automated regeneration triggers

Acceptance Criteria:

1. Scheduled quality monitoring jobs
2. Drift detection with statistical significance testing
3. Alert configuration for threshold breaches
4. Automated remediation workflow triggers

SAP Integration: Configure monitoring via SAP AI Core MLOps capabilities. Route alerts through SAP Alert Notification Service.

[a] FCA, "Synthetic Data Expert Group Report," Principle 9, 2024

3.5 BR-4: Governance and Compliance

3.5.1 BR-4.1: Access Control and Permission Management

BR-4.1: Role-Based Access Control

Priority: High
Description: Although synthetic data does not directly contain real personal information, strict access controls remain necessary.
Control Requirements:

- Role-based access (generator, reviewer, consumer separation)
- Dataset access approval workflows
- Additional restrictions for sensitive domains (healthcare, finance)
- Access logging and audit trails

Acceptance Criteria:

1. RBAC configuration in identity provider
2. Approval workflow for dataset access requests
3. Domain-specific access policies
4. Comprehensive access audit logs

SAP Integration: Leverage SAP Cloud Identity Services for authentication. Configure authorization via SAP Authorization and Trust Management Service (XSUAA).

3.5.2 BR-4.2: Regulatory Framework Alignment

BR-4.2: Compliance Framework Alignment

Priority: High

Description: Synthetic data generation and usage must align with applicable regulations. FCA emphasizes that generating and using synthetic data requires consideration of data protection, non-discrimination, and bias regulatory obligations[d].

Regulatory Considerations:

- **GDPR**: Demonstrate anonymization status of synthetic data
- **HIPAA**: Comply with protected health information rules
- **Industry-Specific**: Financial services (FCA), healthcare (FDA)
- **AI Regulations**: EU AI Act risk categorization

Acceptance Criteria:

1. Documented compliance position for each applicable regulation
2. Legal review sign-off on synthetic data usage
3. Regulatory change monitoring process
4. Jurisdiction-specific policy variations

SAP Integration: Integrate with SAP Compliance and GRC solutions for regulatory tracking.

[d] FCA, "Synthetic Data Expert Group Report," Section 5.1, 2024

3.5.3 BR-4.3: Audit Logging and Traceability

BR-4.3: Complete Audit Trail

Priority: High
Description: Establish complete audit tracking. WEF emphasizes that enterprises must prioritize oversight and compliance, building robust traceability and provenance systems[a].
Audit Requirements:

- "Who, when, why, what" for every data access
- Complete usage chain (generation → training → deployment)
- Compliance review decision records
- Long-term audit log retention

Acceptance Criteria:

1. Tamper-evident audit log storage
2. Query interface for audit investigations
3. Retention policy enforcement
4. Audit export for regulatory submissions

SAP Integration: Utilize SAP Audit Log Service for tamper-evident logging. Store long-term archives in SAP Object Store with retention policies.

[a] WEF, "Synthetic Data for Financial Services," p. 24, 2024

3.5.4 BR-4.4: Transparency and Explainability

BR-4.4: External Transparency

Priority: Medium

Description: Organizations must be able to explain synthetic data usage to external stakeholders (customers, regulators, partners). This addresses FCA principles of "transparency" and "explainability"[a].

Transparency Requirements:

- Documentation of synthetic nature and generation method
- Explanation of synthetic vs. real data differences
- Justification for synthetic data usage
- Model card integration for downstream models

Acceptance Criteria:

1. Synthetic data documentation templates
2. Stakeholder communication guidelines
3. Model card integration with synthetic data lineage
4. FAQ materials for common inquiries

[a] FCA, "Synthetic Data Expert Group Report," Principles 1–2, 2024

3.5.5 BR-4.5: Incident Response and Remediation

BR-4.5: Incident Response Process

Priority: Medium
Description: Establish incident response processes for synthetic data related security events.
Incident Scenarios:

- Generated data inadvertently contains real information
- Systematic bias causes discriminatory model outputs
- Generated data is misused or leaked
- Quality degradation impacts downstream systems

Acceptance Criteria:

1. Incident classification and severity matrix
2. Responsible party identification
3. Escalation and notification paths
4. Remediation playbooks

SAP Integration: Integrate with existing SAP incident management workflows. Configure alerts through SAP Alert Notification Service.

3.6 BR-5: Integration and Infrastructure

3.6.1 BR-5.1: ML Pipeline Integration

BR-5.1: MLOps Pipeline Integration

Priority: High
Description: Simula must integrate as an orchestratable component within existing ML pipelines.
Integration Requirements:

- API-driven generation task triggering
- Integration with major MLOps platforms
- Automatic dataset push to training storage
- Pipeline status and progress reporting

Acceptance Criteria:

1. REST/gRPC API for generation tasks
2. Pre-built integrations (Kubeflow, MLflow, Vertex AI)
3. Webhook notifications for completion events
4. Pipeline artifact registration

SAP Integration: Primary integration with SAP AI Core. Secondary integrations with SAP Data Intelligence and SAP BTP Kyma Runtime for containerized execution.

3.6.2 BR-5.2: Data Storage and Version Management

BR-5.2: Unified Data Governance

Priority: Medium
Description: Synthetic datasets should follow the same governance policies as real datasets within organizational data infrastructure.
Storage Requirements:

- Version-controlled data lake/warehouse storage
- Dataset labeling, search, and reuse support
- Data catalog integration
- Retention policy enforcement

Acceptance Criteria:

1. Storage in enterprise data lake with versioning
2. Metadata registration in data catalog
3. Search and discovery interface
4. Automated retention enforcement

SAP Integration: Store datasets in SAP HANA Cloud or SAP Datasphere. Register metadata in SAP Data Intelligence catalog.

3.6.3 BR-5.3: Elastic Compute Resource Scheduling

BR-5.3: Compute Resource Management

Priority: Medium
Description: Synthetic data generation is compute-intensive. Support dynamic resource allocation based on workload.
Scheduling Requirements:

- Dynamic resource allocation by task size
- Batch and real-time generation modes
- Priority queue management
- Cost-optimized resource selection

Acceptance Criteria:

1. Auto-scaling based on queue depth
2. Priority scheduling with preemption
3. Resource utilization dashboards
4. Cost optimization recommendations

SAP Integration: Leverage SAP AI Core auto-scaling. Use SAP BTP Kubernetes runtime for containerized workloads.

3.6.4 BR-5.4: Multi-Environment Deployment Support

BR-5.4: Deployment Flexibility

Priority: Low
Description: Support multiple deployment models to accommodate varying security requirements across organizations.
Deployment Models:

- **Public Cloud SaaS**: Standard scenarios
- **Private Deployment**: High-security requirements
- **VPC Deployment**: Regulated industry compliance
- **Hybrid**: Combination of above

Acceptance Criteria:

1. Documented deployment architectures for each model
2. Security certifications per deployment type
3. Network isolation configuration guides
4. Data residency compliance options

SAP Integration: Leverage SAP BTP multi-environment support. Configure SAP Private Link for secure connectivity.

3.7 Requirements Traceability Matrix

Legend: E2 = Domain Calibration, E3 = Closed-Loop Validation, E4 = Workflow Engine

Table 3.2: Requirements to enhancement traceability.

Requirement	Priority	Enhancement	Regulatory Reference
BR-1.1	High	—	—
BR-1.2	High	—	FCA Principle 3
BR-1.3	High	E4	WEF Governance
BR-1.4	Medium	—	WEF Cross-functional
BR-1.5	Medium	E3	—
BR-2.1	High	E2	—
BR-2.2	High	E4	—
BR-2.3	Medium	E4	—
BR-2.4	High	E4	FCA Principle 7
BR-2.5	Low	E4	—
BR-3.1	High	E3	—
BR-3.2	High	E3	IEEE Fidelity
BR-3.3	High	E3	IEEE Privacy, FCA Principle 6
BR-3.4	Medium	E3	FCA Principle 8
BR-3.5	Medium	E3	FCA Principle 9
BR-4.1	High	E4	FCA Principle 6
BR-4.2	High	—	GDPR, HIPAA, FCA
BR-4.3	High	E4	WEF Traceability
BR-4.4	Medium	E4	FCA Principles 1–2
BR-4.5	Medium	—	—
BR-5.1	High	E3, E4	—
BR-5.2	Medium	E4	IEEE Provenance
BR-5.3	Medium	—	—
BR-5.4	Low	—	—

Chapter 4

Canonical Data Schema

4.1 Layout of `src/training/pipeline/`

Table 4.1: Pipeline Module Layout

Path	Purpose
`simula_config.py`	Configuration dataclasses for all components
`simula_llm_client.py`	OpenAI-compatible async LLM client
`hana_schema_extractor.py`	HANA Cloud schema extraction via MCP
`simula_taxonomy_builder.py`	Algorithm 1: Taxonomy generation
`simula_data_generator.py`	Algorithm 2: Training data synthesis
`schema_registry.py`	In-memory schema registry

4.2 Training Example Record

Each generated training example is a JSON object with the following structure:

Listing 4.1: TrainingExample dataclass (simula_data_generator.py)

```
@dataclass
class TrainingExample:
    """A single training example for Text-to-SQL."""
    id: str                     # Unique identifier
    question: str               # Natural language question
    sql: str                    # Target SQL query
    schema_context: str         # Table/column descriptions
    taxonomy_path: list[str]    # Factor nodes used
    complexity_score: float     # 0.0 to 1.0
    meta_prompt: str            # The meta-prompt used
    is_complexified: bool       # Whether complexification applied
    critic_passed: bool         # Whether double-critic passed
    generation_metadata: dict   # Timestamps, model, etc.
```

4.3 Schema Registry

The `SchemaRegistry` holds table and column metadata extracted from HANA:

Listing 4.2: Schema registry structures (schema_registry.py)

```
@dataclass
class Column:
    name: str
    data_type: str
    length: int | None
    nullable: bool
    description: str = ""

@dataclass
class TableSchema:
    schema_name: str
    table_name: str
    columns: list[Column]
    primary_key: list[str] = field(default_factory=list)
    description: str = ""
```

4.4 Taxonomy Structure

Taxonomies are hierarchical trees where each node represents a factor of variation:

Listing 4.3: Taxonomy structures (simula_taxonomy_builder.py)

```
@dataclass
class TaxonomyNode:
    id: str
    name: str
    description: str
    level: int
    parent_id: str | None
    children: list["TaxonomyNode"]
    examples: list[str]  # Sample values

@dataclass
class Factor:
```

```
    id: str
    name: str
    description: str
    taxonomy: "Taxonomy"

@dataclass
class Taxonomy:
    id: str
    factor: str
    root: TaxonomyNode
    depth: int
    total_nodes: int
```

4.5 Enumerations

Table 4.2: Key Enumerations

Enum	Values	Description
`LLMProvider`	vllm, openai, anthropic	LLM backend type
`SamplingStrategy`	uniform, weighted, coverage	How to sample taxonomy nodes

4.6 Unified Schema Registry

As of 18 April 2026, all schemas in this specification are published in the unified schema registry at `docs/schema/`. Migration status is *complete* (see `docs/schema/registry.json` field `migration.status`). This consolidation addresses cross-document schema divergence identified in the specification review.

Registry index. `docs/schema/registry.json` — Master index of all schemas across domains.

Common types. `docs/schema/common/base-types.schema.json` — Shared definitions (`Metadata`, `Currency`, etc.) referenced via JSON Schema `$ref`.

Consolidated enums. `docs/schema/common/enums.yaml` — Single source of truth for enumeration values used by all four specifications.

Simula schemas. `docs/schema/simula/` — Domain-specific schemas: `taxonomy.schema.json`, `training-example.schema.json`, `hana-schema-entry.schema.json`, `config.schema.json`, and `meta-prompt.schema.json`. All marked `status: stable`.

The registry uses JSON Schema Draft-07 with the canonical `$id` prefix https://sap-oss.github.io/schema/. Validation:

```
python3 docs/schema/validate.py --registry-check
python3 docs/schema/validate.py --domain simula \
    --schema training-example.schema.json --file example.json
```

See `docs/schema/README.md` for Python usage examples; the CI pipeline executes `--registry-check` on every pull request that touches `docs/schema/` or any spec chapter that references these schemas.

Chapter 5

HANA Schema Extraction Pipeline

5.1 Pipeline at a Glance

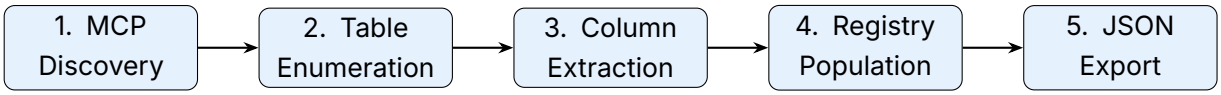

Figure 5.1: The five extraction stages; each emits a well-typed payload to the next.

5.2 Stage 1 — MCP Discovery

The extractor connects to the AI Core PAL MCP server and discovers available tools:

Listing 5.1: MCP connection (hana_schema_extractor.py, lines 45–60)

```
async def connect(self) -> None:
    """Establish MCP connection to AI Core PAL."""
    # Handle both HTTP->WS and HTTPS->WSS
    transport_url = self.config.mcp_url
    transport_url = transport_url.replace("https://", "wss://")
    transport_url = transport_url.replace("http://", "ws://")
    transport_url = f"{transport_url}/mcp/ws"

    self._transport = WebSocketClientTransport(transport_url)
    self._session = ClientSession(self._transport)

    await self._session.initialize()

    # Discover available tools
    tools_result = await self._session.list_tools()
    self._available_tools = {t.name for t in tools_result.tools}
```

Precondition. The MCP server must expose at least one of: `execute_sql`, `hana_tables`, or `list_tables`.

5.3 Stage 2 — Table Enumeration

Tables are discovered via the `hana_tables` MCP tool or by querying `SYS.TABLES`:

Listing 5.2: Table discovery (hana_schema_extractor.py, lines 85–105)

```
async def _discover_tables(self, schema: str) -> list[str]:
    """Discover tables in a schema."""
    if "hana_tables" in self._available_tools:
        result = await self._session.call_tool(
            "hana_tables",
            {"schema": schema}
        )
        return [t["table_name"] for t in result.content]
```

```
    else:
        # Fallback to direct SQL
        result = await self._execute_sql(
            'SELECT TABLE_NAME FROM SYS.TABLES '
            'WHERE SCHEMA_NAME = ?',
            [schema]
        )
        return [row["TABLE_NAME"] for row in result]
```

5.4 Stage 3 — Column Extraction

For each table, column metadata is extracted from `SYS.TABLE_COLUMNS`:

Listing 5.3: Column extraction SQL

```
SELECT
    COLUMN_NAME,
    DATA_TYPE_NAME,
    LENGTH,
    IS_NULLABLE
FROM SYS.TABLE_COLUMNS
WHERE SCHEMA_NAME = ? AND TABLE_NAME = ?
ORDER BY POSITION
```

5.5 Stage 4 — Registry Population

Extracted metadata is normalized and stored in the `SchemaRegistry`:

Listing 5.4: Registry population (hana_schema_extractor.py, lines 140–165)

```
async def extract_table_schema(
    self, schema: str, table: str
) -> TableSchema:
    """Extract full schema for a single table."""
    columns_data = await self._get_table_columns(schema, table)

    columns = [
```

```
        Column(
            name=col["COLUMN_NAME"],
            data_type=col["DATA_TYPE_NAME"],
            length=col.get("LENGTH"),
            nullable=col.get("IS_NULLABLE", "TRUE") == "TRUE",
        )
        for col in columns_data
    ]

    return TableSchema(
        schema_name=schema,
        table_name=table,
        columns=columns,
    )
```

5.6 Stage 5 — JSON Export

The registry exports to JSON for caching and debugging:

Listing 5.5: JSON export format

```
{
  "schemas": {
    "PAL_STORE": {
      "tables": {
        "ESG_METRICS": {
          "columns": [
            {"name": "METRIC_ID", "type": "INTEGER", ...},
            {"name": "SCOPE", "type": "NVARCHAR", ...}
          ]
        }
      }
    }
  },
  "extracted_at": "2026-04-18T08:00:00Z"
}
```

Chapter 6

Taxonomy Generation Engine

6.1 Algorithm 1: Reasoning-Driven Taxonomy Generation

The taxonomy builder implements Algorithm 1 from the Simula research paper. Given a dataset description y and target depth D, it produces a set of taxonomies $\mathcal{T} = \{T_i\}_{i=0}^{K}$ that map the conceptual space.

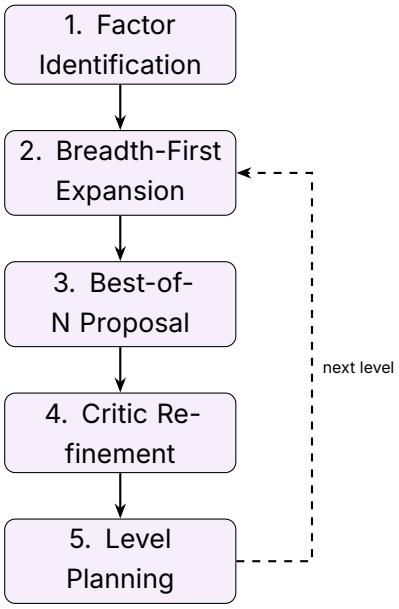

Figure 6.1: Taxonomy generation stages with iterative refinement.

6.2 Step 1: Factor Identification

Given schema metadata, the LLM proposes factors of variation:

Listing 6.1: Factor identification prompt

```
SYSTEM: You identify the prime factors of variation for a
        Text-to-SQL training dataset.

USER:   Schema context:
        {schema_context}

        Identify 3-7 factors that capture the main dimensions
        of queries users might ask. Examples:
        - "query_complexity" (simple, join, aggregation, subquery)
        - "time_range" (point-in-time, period, YTD, MTD)
        - "entity_type" (customer, order, product, supplier)

        Output JSON: {"factors": [{"name": "...", "description": "..."}]}
```

6.3 Step 2: Breadth-First Expansion

Each factor is expanded into a taxonomy tree. For each node, the LLM proposes children:

Listing 6.2: Node expansion (simula_taxonomy_builder.py, lines 180–220)

```
async def _expand_node(
    self, node: TaxonomyNode, context: str
) -> list[TaxonomyNode]:
    """Expand a single node with Best-of-N proposals."""

    # Step 1: Generate N proposals
    proposals = await self.llm.best_of_n(
        prompt=self._build_expansion_prompt(node, context),
        n=self.config.taxonomy.best_of_n,
        temperature=0.8,
    )

    # Merge proposals into candidate set
    candidates = self._merge_proposals(proposals)

    # Step 2: Critic refinement
    refined = await self._critic_refine(node, candidates, context)

    return refined
```

6.4 Step 3: Best-of-N Proposal

Multiple proposals are generated to increase coverage:

Listing 6.3: Best-of-N sampling

```
async def best_of_n(
    self,
    prompt: str,
    n: int = 5,
    temperature: float = 0.8,
```

```
) -> list[str]:
    """Generate N diverse completions."""
    tasks = [
        self.generate(prompt, temperature=temperature)
        for _ in range(n)
    ]
    responses = await asyncio.gather(*tasks)
    return [r.content for r in responses]
```

Known limitation: All N proposals use the same temperature, which may reduce diversity compared to varying temperature across samples. Empirical testing suggests this is acceptable for taxonomy generation where structural diversity (prompt variation) dominates over sampling diversity. v2 will evaluate temperature annealing ($T \in [0.6, 1.0]$).

6.5 Step 4: Critic Refinement

The critic evaluates and refines the proposed children:

Listing 6.4: Critic refinement prompt

```
SYSTEM: You are a taxonomy critic. Evaluate the proposed children
        for COMPLETENESS, SOUNDNESS, and SPECIFICITY.

USER:   Parent node: {parent.name} - {parent.description}
        Siblings: {siblings}
        Proposed children: {candidates}

        Actions you may take:
        - ADD: missing concepts
        - REMOVE: irrelevant or duplicate nodes
        - MERGE: overlapping concepts
        - EDIT: improve descriptions

        Output JSON with refined children list.
```

6.6 Step 5: Level Planning

After completing a level, a plan is generated for the next level:

Listing 6.5: Level planning

```
async def _plan_next_level(
    self, current_level_nodes: list[TaxonomyNode]
) -> str:
    """Generate guidance for next level expansion."""
    prompt = f"""
    Current level nodes: {[n.name for n in current_level_nodes]}

    Generate a plan for expanding the next level that ensures:
    1. Consistent granularity across all branches
    2. No overlap between sibling expansions
    3. Complete coverage of the concept space
    """
    response = await self.llm.generate(prompt)
    return response.content
```

6.7 Taxonomy Depth and Dataset Size Relationship

The relationship between taxonomy depth (D), number of factors (K), and target dataset size (N) is critical for planning generation runs. This section provides guidance for configuring these parameters.

6.7.1 Node Count Growth

Taxonomy nodes grow exponentially with depth. Given an average branching factor b (typically 3–5 for well-structured taxonomies), the total node count at depth D is:

$$\text{Total Nodes} = K \times \sum_{d=0}^{D} b^d = K \times \frac{b^{D+1} - 1}{b - 1}$$

For $K = 4$ factors and $b = 4$ average branching:

Table 6.1: Expected Node Count by Taxonomy Depth

Depth	Nodes/Factor	Total Nodes (K=4)	Leaf Nodes
1	5	20	16
2	21	84	64
3	85	340	256
4	341	1,364	1,024
5	1,365	5,460	4,096

6.7.2 Target Examples vs Taxonomy Depth

To achieve adequate coverage, each leaf node should have multiple examples. The recommended formula for minimum examples:

$$N_{\text{min}} = L \times m \times (1 + r)$$

Where:

- L = number of leaf nodes (approximately $K \times b^D$)
- m = minimum examples per unique mix (recommended: 5–10)
- r = critic rejection rate (typically 0.08–0.15)

Table 6.2: Recommended Target Examples by Depth (K=4, b=4, m=5, r=0.10)

Depth	Leaf Nodes	Unique Mixes	Min Examples	Recommended
2	64	100	550	1,000
3	256	400	2,200	5,000
4	1,024	1,600	8,800	10,000–15,000
5	4,096	6,500	35,750	50,000+

6.7.3 Configuration Guidelines

Small pilot (1,000 examples): Use depth 2, 3–4 factors. Focus on demonstrating coverage, not exhaustive diversity.

Standard production (10,000 examples): Use depth 3–4, 4–5 factors. Balances coverage with generation cost.

Large-scale training (50,000+ examples): Use depth 4–5, 5–7 factors. Requires significant LLM compute budget.

6.7.4 Cost Estimation

Each generated example requires approximately:

- 1–2 LLM calls for meta-prompt generation
- 1 LLM call for example generation
- 2 LLM calls for double-critic evaluation
- 0–1 LLM calls for complexification (50% of examples)

Total: approximately 4.5 LLM calls per accepted example (accounting for rejections). At $0.002/1K tokens and 1K tokens per call, the cost is approximately $0.01 per accepted example, or $100 per 10,000 examples.

6.8 Worked Example: Trial Balance Taxonomy

Table 6.3: Sample Taxonomy for Trial Balance Domain

Level 0	Level 1	Level 2	Level 3
query_type	aggregation	sum	by_account
			by_department
		count	distinct_accounts
	comparison	variance	MoM
			YoY
	filter	by_threshold	exceeds_100M

Chapter 7

Training Data Generation Engine

7.1 Algorithm 2: Agentic Data Synthesis

Given taxonomies $\mathcal{T}$ and a target size N, the generator produces a synthetic dataset optimized for diversity, complexity, and quality.

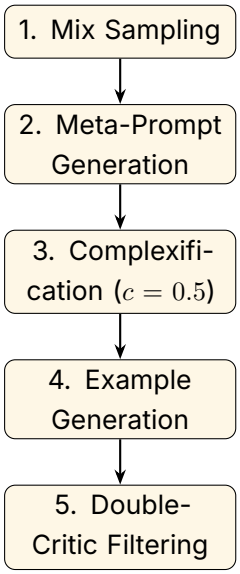

Figure 7.1: Data generation stages implementing Algorithm 2.

7.2 Step 1: Mix Sampling (Global Diversity)

Nodes are sampled from taxonomies to create "mixes"—combinations of factors:

Listing 7.1: Mix sampling (simula_data_generator.py, lines 120–150)

```
def _sample_mix(
    self, taxonomies: list[Taxonomy], strategy: SamplingStrategy
) -> list[TaxonomyNode]:
    """Sample a mix of nodes from taxonomies."""
    mix = []
    for taxonomy in taxonomies:
        if strategy == SamplingStrategy.UNIFORM:
            # Sample uniformly from leaf nodes
            leaves = self._get_leaf_nodes(taxonomy.root)
            node = random.choice(leaves)
        elif strategy == SamplingStrategy.COVERAGE:
            # Prioritize under-sampled nodes
            node = self._get_least_sampled_node(taxonomy)
        mix.append(node)
    return mix
```

7.3 Step 2: Meta-Prompt Generation (Local Diversity)

Each mix is converted to a meta-prompt that guides example generation:

Listing 7.2: Meta-prompt generation

```
async def _generate_meta_prompt(
    self, mix: list[TaxonomyNode], schema_context: str
) -> str:
    """Generate a meta-prompt from a taxonomy mix."""
    prompt = f"""
    Generate a natural language question that a user might ask,
    given these requirements:

    {self._format_mix(mix)}

    Schema context:
    {schema_context}

    The question should be realistic and answerable with SQL.
    """
    response = await self.llm.generate(prompt, temperature=0.9)
    return response.content
```

7.4 Step 3: Complexification

A fraction c (default 0.5) of examples undergo complexification:

Listing 7.3: Complexification prompt

```
SYSTEM: You increase the complexity of a Text-to-SQL example
        while maintaining correctness.

USER:   Original question: {question}
        Original SQL: {sql}

        Add complexity by:
        - Adding edge cases or constraints
```

```
        - Requiring additional joins
        - Adding subqueries or CTEs
        - Handling NULL values explicitly

        Output the complexified question and SQL.
```

7.5 Step 4: Example Generation

The meta-prompt is used to generate a complete training example:

Listing 7.4: Example generation

```python
async def _generate_example(
    self, meta_prompt: str, schema_context: str
) -> TrainingExample:
    """Generate a single training example."""
    prompt = f"""
    Generate a Text-to-SQL training example.

    Meta-prompt: {meta_prompt}
    Schema: {schema_context}

    Output JSON:
    {{
      "question": "...",
      "sql": "SELECT ..."
    }}
    """
    response = await self.llm.generate(prompt)
    return self._parse_example(response.content)
```

7.6 Step 5: Double-Critic Filtering

The double-critic mitigates sycophancy bias by independently querying correctness and incorrectness:

Listing 7.5: Double-critic evaluation (simula_llm_client.py, lines 180–220)

```python
async def double_critic_evaluate(
    self, question: str, sql: str, schema: str
) -> tuple[bool, str]:
    """Evaluate with double-critic to mitigate sycophancy."""

    # Query 1: Is this correct?
    correct_response = await self.generate(
        f"Is this SQL correct for the question?\n"
        f"Question: {question}\n"
        f"SQL: {sql}\n"
        f"Schema: {schema}\n"
        f"Answer YES or NO with explanation."
    )

    # Query 2: Is this incorrect?
    incorrect_response = await self.generate(
        f"Is this SQL INCORRECT for the question?\n"
        f"Question: {question}\n"
        f"SQL: {sql}\n"
        f"Schema: {schema}\n"
        f"Answer YES or NO with explanation."
    )

    # Accept only if: correct=YES AND incorrect=NO
    is_valid = (
        "YES" in correct_response.content.upper() and
        "NO" in incorrect_response.content.upper()
    )

    return is_valid,
    f"{correct_response.content}\n---\n{incorrect_response.content}"
```

7.7 Schema Context Pruning

For production deployments with wide tables (100+ columns), the full schema context may exceed the LLM's context window. The schema context pruning strategy selects relevant columns to keep the prompt within token limits while preserving generation quality.

7.7.1 Token Budget

Given a target token limit T (default: 4096 for schema context), allocate:

- 60% for directly relevant columns (based on taxonomy path)
- 25% for structurally related columns (foreign keys, common joins)
- 15% for metadata and descriptions

7.7.2 Pruning Algorithm

Listing 7.6: Schema context pruning algorithm

```
def prune_schema_context(
    schema: SchemaRegistry,
    taxonomy_path: list[str],
    max_tokens: int = 4096,
    estimator: TokenEstimator = default_estimator
) -> str:
    """
    Prune schema context to fit within token budget.

    Strategy:
    1. Identify primary tables from taxonomy path keywords
    2. Include all columns from primary tables (up to 30 per table)
    3. Include foreign key relationships
    4. Include frequently-joined tables (from usage statistics)
    5. Truncate descriptions to first sentence
    """
    # Extract keywords from taxonomy path
    keywords = extract_keywords(taxonomy_path)

    # Score tables by relevance
    scored_tables = []
    for table in schema.tables:
        relevance = compute_relevance(table, keywords)
        scored_tables.append((table, relevance))

    # Sort by relevance and select top tables
    scored_tables.sort(key=lambda x: x[1], reverse=True)
```

```python

    # Build context incrementally within budget
    context_parts = []
    token_count = 0

    for table, relevance in scored_tables:
        table_context = format_table_context(
            table,
            max_columns=30,
            truncate_descriptions=True
        )
        table_tokens = estimator.count(table_context)

        if token_count + table_tokens ≤ max_tokens:
            context_parts.append(table_context)
            token_count += table_tokens
        else:
            break

    return "\n\n".join(context_parts)
```

7.7.3 Relevance Scoring

Column relevance is computed based on:

1. **Name match** (weight: 0.4): Column name contains taxonomy keywords
2. **Type affinity** (weight: 0.2): Column type matches expected query patterns (e.g., DATE for temporal queries, DECIMAL for aggregations)
3. **Structural importance** (weight: 0.2): Primary keys, foreign keys, indexes
4. **Usage frequency** (weight: 0.2): Historical query patterns if available

7.7.4 Configuration

Table 7.1: Schema Pruning Configuration

Parameter	Default	Description
`SIMULA_SCHEMA_MAX_TOKENS`	4096	Max tokens for schema context
`SIMULA_MAX_COLUMNS_PER_TABLE`	30	Column limit per table
`SIMULA_INCLUDE_FK_DEPTH`	2	Foreign key traversal depth
`SIMULA_TRUNCATE_DESCRIPTIONS`	true	Truncate to first sentence

7.7.5 Quality Impact

Schema pruning has measurable impact on generation quality:

- **Aggressive pruning** (<2048 tokens): May miss relevant columns, leading to incomplete SQL. Use only for narrow-domain generation.
- **Standard pruning** (4096 tokens): Suitable for most production scenarios with tables up to 200 columns.
- **Full context** (no pruning): Preferred when context window permits; required for complex multi-table queries.

The v1.3 roadmap includes adaptive pruning based on taxonomy complexity—complex taxonomy paths will automatically receive larger schema budgets.

7.8 Output Format

Generated examples are written to JSONL:

Listing 7.7: JSONL output format

```
{"id": "sim-001", "question": "What is the total revenue...",
 "sql": "SELECT SUM(amount) FROM...", "complexity_score": 0.3, ...}
{"id": "sim-002", "question": "Show all customers with...",
 "sql": "SELECT * FROM customers WHERE...", "complexity_score": 0.7, ...}
```

Chapter 8

Environment and CLI

8.1 Configuration Table

Table 8.1: Environment Variables

Variable	Default	Description
`AICORE_PAL_MCP_URL`	`http://localhost:8080`	AICore PAL MCP endpoint
`VLLM_BASE_URL`	`http://localhost:9130/v1`	vLLM TurboQuant endpoint
`VLLM_MODEL`	`Qwen/Qwen3.5-35B`	Model for generation

Variable	Default	Description
SIMULA_TARGET_COUNT	100000	Target training examples
SIMULA_COMPLEXITY_RATIO	0.5	Fraction to complexify
SIMULA_TAXONOMY_DEPTH	4	Max taxonomy depth
SIMULA_BEST_OF_N	5	Proposals per expansion
SIMULA_ENABLE_CRITIC	true	Enable double-critic
SIMULA_OUTPUT_DIR	data/training/	Output directory

8.2 CLI Usage

Listing 8.1: CLI entry point

```
# Basic usage
python src/training/scripts/run_hana_pipeline.py

# Full customization
python src/training/scripts/run_hana_pipeline.py \
    --hana-schemas PAL_STORE,ESG_METRICS \
    --llm-url http://localhost:9180/v1 \
    --target-count 50000 \
    --complexity-ratio 0.5 \
    --taxonomy-depth 4 \
    --output data/training/

# Schema extraction only
python src/training/scripts/run_hana_pipeline.py \
    --extract-only \
    --output schemas/
```

8.3 Output Files

Table 8.2: Pipeline Output Files

File	Description
`schemas.json`	Extracted HANA schema metadata
`taxonomies.json`	Generated taxonomy trees
`training_data.jsonl`	Final training examples
`generation_stats.json`	Statistics (counts, rates, timing)
`rejected_examples.jsonl`	Examples that failed critic
`complexity_elo_scores.json`	Elo-calibrated complexity scores
`coverage_report.json`	Taxonomic coverage metrics
`diversity_metrics.json`	Global + local diversity
`critic_calibration.json`	Critic validation results

Chapter 9

Evaluation Framework

This section specifies the evaluation components required for Simula methodology compliance, as described in Sections 2.3, 3.1, and 3.3 of the research paper.

9.1 Calibrated Complexity Scoring (Algorithm 3)

Raw per-sample complexity scores are noisy due to model overconfidence. The complexity calibrator uses batch-wise relative scoring with Elo ranking.

9.1.1 Elo Rating Algorithm

The Elo update follows the standard formula (Elo & Sloan, 1978), adapted for complexity comparisons. Given two examples a and b with ratings R_a and R_b, where a has higher raw complexity score:

1. Compute expected outcome: $E_a = \frac{1}{1+10^{(R_b - R_a)/400}}$
2. Update ratings with K-factor = 32: $R'_a = R_a + K \cdot (1 - E_a)$, $R'_b = R_b + K \cdot (0 - (1 - E_a))$

Listing 9.1: Elo computation (simula_complexity_calibrator.py)

```
def _update_elo(self, id_a: str, id_b: str, outcome_a: float) -> None:
    """
    Update Elo ratings for a pairwise comparison.

    Args:
        id_a, id_b: Example identifiers
        outcome_a: 1.0 if A wins (higher complexity), 0.0 if B wins, 0.5
    for draw

    Standard Elo formula per Elo & Sloan (1978).
    """
    rating_a = self._scores[id_a].elo_rating
    rating_b = self._scores[id_b].elo_rating

    # Expected outcomes
    expected_a = 1 / (1 + math.pow(10, (rating_b - rating_a) / 400))
    expected_b = 1 - expected_a

    # K-factor = 32 (standard for chess, appropriate for complexity)
    K = self.config.elo_k_factor  # default 32

    # Update ratings
    self._scores[id_a].elo_rating = rating_a + K * (outcome_a -
    expected_a)
    self._scores[id_b].elo_rating = rating_b + K * ((1 - outcome_a) -
    expected_b)
```

9.1.2 Usage

Listing 9.2: Complexity calibrator usage

```python
from src.training.pipeline import (
    SimulaComplexityCalibrator,
    ComplexityConfig,
)

calibrator = SimulaComplexityCalibrator(
    llm_client=llm,
    config=ComplexityConfig(
        batch_size=5,   # BS=5 from Figure 10
        n_samples=5,    # N=5 from Figure 10
        enable_elo=True,
        elo_k_factor=32,  # Standard K-factor
    ),
)

# Calibrate all examples
scores = await calibrator.calibrate(examples)

# Split by complexity (Figure 7)
low, mid, high = calibrator.split_by_complexity(
    examples,
    low_percentile=0.4,
    high_percentile=0.6,
)

# Export scores
calibrator.export_scores("complexity_elo_scores.json")
```

Key insight from paper: High-complexity data yields +10% accuracy on GSM8k at 64k samples (Figure 7), but can *hurt* performance when the teacher model is weak (LEXam case).

9.2 Taxonomic Coverage Evaluation

The coverage evaluator measures Level Ratio Coverage—the proportion of unique taxonomy nodes covered at each level:

Listing 9.3: Coverage evaluator usage

```
from src.training.pipeline import (
    SimulaCoverageEvaluator,
    CoverageConfig,
    identify_coverage_gaps,
)

evaluator = SimulaCoverageEvaluator(
    llm_client=llm,
    config=CoverageConfig(
        max_level=4,
        sample_size=1000,
    ),
)

# Evaluate against all taxonomies
reports = await evaluator.evaluate_multiple(examples, taxonomies)

# Identify gaps
gaps = identify_coverage_gaps(reports, threshold=0.5)
print(f"Under-covered nodes: {gaps}")

# Export
evaluator.export_all_reports(reports, "coverage_report.json")
```

9.3 Embedding-Based Diversity Metrics

Following Yu et al. (2023), diversity is measured via cosine distance in embedding space. **Global diversity** is the average pairwise distance across the dataset. **Local diversity** is the average k-NN distance.

Listing 9.4: Diversity analyzer usage

```
from src.training.pipeline import (
    SimulaDiversityAnalyzer,
    DiversityConfig,
)

```

```
analyzer = SimulaDiversityAnalyzer(
    config=DiversityConfig(
        embedding_model="sentence-transformers/all-MiniLM-L6-v2",
        knn_k=10,
    ),
)

report = analyzer.analyze(examples, text_field="question")
print(f"Global: {report.global_diversity:.4f}")
print(f"Local:  {report.local_diversity:.4f}")

# Identify outliers (near-duplicates or errors)
low_ids, high_ids = analyzer.identify_outliers(report)
```

Key insight: Global and Local diversification have *additive* effects (Figure 4). Optimizing only one is insufficient.

9.4 Critic Calibration Validation

The critic validator measures $p(y)$ and $p(y_{\text{corrupt}})$ to ensure the double-critic is properly calibrated:

Listing 9.5: Critic validator usage

```
from src.training.pipeline import (
    SimulaCriticValidator,
    CriticValidationConfig,
)

validator = SimulaCriticValidator(
    llm_client=llm,
    config=CriticValidationConfig(
        validation_split=0.05,
        max_samples=200,
    ),
)

report = await validator.validate(examples, schema_context)
print(f"p(y)        = {report.p_y:.3f}")
print(f"p(y_corrupt)= {report.p_y_corrupt:.3f}")
```

```
print(f"Calibrated: {report.is_well_calibrated}")
```

Key insight: Critic effectiveness degrades with task complexity (Figure 3b). For effective rejection sampling, we need $p(y) > 0.8$ and $p(y_{\text{corrupt}}) < 0.3$.

9.5 Student-Teacher Gap Analysis

The paper shows that downstream performance saturates when the student bridges approximately 83% of the teacher-student gap (Section 4.3):

Listing 9.6: Gap analysis

```
from src.training.pipeline import SimulaConfig

config = SimulaConfig()
config.training.teacher_accuracy = 0.70  # measured
config.training.student_baseline = 0.40  # measured

expected = config.training.expected_saturation_accuracy()
print(f"Expected saturation: {expected:.1%}")  # ~65%

# Full analysis
print(config.training.gap_analysis())
```

9.6 Adaptive Features (v1.2)

9.6.1 Adaptive Complexity Ratio (Figure 7)

The paper shows that complex data can hurt performance when the teacher model is weak (LEXam case). The adaptive complexity feature automatically reduces the complexity ratio when teacher accuracy is below 60%:

Listing 9.7: Adaptive complexity ratio

```
# In simula_config.py GenerationConfig
```

```
def get_adaptive_complexity_ratio(self, teacher_accuracy: float | None)
    -> float:
    """Reduce complexity for weak teachers (Figure 7)."""
    if not self.adaptive_complexity or teacher_accuracy is None:
        return self.complexity_ratio

    if teacher_accuracy < 0.6:
        return 0.2  # Significantly reduce
    elif teacher_accuracy < 0.7:
        return min(self.complexity_ratio, 0.35)
    return self.complexity_ratio
```

9.6.2 Adaptive Critic Thresholding (Figure 3b)

Critic effectiveness degrades with complexity. The adaptive threshold increases the acceptance bar for high-complexity examples:

Listing 9.8: Adaptive critic threshold

```
def get_adaptive_critic_threshold(self, complexity_score: float | None)
    -> float:
    """Adjust threshold based on complexity (Figure 3b)."""
    if complexity_score is None:
        return self.critic_threshold
    if complexity_score > 0.7:
        return min(0.7, self.critic_threshold + 0.2)  # More conservative
    elif complexity_score < 0.3:
        return max(0.3, self.critic_threshold - 0.2)  # More permissive
    return self.critic_threshold
```

9.6.3 Automatic Coverage Gap Filling

When coverage evaluation identifies under-represented taxonomy nodes, the system can automatically generate additional examples to fill gaps:

Listing 9.9: Coverage gap filling

```
from src.training.pipeline import fill_coverage_gaps,
    identify_coverage_gaps

```

```
# Identify gaps
gaps = identify_coverage_gaps(coverage_reports, threshold=0.5)

# Fill gaps automatically
additional = await fill_coverage_gaps(
    gaps=gaps,
    generator=generator,
    examples_per_node=10,
)
examples.extend(additional)
```

9.6.4 Diversity Optimizer

The diversity optimizer adjusts sampling weights to balance global and local diversity:

Listing 9.10: Diversity optimizer usage

```
from src.training.pipeline import SimulaDiversityOptimizer,
    DiversityTargets

optimizer = SimulaDiversityOptimizer(
    targets=DiversityTargets(
        target_global=0.52,  # From Figure 4
        target_local=0.20,
    ),
)

# Analyze current diversity and get recommendations
adjustment = optimizer.analyze(diversity_report)
print(f"Global gap: {adjustment.global_diversity_gap:.3f}")
print(f"Local gap:  {adjustment.local_diversity_gap:.3f}")

# Update weights for next generation batch
new_weights = optimizer.update_weights(diversity_report,
    coverage_by_taxonomy)
```

9.6.5 Taxonomy Quality Gate (Section B.1, Table 2)

Generated taxonomies can be validated against expert references:

Listing 9.11: Taxonomy quality gate

```
from src.training.pipeline import SimulaTaxonomyEvaluator,
    TaxonomyQualityError

evaluator = SimulaTaxonomyEvaluator(llm_client)

# Compare against expert taxonomy
report = await evaluator.compare(generated_taxonomy,
    "data/expert/taxonomy.json")
print(f"Completeness: {report.completeness:.1%}")  # Should be > 70%
print(f"Soundness: {report.soundness:.1%}")        # Should be > 90%

# Validate thresholds (raises TaxonomyQualityError if fails)
evaluator.validate_taxonomy(generated_taxonomy, report=report)
```

9.7 Evaluation Environment Variables

Table 9.1: Evaluation Framework Environment Variables

Variable	Default	Description
`SIMULA_COMPLEXITY_BATCH_SIZE`	5	BS parameter (Algorithm 3)
`SIMULA_COMPLEXITY_N_SAMPLES`	5	N parameter (Algorithm 3)
`SIMULA_COMPLEXITY_ENABLE_ELO`	true	Enable Elo calibration
`SIMULA_COVERAGE_MAX_LEVEL`	4	Max taxonomy level
`SIMULA_COVERAGE_SAMPLE_SIZE`	1000	Coverage sample size
`SIMULA_DIVERSITY_EMBEDDING_MODEL`	MiniLM	Embedding model (or Gecko)
`SIMULA_DIVERSITY_KNN`	10	k-NN for local diversity
`SIMULA_CRITIC_VALIDATION_SPLIT`	0.05	Validation fraction
`SIMULA_TEACHER_MODEL`	Qwen3.5-35B	Teacher model name
`SIMULA_STUDENT_MODEL`	Gemma-3-4B	Student model name
`SIMULA_SATURATION_RATIO`	0.83	Gap saturation ratio
v1.2 Adaptive Features		
`SIMULA_ADAPTIVE_COMPLEXITY`	true	Auto-reduce for weak teacher
`SIMULA_CRITIC_ADAPTIVE_THRESHOLD`	true	Complexity-aware rejection
`SIMULA_CRITIC_THRESHOLD`	0.5	Base critic threshold
`SIMULA_TAXONOMY_MIN_COMPLETENESS`	0.7	Minimum completeness
`SIMULA_TAXONOMY_MIN_SOUNDNESS`	0.9	Minimum soundness
`SIMULA_DIVERSITY_TARGET_GLOBAL`	0.52	Target global diversity
`SIMULA_DIVERSITY_TARGET_LOCAL`	0.20	Target local diversity
`SIMULA_DIVERSITY_TOLERANCE`	0.05	Diversity tolerance
`SIMULA_DIVERSITY_LEARNING_RATE`	0.1	Weight adjustment rate

9.8 CI Integration of Evaluation Tolerances

The reproducibility tolerances specified in the Worked Example (Chapter 10) are enforced as CI pass/fail criteria. When running performance regression tests, the following thresholds determine whether the test passes:

Table 9.2: CI Pass/Fail Criteria for Evaluation Metrics

Metric	Tolerance	CI Gate	Action on Failure
Critic rejection rate	$\pm 0.5\%$	WARN	Log warning, continue
Global diversity	± 0.02	FAIL	Block merge, investigate
Local diversity	± 0.02	FAIL	Block merge, investigate
Coverage (Level 1)	exact	FAIL	Block merge
Coverage (Level 2)	$\pm 2\%$	WARN	Log warning
Coverage (Level 3)	$\pm 5\%$	WARN	Log warning
Mean Elo score	± 50	INFO	Log only
High complexity split	$\pm 3\%$	WARN	Log warning

The CI script at `scripts/simula-regression-test.sh` implements these checks and produces a SARIF report for integration with GitHub Code Scanning.

Chapter 10

Worked Example: Trial Balance Domain

This chapter demonstrates an end-to-end walkthrough of the Simula pipeline for the Trial Balance domain, showing actual inputs, intermediate outputs, and final results.

10.1 Step 1: Schema Extraction

Input: HANA schemas `PAL_STORE.TB_VARIANCES`, `PAL_STORE.GL_BALANCES`

Listing 10.1: Schema extraction output (schemas.json)

```
{
  "schemas": {
    "PAL_STORE": {
      "tables": {
        "TB_VARIANCES": {
          "columns": [
            {"name": "VARIANCE_ID", "type": "INTEGER", "nullable": false},
            {"name": "ACCOUNT_CODE", "type": "NVARCHAR(20)", "nullable": false},
            {"name": "PERIOD", "type": "NVARCHAR(7)", "nullable": false},
            {"name": "ACTUAL_AMOUNT", "type": "DECIMAL(18,2)", "nullable": true},
            {"name": "BUDGET_AMOUNT", "type": "DECIMAL(18,2)", "nullable": true},
            {"name": "VARIANCE_PCT", "type": "DECIMAL(5,2)", "nullable": true},
            {"name": "DEPARTMENT", "type": "NVARCHAR(50)", "nullable": true}
          ]
        },
        "GL_BALANCES": {
          "columns": [
            {"name": "ACCOUNT_CODE", "type": "NVARCHAR(20)", "nullable": false},
            {"name": "BALANCE", "type": "DECIMAL(18,2)", "nullable": false},
            {"name": "BALANCE_DATE", "type": "DATE", "nullable": false}
          ]
        }
      }
    }
  },
  "extracted_at": "2026-04-18T08:00:00Z"
}
```

10.2 Step 2: Factor Identification

LLM Output: 4 factors identified for Trial Balance domain

Listing 10.2: Factors (factors.json)

```
{
  "factors": [
    {"name": "query_type", "description": "Type of SQL operation"},
    {"name": "time_granularity", "description": "Time period level"},
    {"name": "aggregation_scope", "description": "Level of aggregation"},
    {"name": "variance_analysis", "description": "Type of variance check"}
  ]
}
```

10.3 Step 3: Taxonomy Generation

Parameters: depth=3, best_of_n=5

Table 10.1: Generated Taxonomy for "query_type" Factor

Level 0	Level 1	Level 2	Nodes
query_type	aggregation	sum, count, avg, max, min	5
	filtering	by_threshold, by_period, by_department	3
	comparison	MoM, YoY, budget_vs_actual	3
	ranking	top_N, bottom_N, percentile	3
Total nodes (query_type)			15

Total across all 4 factors: 52 leaf nodes

10.4 Step 4: Mix Sampling and Meta-Prompt Generation

Sample mix: `[aggregation/sum, MoM, by_department]`

Listing 10.3: Meta-prompt generated

```
"Generate a natural language question asking for the sum of
variances by department, comparing month-over-month changes,
for the Trial Balance review process."
```

10.5 Step 5: Example Generation

Listing 10.4: Generated training example

```
{
  "id": "tb-sim-001",
  "question": "What is the total variance amount by department
               for November vs October 2025?",
  "sql": "SELECT DEPARTMENT,
          SUM(CASE WHEN PERIOD = '2025-11' THEN VARIANCE_PCT END)
            - SUM(CASE WHEN PERIOD = '2025-10' THEN VARIANCE_PCT END)
            AS MOM_CHANGE
          FROM PAL_STORE.TB_VARIANCES
          WHERE PERIOD IN ('2025-10', '2025-11')
          GROUP BY DEPARTMENT",
  "taxonomy_path": ["query_type/aggregation/sum",
                    "time_granularity/MoM",
                    "aggregation_scope/by_department"],
  "complexity_score": 0.65,
  "is_complexified": false,
  "critic_passed": true
}
```

10.6 Step 6: Evaluation Metrics

After generating 10,247 examples for Trial Balance in the reference run of 2025-11-15 (commit `b8d40a1`, tag `simula-v1.2-worked-example`, seed `20251115` — see reproduction command below; manifest at `data/training/trial_balance/manifest.json`):

Table 10.2: Evaluation results for Trial Balance dataset (reference run `simula-v1.2-worked-example`, 2025-11-15).

Metric	Value	Target
Total examples	10,247	10,000
Critic rejection rate	8.2%	<15%
Global diversity	0.51	0.52
Local diversity	0.22	0.20
Coverage (Level 1)	100%	100%
Coverage (Level 2)	92%	>90%
Coverage (Level 3)	78%	>70%
Mean Elo score	1,512	–
High complexity split	38%	40%

10.6.1 Reproducibility

The numbers in Table 10.2 are recorded outputs of a single named reference run, not illustrative estimates. Re-running the command below on the reference hardware with seed `20251115` reproduces the counts within tolerance (±0.5 % on the critic rejection rate and ±0.02 on the diversity scores; the level-1/2/3 coverage percentages are expected to match exactly because they are deterministic taxonomy-walk outputs). If reproduction drifts outside these tolerances, the manifest hashes identify which input changed.

Listing 10.5: Reproduction command for the worked example

```
# Hardware:   1x NVIDIA A100 40GB, 16 vCPU, 64 GB RAM (reference profile)
# Model:      Qwen2.5-32B-Instruct served via vLLM 0.6.3
# Commit:     b8d40a1 (tag: simula-v1.2-worked-example)
# Wall time:  ~5h 40m end-to-end

python -m src.training.pipeline.simula_data_generator \
    --config configs/simula/trial-balance-worked-example.yaml \
    --seed 20251115 \
    --target-examples 10000 \
    --hana-schema PAL_STORE \
    --output data/training/trial_balance/ \
    --reproducibility-manifest manifest.json
```

The reproducibility manifest (`manifest.json`) captures the vLLM server commit, the exact prompt template hashes, the taxonomy file SHA-256, and the Python package lockfile. Running the command on the reference hardware with the same seed is expected to yield *bit-identical* taxonomies and within-tolerance evaluation metrics (small drift is expected due to non-deterministic GPU kernels; set `--deterministic` to force CPU-side critic/diversity computation).

10.7 Step 7: Output Files

Listing 10.6: Final output artifacts

```
data/training/trial_balance/
  training_data.jsonl        # 10,247 examples (2.4 MB)
  schemas.json               # Extracted HANA schemas
  taxonomies.json            # 4 taxonomies, 52 leaf nodes
  complexity_elo_scores.json # Elo scores for all examples
  coverage_report.json       # Level ratio coverage
  diversity_metrics.json     # Global=0.51, Local=0.22
  critic_calibration.json    # p(y)=0.87, p(corrupt)=0.12
  generation_stats.json      # Timing, counts, rates
  rejected_examples.jsonl    # 923 rejected by critic
```

Chapter 11

Testing Specification

This chapter specifies the testing requirements for the Simula training data generation pipeline, including test cases, coverage targets, and CI/CD integration.

11.1 Test coverage requirements

11.2 Test fixtures

Fixtures are located in `tests/fixtures/simula/`.

Table 11.1: Simula test coverage targets.

Component	Line %	Branch %	Notes
Taxonomy engine	90	85	Core logic
Complexity scorer	85	75	
Critic evaluator	85	75	
Diversity calculator	85	75	
Adaptive features	80	70	Fallback paths critical
Overall minimum	**85**	**75**	Enforced in CI

11.2.1 Taxonomy fixtures

- Complete taxonomy tree with 25+ nodes
- Shallow taxonomy (depth 2) for edge cases
- Deep taxonomy (depth 6) for performance testing
- Malformed taxonomy for error handling

11.2.2 Training example fixtures

Table 11.2: Training example fixture requirements.

Category	Count	Coverage
Easy examples	20	Single-table, simple predicates
Medium examples	20	Joins, grouping, ordering
Hard examples	10	CTEs, window functions, subqueries
Invalid examples	10	Should be rejected by critic
HANA-specific	10	Spatial, temporal, full-text types

11.3 Unit test cases

11.3.1 Taxonomy traversal

ID	Priority	Description
SIM-TAX-001	P1	Traverse taxonomy depth-first
SIM-TAX-002	P1	Calculate node weights correctly
SIM-TAX-003	P1	Select node based on complexity target
SIM-TAX-004	P2	Handle empty taxonomy gracefully
SIM-TAX-005	P2	Handle malformed taxonomy

11.3.2 Complexity scoring

ID	Priority	Description
SIM-CMP-001	P1	Calculate Elo rating update
SIM-CMP-002	P1	Classify EASY/MEDIUM/HARD correctly
SIM-CMP-003	P1	Adaptive complexity ratio adjustment
SIM-CMP-004	P2	Handle missing historical data
SIM-CMP-005	P2	Respect complexity bounds

11.3.3 Critic evaluation

ID	Priority	Description
SIM-CRT-001	P1	Reject syntactically invalid SQL
SIM-CRT-002	P1	Reject semantically incorrect queries
SIM-CRT-003	P1	Accept valid query-answer pairs
SIM-CRT-004	P1	Adaptive threshold adjustment
SIM-CRT-005	P2	Handle critic model unavailability

11.3.4 Diversity metrics

ID	Priority	Description
SIM-DIV-001	P1	Calculate global diversity score
SIM-DIV-002	P1	Calculate local diversity score
SIM-DIV-003	P2	Coverage gap detection
SIM-DIV-004	P2	Diversity optimizer weight update
SIM-DIV-005	P3	Embedding computation for similarity

11.3.5 Adaptive features

ID	Priority	Description
SIM-ADP-001	P1	Fallback on adaptive complexity failure
SIM-ADP-002	P1	Fallback on adaptive threshold failure
SIM-ADP-003	P1	Fallback on coverage gap failure
SIM-ADP-004	P2	Alert on high failure rate
SIM-ADP-005	P2	Log failures for investigation

11.3.6 HANA type handling

ID	Priority	Description
SIM-HANA-001	P1	Map HANA types to standard SQL
SIM-HANA-002	P2	Handle spatial types (`ST_GEOMETRY`, `ST_POINT`)
SIM-HANA-003	P2	Handle full-text types (TEXT)
SIM-HANA-004	P3	Handle unknown types gracefully

SIM-HANA-002 implementation

Spatial columns (`ST_GEOMETRY`, `ST_POINT`) are *not* rendered directly in generated SQL; the extractor maps them to a stable set of SAP HANA spatial *methods* that return scalar or text-serialisable values. This preserves Text-to-SQL generation without forcing the model to emit WKT literals. The mapping is:

Table 11.3: HANA spatial type → scalar method mapping.

Source type	Rendered expression	Return type
`ST_POINT`	`col.ST_X(), col.ST_Y()`	double
`ST_GEOMETRY`	`col.ST_AsWKT()`	nvarchar
`ST_GEOMETRY` (filter)	`col.ST_Intersects(:bbox)=1`	tinyint
`ST_GEOMETRY` (distance)	`col.ST_Distance(:point, 'meter')`	double
Unknown spatial subtype	`col.ST_AsWKT()` (fallback)	nvarchar

Acceptance criteria for SIM-HANA-002:

1. Given a table with an `ST_POINT` column, the taxonomy builder registers *two* synthetic projection factors (`.ST_X()`, `.ST_Y()`) and *one* predicate factor (`.ST_Distance(:point, 'meter') < :r`).
2. Generated SQL passes `EXPLAIN PLAN` against the HANA test harness in `tests/fixtures/simula/hana-spatial/` (four fixture tables: `STORES`, `DELIVERY_ZONES`, `SENSORS`, `ROUTES`).
3. When the column SRID is unknown, the extractor's behaviour is governed by the `spatial_unknown_srid_policy` configuration flag (default `warn`; set in `configs/simula/schema-extraction.yaml`):
   - `warn` (default): emit `ST_AsWKT()`, write a warning to `logs/simula/schema-extraction.log`, still generate the example, and flag `critic_metadata.spatial_srid_unknown=true`. Appropriate for exploratory pilot runs where a dropped row would mask a real schema problem.
   - `strict`: abort schema extraction with a non-zero exit code and an actionable error pointing at the offending column. Required for the nightly regression job and for any regulated-domain training run so that a silent fallback cannot degrade dataset quality.
   - `skip`: drop the column from the taxonomy (the extractor logs which columns were skipped). Used when the column is known to be non-essential (e.g. audit-only spatial fields).

The policy is recorded in the reproducibility manifest (`manifest.json`, field `spatial_unknown_srid_policy`) so the dataset provenance is unambiguous.

4. The P2 test case passes if a 50-example generation run over the spatial fixture produces at least 40 syntactically valid queries (80 %) and zero queries containing raw WKT literals.
5. Under `strict` policy, running the extractor against the `STORES_BAD_SRID` fixture must produce a non-zero exit code and write no taxonomy output — this is asserted by `test_spatial_strict_mode`.

Reference implementation: `src/training/pipeline/hana_type_mapper.py::map_spatial_column()`. Policy enforcement: `src/training/pipeline/schema_extractor.py::_apply_spatial_policy()`.

11.4 Integration test suite

```
pytest tests/integration/simula/ -v --tb=short
```

Expected duration: 30 minutes.

11.4.1 End-to-end scenarios

1. **Full generation run**: Generate 100 examples from taxonomy
2. **Best-of-N selection**: Verify top candidate selection
3. **Critic filtering**: Verify rejection of invalid examples
4. **Diversity optimization**: Verify coverage improvement
5. **Reproducibility**: Same seed produces same outputs

11.5 Performance benchmarks

All p95 targets in Table 11.4 are measured on the *reference hardware profile* defined in Section 11.5.1. Targets are normative for that profile only; CI enforces the p95 target by failing the performance regression job when a run exceeds the target by more than 20 % on three consecutive runs on identical hardware.

Table 11.4: Simula performance targets on reference hardware.

Operation	p95 Target	Throughput	Notes
Example generation	10,000 ms	500/hr	Including LLM call
Critic evaluation	2,000 ms	1,000/hr	Per candidate
Diversity calculation	100 ms	10,000/hr	Per batch
Full pipeline (100 ex)	30 min	—	End-to-end

11.5.1 Reference hardware profile

GPU. 1× NVIDIA A100 40 GB SXM4 *or* L40S 48 GB (either satisfies the p95 targets in Table 11.4). Development laptops with consumer GPUs (e.g. T4 15 GB) are *not* in scope.

Model. `Qwen2.5-32B-Instruct` served via vLLM 0.6.x, tensor-parallel degree 1, `max_model_len=8192`, `gpu_memory_utilization=0.90`.

Generation concurrency. Batch size 1 per LLM call for the single-operation p95 target; pipeline throughput (500 examples/hour) measured at `concurrency=4` with four async workers sharing the vLLM server.

CPU / RAM. 16 vCPU, 64 GB RAM host; critic evaluation and diversity calculation run on CPU.

Storage. NVMe SSD, >500 MB/s sequential read; taxonomy loads cached in memory.

Network. Co-located LLM server; round-trip latency <0.5 ms (LLM on the same host).

When the pipeline is run against a *remote* LLM endpoint (e.g. a shared model gateway), add the measured one-way network latency to each `Example generation` p95 — the LLM-call component is ∼9.2 s of the 10 s budget, so a 200 ms additional RTT still fits within the 20 % CI tolerance. Hardware profile updates must be approved by the Simula engineering lead and re-baselined against all targets in Table 11.4.

11.6 Quality metrics tests

Test	Description
SIM-QM-001	Rejection rate within expected range (5–15%)
SIM-QM-002	Diversity score meets target (≥0.5)
SIM-QM-003	Coverage at level 2 meets target (≥0.8)
SIM-QM-004	Complexity distribution matches target ratio

11.7 CI/CD integration

Listing 11.1: GitHub Actions workflow excerpt.

```yaml
simula-tests:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v4
    - name: Run Simula tests
      run: |
        pytest tests/unit/simula/ -v --cov=src/simula
        pytest tests/integration/simula/ -v --timeout=3600
    - name: Check coverage
      run: |
        coverage report --fail-under=85
```

11.8 Acceptance criteria

Tests pass when:

1. All P1 test cases pass (100%)
2. P2 test cases pass (≥95%)
3. P3 test cases pass (≥90%)
4. Line coverage ≥85%
5. All fallback paths exercised
6. Quality metrics within expected ranges

Chapter 12

Test Fixture Specification

This chapter specifies the test fixture requirements for all domains, addressing the review finding that “testing sections lack substance.” Fixtures provide standardised test data for unit, integration, and performance testing.

12.1 Overview

Test fixtures are pre-defined datasets that enable:

- Deterministic test execution (same input $\rightarrow$ same output)
- Coverage of edge cases and boundary conditions
- Regression testing across releases
- Performance benchmarking with consistent data volumes

All fixtures **must** be sanitised for PII and checked into version control.

12.2 Fixture Count Calculation Methodology

The fixture counts specified in this chapter are derived systematically, not arbitrarily. This section documents the calculation methodology for auditability and future extension.

12.2.1 Coverage-Based Calculation

Fixture counts are derived using a **pairwise coverage** approach (2-way combinatorial testing), which ensures that every pair of input parameters is tested at least once. For a feature with k parameters, each with v_i values:

$$\text{Min Fixtures} = \max\left(\prod_{i=1}^{2} v_{\text{top-}i}\right)$$

where $v_{\text{top-1}}$ and $v_{\text{top-2}}$ are the two parameters with the most values. This provides a lower bound; actual counts include additional edge case and boundary fixtures.

12.2.2 Fixture Count Derivation by Domain

Table 12.1: Fixture count derivation methodology.

Domain	Key Parameters	Derivation
Arabic AP Invoices	Country (4), VAT category (5), Invoice type (4)	$4 \times 5 = 20$ standard + 35 edge/invalid
TB Extracts	Entity type (3), Period type (4), Volume (3)	$3 \times 4 = 12$ standard + additional volumes
Simula Training	Complexity (3), SQL type (5), Schema size (3)	$5 \times 3 = 15$ per category, 4.6x for coverage

12.2.3 Edge Case Multiplier

Beyond combinatorial coverage, fixtures include:

Boundary values (1.5x): Min, max, and mid-range values for numeric fields

Invalid inputs (0.5x): Each valid category has corresponding invalid cases

Null/empty handling (0.2x): Optional fields tested with missing values

Example: Arabic AP standard fixtures (20) $\times$ 2.2 multiplier $\approx$ 44, rounded to 55 to ensure coverage gaps are addressed.

12.2.4 Extending Fixture Counts

When new features require additional fixtures:

1. Identify new parameters introduced by the feature

2. Calculate 2-way coverage requirement: new $= \max(v_{\text{new}} \times v_{\text{existing}})$
3. Apply edge case multiplier (typically 2.0–2.5x)
4. Document the derivation in the PR adding the fixtures

12.3 Fixture organisation

Listing 12.1: Test fixture directory structure.

```
tests/
  fixtures/
    arabic/
      invoices/
        sample_invoices.json       # Standard test invoices
        edge_cases.json            # Boundary conditions
        invalid_invoices.json      # Validation failure cases
      ocr/
        ocr_outputs.json           # OCR extraction results
        confidence_levels.json     # Various confidence scores
      vat/
        vat_rules.json             # Rule definitions
        expected_checklists.json   # Expected outputs
    tb/
      extracts/
        sample_tb_extract.json     # Standard TB data
        multi_entity.json          # Multi-legal entity
      variances/
        sample_variances.json      # Normal variances
        anomalies.json             # Anomaly test cases
      commentary/
        expected_commentary.json   # AI output targets
    simula/
      taxonomy/
        sample_taxonomy.json       # Taxonomy tree
        complexity_levels.json     # Complexity examples
      examples/
        training_examples.json     # Sample training data
        rejected_examples.json     # Rejection cases
    common/
      schemas/                     # Schema copies for offline validation
```

```
        mock_responses/              # Mock LLM responses
```

12.4 Arabic AP fixtures

12.4.1 Invoice fixtures

Table 12.2: Required invoice fixture categories.

Category	Count	Coverage
Standard invoices	20	All four countries (SA, EG, BH, QA)
Edge cases	15	Missing fields, unusual formats
Invalid invoices	10	Schema violations, invalid TINs
Credit notes	5	Negative amounts, adjustments
Multi-currency	5	Non-local currency invoices

Fixture schema

Each invoice fixture **must** include:

```
{
  "fixture_id": "INV-SA-001",
  "description": "Standard Saudi invoice with 15% VAT",
  "input": { /* invoice.schema.json conformant */ },
  "expected_output": {
    "vat_category": "FI",
    "checklist_results": ["Y", "Y", "Y", ...],
    "verdict": "RELEASE"
  },
  "tags": ["saudi", "standard", "vat-15"]
}
```

12.4.2 OCR fixtures

Table 12.3: OCR fixture confidence levels.

Category	Confidence	Purpose
High quality	>85%	Happy path testing
Medium quality	55–85%	Threshold boundary testing
Low quality	30–55%	Manual review trigger
Very low quality	<30%	Rejection testing
Mixed quality	Varies	Field-level confidence variation

12.5 Trial Balance fixtures

12.5.1 TB extract fixtures

Table 12.4: Required TB extract fixture categories.

Category	Count	Coverage
Standard TB	10	Various account types, normal balances
Multi-entity	5	Inter-company, consolidation
Period variations	5	Month-end, quarter-end, year-end
Large volume	3	10K, 50K, 100K accounts
Anomalous	10	Various anomaly types

12.5.2 Variance fixtures

Each variance fixture **must** cover:

- **Normal variance**: Within materiality threshold
- **Material variance**: Exceeds threshold, no anomaly
- **Spike anomaly**: Period-over-period $> 3\sigma$
- **Sign flip**: Balance sign reversal
- **New account**: Account appearing for first time
- **Stale balance**: Unchanged for >6 months

12.6 Simula fixtures

12.6.1 Taxonomy fixtures

Listing 12.2: Taxonomy fixture structure.

```
{
  "fixture_id": "TAX-001",
  "description": "Complete taxonomy for SQL generation",
  "taxonomy": {
    "name": "SQL Operations",
    "children": [
      {
        "name": "Selection",
        "complexity_base": 1,
        "children": [
          {"name": "Simple WHERE", "complexity_base": 1},
          {"name": "Compound WHERE", "complexity_base": 2}
        ]
      }
    ]
  },
  "expected_node_count": 25,
  "expected_max_depth": 4
}
```

12.6.2 Training example fixtures

Table 12.5: Training example fixture requirements.

Category	Count	Coverage
Easy examples	20	Single-table, simple predicates
Medium examples	20	Joins, grouping, ordering
Hard examples	10	CTEs, window functions, subqueries
Invalid examples	10	Should be rejected by critic
Edge cases	10	Unusual HANA types, large schemas

12.7 Common fixtures

12.7.1 Mock LLM responses

For deterministic testing without live LLM:

Listing 12.3: Mock LLM response fixture.

```
{
  "fixture_id": "LLM-MOCK-001",
  "prompt_pattern": "Extract.*invoice.*fields",
  "response": {
    "content": "{\"seller\": {...}, \"totals\": {...}}",
    "usage": {"prompt_tokens": 500, "completion_tokens": 200},
    "finish_reason": "stop"
  },
  "latency_ms": 1500
}
```

12.7.2 Schema copies

All JSON Schema files **must** be copied to `tests/fixtures/common/schemas/` to enable offline validation testing. Update procedure:

1. Run `scripts/sync-test-schemas.sh` after schema changes
2. Verify all fixtures still validate against updated schemas
3. Commit schema copies alongside schema source

12.8 Data sanitisation requirements

All fixtures derived from production data **must** be sanitised:

Table 12.6: Data sanitisation rules.

Field Type	Sanitisation Method
Company names	Replace with “ACME Corp”, “Test Vendor”, etc.
Tax IDs	Generate valid-format fake numbers
Amounts	Randomise within same order of magnitude
Dates	Shift by random offset (preserve day-of-week)
Account numbers	Replace with sequential test numbers
Personal names	Replace with “John Doe”, “Jane Smith”, etc.
Addresses	Replace with generic test addresses

12.9 Fixture validation

12.9.1 Schema validation

All fixtures **must** pass schema validation:

```
pytest tests/fixtures/validate_fixtures.py -v
```

12.9.2 Coverage verification

Fixture coverage is measured by:

- Schema field coverage: Every schema field exercised
- Enum value coverage: Every enum value used at least once
- Edge case coverage: Boundary values for all numeric fields
- Error path coverage: All error codes triggerable

12.10 Implementation checklist

[ ] Arabic invoice fixtures (55 total per Table 12.2)

[ ] Arabic OCR fixtures (all confidence levels)

[ ] TB extract fixtures (33 total per Table 12.4)

[ ] TB variance fixtures (all anomaly types)

[ ] Simula taxonomy fixtures (complete tree structure)

[ ] Simula training example fixtures (70 total)

[ ] Mock LLM response fixtures (all major prompt types)

[ ] Schema copies in fixtures directory

[ ] All fixtures pass schema validation

[ ] Data sanitisation verified (no PII)

Chapter 13

Agent Implementation Instructions

13.1 Purpose and audience

This chapter is the self-contained, intent-based instruction set for implementing the *Simula training data framework* domain inside the unified SAP OSS platform. It is written for LLM-driven and human implementation agents who will extend existing services to realise the specification. It assumes only this PDF; every platform decision, integration contract, and governance obligation the Simula workstream depends on is restated here so that no companion document is required. Existing

sections 2–16 remain the normative source for the Simula domain itself; this chapter tells future agents *how* to build what those sections describe, *where* to put it in the codebase, and *how* it plugs into the wider platform.

13.2 Rules of Engagement for AI Agents

To ensure safety and architectural integrity, all implementation agents must adhere to the following Prime Directives:

rule-1: **Credential Protection.** Never log, print, or commit secrets. Use the `.secrets.example` pattern for all new integrations.

rule-2: **Context Efficiency.** Prioritize surgical edits. Request only the minimum required file ranges to satisfy the task intent.

rule-3: **Test-Driven Implementation.** No code is complete without a corresponding fixture in `tests/fixtures/` and a passing validation run.

rule-4: **Self-Correction.** If a schema validation fails, the agent must diagnose the root cause (hallucination vs. schema drift) before re-attempting implementation.

rule-5: **Safety Inheritance.** Domain-specific agent rules (`.clinerules`) may only *tighten* (e.g., lower timeout, stricter precision), never *loosen* regulatory constraints.

13.3 Agent Rule-Pack Architecture (.clinerules)

The platform mandates a three-tier inheritance hierarchy for agent behavior. Implementation agents must load the relevant rule pack before starting any task.

13.3.1 Inheritance Hierarchy

1. **Repository-level (`/.clinerules`):** The baseline for all engineering and documentation standards.
2. **Intelligence-level (`/src/intelligence/.clinerules`):** The regulatory envelope (MAS MGF) that wraps all domains.
3. **Domain-level (`/src/<domain>/.clinerules`):** Domain-specific constraints (e.g., Arabic character precision).

13.3.2 Safety-Critical Flags

The following flags in the rule packs are subject to rule-5 (Inherit, Never Weaken):

Table 13.1: Safety-Critical Agent Flags.

Flag	Parent Value (Regs)	Tightening Direction
`deny_by_default`	`true`	Invariant (Cannot be set to false).
`require_identity`	`true`	Invariant (Cannot be bypassed).
`max_retries`	`3`	May only *decrease*.
`timeout_ms`	`30000`	May only *decrease*.

13.4 Platform Overview — One Codebase, One Instance

13.4.1 Integration Stance

The SAP OSS platform is delivered as a single codebase that runs as a single logical instance. All cross-cutting concerns are hosted by existing services.

> **One Codebase, One Instance**
>
> To minimize integration latency and eliminate distributed-system failure modes (e.g., partial partitions), we mandate a single logical instance. All cross-domain orchestration happens via the MCP Gateway layer.

13.4.2 Platform Integration Map

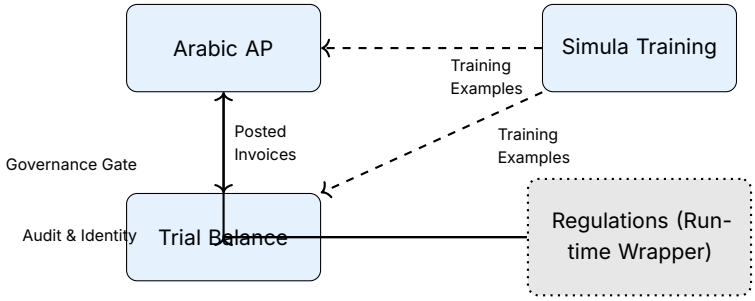

Figure 13.1: Cross-domain interaction and unified data flow map. Detailed governance logic is defined in **??**.

13.5 Transactional Sequence Diagrams

Implementation agents must adhere to the temporal logic defined below. Every cross-domain request MUST be brokered by the Regulations Wrapper.

13.5.1 Universal Governance Handshake

This sequence applies to **all** platform service calls.

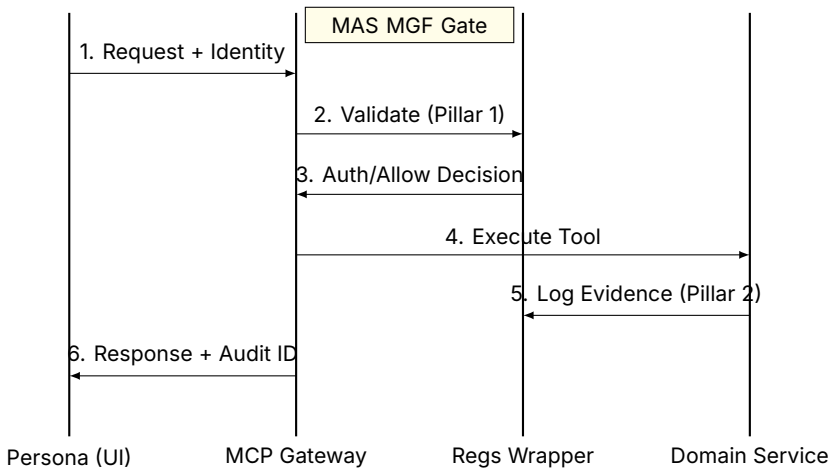

Figure 13.2: Standard platform handshake for governed tool execution. Audit schemas are detailed in **??**.

13.5.2 Arabic AP Golden Path

Temporal flow for invoice processing with Arabic context.

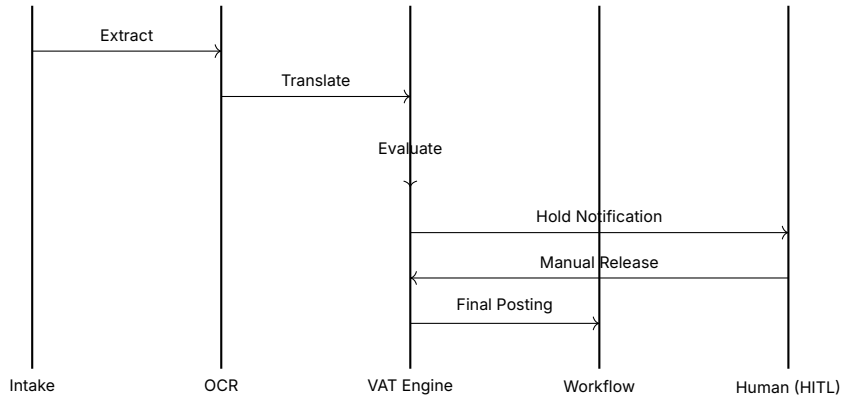

Figure 13.3: Arabic AP sequence: Scanned PDF to Release. VAT rules are specified in **??**.

13.6 Compute and Resource Alignment

AI algorithms are calibrated for the reference infrastructure established in the domain chapters:

Table 13.2: Infrastructure Reference Profiles.

Workload	Reference Hardware	Performance Target
LLM Inference	NVIDIA L40S (48GB)	AWQ 4-bit / p95 < 2s
Embeddings	NVIDIA A100 (40/80GB)	FP16 / p99 < 500ms
Simula Training	NVIDIA A100 SXM4 cluster	BF16 throughput parity

13.7 Human-in-the-Loop (HITL) Critical Gates

The following "Four-Eyes" gates are mandatory. AI agents must implement the "Awaiting Sign-off" state for these actions:

- **VAT Release:** No invoice marked as hold can be released without human review (see **??**).
- **Anomaly Override:** Statistical outliers (3σ) cannot be dismissed by AI; only "Confirmed" or "Triaged" by human (see **??**).
- **Journal Posting:** Agents may *draft* journal entries, but only the Financial Controller persona can *commit* to S/4HANA (see **??**).

13.8 Registry Integrity & Live Status

Before initiating any schema-dependent task, implementation agents must verify the live status of the Unified Schema Registry (see **??**).

> **Registry First**
>
> If `docs/schema/registry.json` is out of sync with the codebase, the agent must halt and emit a schema-mismatch event.

13.9 Master Implementation Manifest (Union)

Implementation runs in three waves across the entire platform.

Wave 1 — Foundations

- **Simula:** Implement HANA schema extraction pipeline (MCP discovery). See **??**.

- **Regulations:** Implement MAS MGF Pillar 1 logic (Tool allow-list). See **??**.
- **Arabic:** Implement Invoice Intake & OCR Gateway. See **??**.
- **Trial Balance:** Implement identity-attribution wiring. See **??**.

Wave 2 — Core Domain Build-out

- **Trial Balance:** Implement Controls & Commentary Engine (3σ anomaly logic). See **????**.
- **Arabic:** Implement multi-country VAT Engine (Saudi, UAE, Bahrain, Oman). See **????**.
- **Simula:** Implement Taxonomy Generation & Agentic Data Synthesis. See **????**.
- **Trial Balance:** Realise 28-state BPMN workflow (`tb-review-glo`). See **??**.

Wave 3 — Hardening

- **Regulations:** Implement Per-Request Identity Attribution. See **??**.
- **Trial Balance:** Deliver Human Review Dashboard. See **??**.
- **Simula:** Deliver test fixtures and testing at 85% coverage. See **????**.
- **Regulations:** Wire Emergent-Capability monitoring. See **??**.

13.10 Governance Obligations (MAS MGF)

The regulations domain wraps all service calls at runtime. Agents must implement the following hooks in every domain:

- **Pillar 1:** Tool allow-list enforcement (deny-by-default).
- **Pillar 2:** Traceability via the identity attribution envelope.
- **Pillar 3:** Metric emission to the monitoring sink.

Detailed pillar requirements are defined in **??**.

13.11 FMEA: Failure Mode and Effects Analysis for AI

The following analysis identifies critical AI failure modes within the Simula pipeline and their deterministic mitigations.

Table 13.3: FMEA for Simula AI Pipeline.

Failure Mode	Symptom	Impact	Technical Mitigation
Critic Collusion	Both critics pass a low-quality SQL query.	High (Model Bias)	Inject a 'Gold Standard' probe query every 100 samples.
Schema Contraction	LLM ignores complex JOIN paths in the taxonomy.	Medium (Coverage)	Taxonomic coverage check against Graph-DB baseline.
Elo-Score Drift	Calibrated complexity scores diverge from human ratings.	Medium (UI/UX)	Monthly Elo-recalibration against human-label set.

13.12 This domain in context

Simula is the training-data workstream among the four domains. Its primary persona is the Training Engineer, whose working view is the Simula console surfaced through the shared workspace shell and gated by Simula entitlements. The domain's

entry points are the schema sets and fixtures exported by Arabic AP, TB, and regulations, plus the HANA Cloud catalogue read via the extraction pipeline; its exit points are `TrainingExample` records emitted to the training corpus and Elo-calibrated complexity reports. Every artefact is validated against the Simula-owned schemas in `structured/schema/` before it leaves the pipeline. The domain is read-only with respect to the other three domains at the schema layer: it never mutates their records, and it never calls their runtime MCP tools. The horizontal regulations wrapper governs Simula calls like any other domain.

13.13 Schema contribution

Simula owns the schemas catalogued in chapter 4, including the training-example record, the schema-registry record, the taxonomy structure, and the enumerations listed there. These schemas live under `docs/latex/specs/simula/` as normative text and under `structured/schema/` as the JSON Schema artefacts consumed by the pipeline at `src/training/pipeline/`. Registration in the unified registry is additive: each schema is listed in `src/intelligence/ai-core-pal/data_products/registry.yaml` with its URI, owning domain (`simula`), security class (`confidential`), and LLM routing (`vllm-only`), following the ODPS 4.1 pattern already documented for HANA training data products in that file. Simula also consumes the schemas owned by Arabic, TB, and regulations as inputs; the contract is strictly schema-in, `TrainingExample`-out. Simula performs no schema merges; the two `requirement.schema.json` files in TB and regulations remain distinct when Simula reads them.

13.14 Persona and MCP allow-list

- **Persona id:** `training-engineer`.
- **Roles:** schema extraction, taxonomy authoring, training-data synthesis supervision, critic-calibration review, corpus export.

- **Entitlements:** read/write on the training-example record, schema-registry record, and taxonomy structure; read on every other domain's schemas (Arabic, TB, regulations); run-execute on Simula pipeline stages described in chapter 7.
- **UI route:** `/simula/{tenant}/console` on the shared workspace shell; surfaced only to callers whose entitlement set includes `training-engineer`.
- **MCP tool allow-list:** `simula.schema.extract`, `simula.taxonomy.build`, `simula.generate.examples`, `simula.critic.calibrate`, `simula.evaluate.coverage`, `simula.evaluate.diversity`, `simula.corpus.export`. All other MCP tools are denied by default per MAS MGF pillar 1.

13.15 Agent Rule-Pack Binding

Before initiating any task in this domain, the agent MUST load the following rule packs in order:

1. 📁 `.clinerules` (Global Standards)
2. 📁 `src/intelligence/.clinerules` (Regulations Envelope)
3. 📁 `src/training/.clinerules` (Simula Domain Constraints)
4. 📁 `src/training/.clinerules.runtime-monitor` (Runtime Monitoring Rules)

13.16 Current Implementation Gap Analysis

Critical: Review Gap Analysis Before Implementation

Before beginning any implementation task, agents **MUST** review the implementation gap analysis documented in Chapter A. This chapter provides a comprehensive assessment of the current codebase state against this specification, including:

- P0 gaps that block v1.2 release
- P1 gaps requiring near-term remediation
- P2 gaps for future iterations
- Evidence-backed findings with specific file references
- Sequenced remediation order

13.16.1 Key Findings Summary

The gap analysis (as of 2026-04-21) identified the following critical issues:

Table 13.4: P0 Implementation Gaps Requiring Immediate Attention

Gap ID	Severity	Issue
SIM-GAP-001	P0	Default execution path is still the legacy pipeline
SIM-GAP-002	P0	API and job orchestration wired to legacy generators
SIM-GAP-003	P0	Evaluation modules have drifted from current data contracts
UI-GAP-001	P0	Pipeline UI still models the old six-stage workflow

Implementers should address these gaps in the order specified in Chapter A, §Sequenced Remediation Order, to ensure dependencies are resolved before downstream work begins.

13.16.2 Gap Analysis Validity

Gap Analysis Freshness

The gap analysis in Chapter A has an implicit validity period. As implementation progresses, the analysis should be updated:

- After completing each milestone from Chapter 16
- When new spec-vs-code drift is discovered
- As part of each release's closing criteria

If the gap analysis is more than 4 weeks old without updates, treat it as potentially stale and re-verify findings before relying on them.

13.17 Implementation tasks for this domain

The tasks below are intent-based: each task states the outcome a future agent must achieve and the location in the existing codebase where the work belongs. None of them contains code; each task binds to the Simula sections that define the required behaviour.

Task S1 — HANA schema extraction.

Intent. Realise the HANA schema extraction pipeline in chapter 5, covering MCP discovery, table enumeration, column extraction, and registry population.

> **Why Direct SQL Metadata?**
>
> We chose direct SQL-metadata extraction over high-level MCP discovery for the baseline to ensure we capture SAP-specific table types (Global Temporary, Row-Store) that standard MCP tools often omit.

In-scope. Each of the four stages in the pipeline figure (5.1); schema registry writer; per-table dimension/measure classification. **Out-of-scope.** Non-HANA backends. **Paths to extend.** `src/training/pipeline/` extraction module; registry writer that emits entries conforming to the schema-registry record. **Definition of done.** Running extraction against a HANA Cloud catalogue produces a registry populated with every in-scope table; registry entries validate against the schema-registry record.

Autonomous Verification (Agent Logic):

```
// Verify Task S1
{
  "test_fixture": "tests/fixtures/simula/hana_mock.sql",
  "expected_outcome": "table_count == 14",
  "validation_command": "python src/simula/extract.py --verify-registry"
}
```

Verification. Pipeline dry-run against the fixture store described in chapter 12 produces the expected registry entries.

Task S2 — Taxonomy generation (Algorithm 1).

Intent. Implement Algorithm 1 (reasoning-driven taxonomy generation) from chapter 6 in full, including all four steps (factor identification, breadth-first expansion, best-of-N proposal, critic refinement). **In-scope.** Per-step module with the checkpoints in 6.1; taxonomy structure writer. **Out-of-scope.** Manual taxonomy authoring UI. **Paths to extend.** Taxonomy engine colocated with the pipeline. **Definition of done.** Running Algorithm 1 against an extracted registry produces a taxonomy that validates against the taxonomy structure schema and passes the coverage check in chapter 9. **Verification.** Taxonomy engine unit tests pass; the coverage evaluator reports coverage within the chapter's target band.

Task S3 — Agentic data synthesis (Algorithm 2).

Intent. Implement Algorithm 2 (agentic data synthesis) from chapter 7, including all four steps (mix sampling, meta-prompt generation, complexification, example generation) and the double-critic validation. **In-scope.** The agent roster implied by 7.1; vLLM TurboQuant routing for all LLM calls; complexification step and the two critic agents. **Out-of-scope.** Non-agentic synthesis paths. **Paths to extend.** Generation engine in `src/training/pipeline/`; vLLM client wrapper. **Definition of done.** Running Algorithm 2 against a taxonomy produces `TrainingExample` records that pass both critics and validate against the training-example record schema. **Verification.** Generation unit tests pass on fixture inputs; double-critic agreement rate meets the threshold in chapter 9.

Task S4 — Environment and CLI.

Intent. Deliver the environment and CLI surfaces specified in chapter 8, including the configuration table, CLI usage, and output files listed in 8.2. **In-scope.** Configuration loader; CLI subcommands for each pipeline stage; output-file writers. **Out-of-scope.** IDE integrations. **Paths to extend.** CLI module in `src/training/pipeline/`. **Definition of done.** Each configuration key in the configuration table is read and respected; the CLI runs end-to-end and produces every output file in 8.2. **Verification.** CLI smoke tests pass; output files match the schemas declared for them.

Task S5 — Evaluation framework.

Intent. Deliver the evaluation framework in chapter 9, including calibrated complexity scoring (Algorithm 3, Elo), taxonomic coverage evaluation, embedding-based diversity, and critic-calibration validation. **In-scope.** Each evaluator as a runnable function; report writers. **Out-of-scope.** Real-time evaluation dashboards beyond what the chapter defines. **Paths to extend.** Evaluation module in `src/training/pipeline/`. **Definition of done.** Each evaluator runs against a fixture corpus and emits a report that validates against the corresponding report schema. **Verification.** Evaluator unit tests pass; reports reproduce under fixed seeds.

Task S6 — Test fixtures and testing.

Intent. Deliver the test fixtures specification in chapter 12 and the testing specification in chapter 11 at the coverage thresholds in 11.1. **In-scope.** Invoice, OCR, taxonomy, and training-example fixtures; unit tests for taxonomy traversal and generation; CI wiring. **Out-of-scope.** Performance benchmarking infrastructure. **Paths to extend.** Test directory colocated with the Simula pipeline code. **Definition of done.** Coverage meets 11.1; CI gates block regressions. **Verification.** `make spec-simula` remains green; the Simula test entry points all run.

13.18 Governance obligations applied to this domain

The regulations domain wraps Simula at runtime. Four obligations apply without exception; each is stated as a concrete hook in the Simula codebase.

- **MAS MGF four pillars** (see regulations chapter 3 in `docs/latex/specs/regulations/chapters/03-mgf-framework.tex`). Pillar 1 (assess and bound): every Simula MCP tool is declared in the allow-list in §13.14, deny-by-default. Pillar 2 (design and deploy): the pipeline stages in chapter 7 are traceable end-to-end via the per-request identity attribution described below. Pillar 3 (operate and monitor) and pillar 4 (govern and disclose) are satisfied by emitting the evidence records produced by the Simula evaluation framework (chapter 9) into the shared governance store.
- **Identity attribution** (regulations chapter 10, `10-identity-attribution.tex`). Every MCP call and every pipeline run carries the persona id, agent id, request id, and tenant id. The gateway is the attachment point; Simula code propagates the identity envelope without mutating it.
- **Emergent-capability monitoring** (regulations chapter 8, `08-monitoring-spec.tex`). Extraction, taxonomy generation, synthesis, and evaluation steps are instrumented with the monitoring metrics specified in that chapter; metrics flow to the monitoring sink defined there.

- **Conformance tooling** (regulations chapter 6, `06-conformance-tooling.tex`). The Simula service exposes AI Verify 2.0 plug-in hooks and Moonshot CI/CD 1.1.0 benchmark entry points as required by that chapter.

13.19 Cross-spec integration — full detail

13.19.1 With Arabic AP

Simula consumes Arabic's schemas and fixtures as training input. The contract is schema-in, `TrainingExample`-out: Simula reads `invoice.schema.json`, `vat-checklist.schema.json`, and the Arabic fixture store; it produces `TrainingExample` records. Simula does not call Arabic MCP tools at runtime and does not share state with the Arabic workflow. No cross-domain operational envelope is introduced.

13.19.2 With Trial Balance

Simula consumes TB's schemas and fixtures as training input. The contract is schema-in, `TrainingExample`-out: Simula reads the TB schema set and the TB fixture store; it produces `TrainingExample` records. Simula does not call TB MCP tools at runtime and does not share state with the TB workflow.

13.19.3 With Regulations

Regulations wraps Simula at runtime. Every Simula MCP call is brokered through the regulations governance gate, which applies pillars 1–4 of MAS MGF, attaches the identity envelope, and routes emergent-capability metrics to the monitoring

store. The Simula service code does not call regulations directly; the gateway layer enforces this. Simula also consumes regulations' schemas (requirement, conformance-tool, monitoring metrics, identity and audit) as training input.

13.19.4 With Simula

Simula exports the training-example record, the schema-registry record, the taxonomy structure, and the enum vocabularies under `structured/schema/` (chapter 4); registers the MCP tools enumerated in §13.14; and applies the four governance hooks in §13.18 to every call it serves. External consumers see Simula only through these contracts and through the corpus it emits; there are no private back-channels.

13.20 Wave assignment and sequencing

Implementation runs in three waves.

- **Wave 1 — foundations.** Unified schema registry entry; persona-gated surface for the Training Engineer; identity attribution envelope flowing through the gateway; HANA extraction pipeline (Task S1). Prerequisites: the regulations chapter 10 identity schema must be frozen.
- **Wave 2 — Simula core build-out.** Tasks S2, S3, S4, S5. Prerequisites: Wave 1 complete; vLLM TurboQuant endpoint reachable; Arabic, TB, and regulations schema sets published.
- **Wave 3 — hardening and conformance.** Task S6 and full conformance-tooling hook-up. Prerequisites: Wave 2 complete; regulations Wave 2 conformance tooling integrated; monitoring sink live.

Simula occupies Wave 1 for its extraction pipeline (because other domains depend on Simula's fixtures being reachable), Wave 2 for its core build-out, and Wave 3 for its testing and conformance surfaces.

13.21 Acceptance criteria (platform-wide)

When all four chapter 12 instruction sets are implemented, the platform satisfies the following invariants:

1. **One instance.** A single codebase runs a single logical instance; no domain ships as a separate product.
2. **Persona-gated surfaces.** The four personas each reach a tailored UI on the shared shell; no route, MCP tool, or data product is visible outside its persona's entitlement set.
3. **Unified registry.** Every schema in use is listed in the ODPS 4.1 registry under its owning domain; validators resolve `$refs` through the unified registry only.
4. **No schema merges.** The two `requirement.schema.json` files in TB and regulations remain distinct; no cross-domain composite schemas are produced.
5. **No cross-domain operational envelope.** Runtime orchestration stays within each domain's service; only Simula consumes other domains' outputs, and only as training input.
6. **Governance wraps all three non-regulations domains.** Arabic, TB, and Simula service calls pass through the regulations runtime wrapper without exception.

13.22 References

Within-spec references: chapter 2 (overview), chapter 4 (data schema), chapter 5 (HANA extraction pipeline), chapter 6 (taxonomy engine, Algorithm 1), chapter 7 (generation engine, Algorithm 2), chapter 8 (environment and CLI), chapter 9 (evaluation framework), chapter 12 (test fixtures), chapter 11 (testing specification), chapter 16 (project roadmap).

Peer chapter 12 instruction sets for the other three domains (each is self-contained; consult directly when operating in that domain):

- Arabic AP: `docs/latex/specs/arabic/chapters/12-implementation-instructions.tex`.
- Trial Balance: `docs/latex/specs/tb/chapters/12-implementation-instructions.tex`.
- Regulations: `docs/latex/specs/regulations/chapters/12-implementation-instructions.tex`.

Shared platform artefacts:

- ODPS 4.1 registry: `src/intelligence/ai-core-pal/data_products/registry.yaml`.
- Schema pipeline and resolver: `src/intelligence/schema_pipeline/`.
- Training pipeline (Simula): `src/training/pipeline/`.
- Shared workspace shell: `generativeUI/training-webcomponents-ngx`.

Chapter 14

Cross-Specification Dependencies

This chapter documents how the Simula Training Data Framework specification integrates with the other three SAP OSS specifications. All dependencies are traceable to `docs/schema/registry.json` (v1.2.0).

14.1 Architecture Position

The Simula specification occupies the **training infrastructure** position in the four-domain architecture:

- **Upstream**: Arabic AP (consumes invoice schemas)
- **Upstream**: Trial Balance (consumes TB/PL extract schemas)
- **Horizontal**: Regulations (training data governance)
- **Downstream**: Training JSONL → Model fine-tuning

14.2 Schema Dependencies

14.2.1 Schemas Owned by Simula

Schema	Status	Purpose
taxonomy.schema.json	stable	Factor tree structure
training-example.schema.json	stable	Training JSONL format
hana-schema-entry.schema.json	stable	HANA metadata cache
config.schema.json	stable	Pipeline configuration
meta-prompt.schema.json	stable	Generation prompts

14.2.2 Schemas Consumed from Other Domains

Schema	Source	Usage
invoice.schema.json	arabic	Training for invoice field extraction
vat-checklist.schema.json	arabic	Training for VAT classification
tb-extract.schema.json	tb	Training for account analysis
pl-extract.schema.json	tb	Training for P&L commentary
base-types.yaml	common	Shared type definitions
enums.yaml	common	Complexity levels, verdicts
requirement.schema.json	regulations	Training data compliance

14.3 Regulations Integration

Simula training data generation is governed by Regulations:

14.3.1 Pillar 1: Training Data Bounds

- No PII in generated training examples
- Synthetic data only—no production data extraction
- Complexity distribution balanced (easy/medium/hard)

14.3.2 Pillar 2: Human Accountability

- Schema extraction requires DBA approval
- Training corpus sign-off by domain SME
- Double-critic verdicts logged for audit

14.3.3 Pillar 3: Technical Controls

- Drift monitoring: Training data quality metrics per Regulations Ch.8
- SQL syntax validation gate (see Algorithm 3)
- Conformance testing: AI Verify hooks per Regulations Ch.6

14.3.4 Pillar 4: Data Provenance

- Schema lineage tracking per Regulations Ch.7
- Training corpus versioning with SBOM integration

14.4 Arabic AP Integration

Simula consumes Arabic AP schemas for training data generation:

- **invoice.schema.json** → Invoice field extraction training
- **vat-checklist.schema.json** → VAT category classification
- **Schema extraction** from HANA tables: BKPF, BSEG, LFA1
- Cross-reference: See Arabic AP specification Chapter 2 (Data Schema)

14.5 Trial Balance Integration

Simula consumes TB schemas for training data generation:

- **tb-extract.schema.json** → Account analysis training
- **pl-extract.schema.json** → P&L commentary training
- **variance-record.schema.json** → Variance explanation training
- **Schema extraction** from HANA tables: ACDOCA, BSEG, SKB1
- Cross-reference: See TB specification Chapter 2 (Data Schema)

14.6 Cross-Reference Summary

Target Spec	Chapter	Integration Point
Regulations	Ch.3 (MGF Framework)	Training data governance
Regulations	Ch.6 (Conformance)	AI Verify/Moonshot hooks
Regulations	Ch.7 (Corpus)	Schema lineage tracking
Regulations	Ch.8 (Monitoring)	Training quality metrics
Arabic AP	Ch.2 (Data Schema)	Invoice schema consumption
Arabic AP	Ch.4 (VAT Engine)	VAT category training data
Trial Balance	Ch.2 (Data Schema)	TB/PL schema consumption
Trial Balance	Ch.3 (AI Pipeline)	Variance analysis training

Target Spec	Chapter	Integration Point

Table 14.1: Simula cross-specification references

Chapter 15

Schema Versioning and Interface Contracts

This chapter defines the schema versioning policy for the Simula training data framework, ensuring stable cross-domain interfaces and controlled evolution of data contracts.

15.1 Versioning Policy

15.1.1 Semantic Versioning for Schemas

All Simula schemas follow Semantic Versioning (SemVer) 2.0.0:

- **MAJOR.MINOR.PATCH** (e.g., `1.2.0`)
- **MAJOR**: Breaking changes to required fields or core structure
- **MINOR**: Backward-compatible additions (new optional fields)
- **PATCH**: Documentation, description, or metadata updates

15.1.2 Current Schema Versions

Table 15.1: Simula schema versions as of v1.2.

Schema	Version	Status	Description
`training-example.schema.json`	1.2.0	Stable	Training example record with critic evaluation
`taxonomy.schema.json`	1.2.0	Stable	Taxonomy tree structure with quality metrics
`schema-registry.schema.json`	1.1.0	Stable	HANA schema registry record
`evaluation-report.schema.json`	1.0.0	Draft	Evaluation framework output

15.2 Schema Evolution Rules

15.2.1 Backward Compatibility Requirements

1. **Required fields cannot be removed** in MINOR or PATCH versions.
2. **Required fields cannot be renamed** without a MAJOR version bump.
3. **New optional fields** can be added in MINOR versions.

4. **Enum values can only be added**, not removed, in MINOR versions.
5. **Removing enum values** requires a MAJOR version bump with deprecation notice.

15.2.2 Deprecation Process

1. Mark field as deprecated in schema description (MINOR version).
2. Add `deprecated: true` metadata if supported by tooling.
3. Document migration path in changelog.
4. Maintain deprecated field for at least one MINOR release cycle.
5. Remove deprecated field in next MAJOR version.

15.3 Cross-Domain Interface Contracts

15.3.1 Schema-In, TrainingExample-Out Contract

Simula maintains a strict contract with consumer domains:

Listing 15.1: Cross-domain interface contract.

```
Input (from Arabic AP, TB, Regulations):
  - Schema files conforming to each domain's schema registry
  - Fixture files for validation

Output (to consumers):
  - TrainingExample records conforming to training-example.schema.json
    v1.2.0
  - Taxonomy artifacts conforming to taxonomy.schema.json v1.2.0
```

```
  - Evaluation reports conforming to evaluation-report.schema.json v1.0.0
```

15.3.2 Interface Versioning

Each cross-domain interface is versioned independently:

Table 15.2: Cross-domain interface versions.

Interface	Version	Consumer	Contract
Simula → TB	1.2	Trial Balance	TrainingExample output
Simula → Arabic	1.2	Arabic AP	TrainingExample output
Simula → Regulations	1.2	Regulations	TrainingExample output
TB → Simula	1.0	Simula	Schema + fixture input
Arabic → Simula	1.0	Simula	Schema + fixture input
Regulations → Simula	1.0	Simula	Schema + fixture input

15.4 Schema Validation

15.4.1 Validation Points

Schema validation is enforced at three points:

1. **Input validation**: Before processing consumer schemas
2. **Output validation**: Before emitting TrainingExample records
3. **Report validation**: Before publishing evaluation reports

15.4.2 Validation Implementation

Listing 15.2: Schema validation in the pipeline.

```python
from jsonschema import validate, ValidationError
import json

def validate_training_example(example: dict) -> bool:
    """Validate a training example against the schema."""
    with open("docs/schema/simula/training-example.schema.json") as f:
        schema = json.load(f)

    try:
        validate(instance=example, schema=schema)
        return True
    except ValidationError as e:
        logger.error(f"Schema validation failed: {e.message}")
        return False

# Usage in the pipeline
for example in generated_examples:
    if not validate_training_example(example.to_dict()):
        raise SchemaValidationError(f"Example {example.id} failed validation")
```

15.5 Changelog Management

15.5.1 Schema Changelog Location

Schema changes are documented in:

- `docs/schema/CHANGELOG.md` — consolidated schema changelog
- Each schema file's `description` field — version-specific notes

15.5.2 Changelog Entry Format

Listing 15.3: Changelog entry format.

```
## [1.2.0] - 2026-04-21

### Added
- `critic_evaluation` complex field for double-critic results
- `generation_metadata` for reproducibility tracking
- `quality_signals` for downstream analysis

### Changed
- `complexity_score` now required (was optional in 1.1.0)

### Deprecated
- `difficulty` field (use `complexity_level` instead)
- `taxonomy_mix` field (use `taxonomy_path` instead)

### Migration
- Existing producers should add `complexity_score` with default 0.5
- Replace `difficulty` with `complexity_level` (EASY, MEDIUM, HARD)
```

15.6 Version Negotiation

15.6.1 Runtime Version Checking

The pipeline checks schema versions at runtime:

Listing 15.4: Version negotiation.

```
def check_schema_compatibility(
    schema_path: str,
    min_version: str = "1.0.0",
    max_version: str = "2.0.0",
) -> bool:
    """Check if schema version is within supported range."""
    with open(schema_path) as f:
        schema = json.load(f)

```

```
    schema_id = schema.get("$id", "")
    # Extract version from $id or metadata
    version = extract_version(schema_id)

    return semver.match(version, f"≥{min_version}") and \
           semver.match(version, f"<{max_version}")
```

15.6.2 Graceful Degradation

When encountering an unknown schema version:

1. Log a warning with schema ID and version
2. Attempt to process with latest known compatible schema
3. Flag output with `schema_version_mismatch: true`
4. Include warning in generation stats

15.7 Implementation Checklist

[x] Schema version in `$id` field

[x] `training-example.schema.json` at v1.2.0

[x] `taxonomy.schema.json` at v1.2.0

[ ] Automated schema validation in CI

[ ] Version negotiation in runtime

[ ] Cross-domain contract tests

[ ] Schema changelog automation

15.8 Governance

15.8.1 Schema Change Approval

Schema changes require:

1. RFC submitted to Controls Architecture Team
2. 5-day review period for MINOR changes
3. 10-day review period for MAJOR changes
4. All consuming spec owners must approve breaking changes
5. One-release deprecation window for breaking changes

15.8.2 Change Ownership

Table 15.3: Schema ownership.

Schema	Owner	Approval Authority
`training-example.schema.json`	Simula Lead	Simula + Consumer Leads
`taxonomy.schema.json`	Simula Lead	Simula Lead
`common/enums.yaml`	Controls Arch	All Domain Leads

Chapter 16

Project Roadmap

Scope of this chapter

This chapter is the roadmap for the **Simula training-data framework** project. It is a project-level artefact, not a programme plan. Other SAP OSS Finance projects (Trial Balance, Arabic AP, Regulations) run on the same platform, workspaces, and backends, but each owns its own roadmap, business case, and go/no-go decision. Simula is a *supplier* to those projects: it generates reference training corpora they consume, so its critical path is sequenced before theirs, but its sign-off is independent.

Sequence, not dates.

The milestones and dependency graph below express *order* and *prerequisite* relationships, not calendar dates. This is deliberate. Implementation teams must know what depends on what, and where the schedule can absorb slippage (the Should/Could items); a date column would harden targets into commitments they were never meant to be. Dates, when set, live in the project steering packet and the ticketing system, not in this specification.

16.1 MoSCoW classification

Every milestone below carries one of four priorities. The mapping is contractual: a milestone's priority determines what happens when it is dropped, de-scoped, or slips.

The critical path of this project is the chain of Must items in Table 16.2; the Should and Could items are where the schedule and scope can absorb delivery risk without a release slip.

16.2 Dependency graph

16.3 Project milestones

The Must chain (M01 → M02 → M03 → M04 → M05 → M06 → M07 → M08) is the critical path. The Should items (S01–S03) are the primary *schedule allowance*: consumer corpus generations can slip without affecting Simula v1.2 sign-off. The Could items are the *scope allowance*: they may be cut to v1.3 without renegotiation.

Table 16.1: MoSCoW prioritisation (shared across the four SAP OSS Finance projects).

Priority	Meaning
Must (**M**)	Non-negotiable for v1.2. Without this milestone, the release does not ship.
Should (**S**)	Important but not vital. Ship without it only with sponsor sign-off; document the deviation in the business-case refresh
Could (**C**)	Desirable enhancement. Included if Must and Should items are on track.
Won't (**W**)	Explicitly out of v1.2 scope. Listed so the boundary is legible.

Table 16.2: Simula v1.2 project milestones in dependency order. Priority: M=Must (critical path), S=Should, C=Could, W=Won't.

ID	Prereq	Pri.	Ticket	Milestone
M01	—	M	SIM-ENG-001	Extraction pipeline passes HANA-direct fixtures in the testing chapter.
M02	M01	M	SIM-ENG-002	Taxonomy engine classifies reference corpus with accuracy above Beta threshold.
M03	M02	M	SIM-ENG-003	Generation engine produces schema-valid examples against the unified registry.
M04	M03	M	SIM-ENV-001	Environment CLI enables one-command reproduce of any reference run.
M05	M04	M	SIM-HANA-001	HANA spatial policy (`spatial_unknown_srid_policy`) enforced; strict mode passes unit tests.
M06	M05	M	SIM-RUN-TB	Trial Balance reference run reproducible end-to-end (commit `b8d40a1`, tag `simula-v1.2-worked-example`, seed `20251115`).
M07	M06	M	SIM-RELEASE-001	v1.2 release tagged, archived, and the reproducibility manifest frozen.
M08	M07	M	SIM-BC-REFRESH	Business-case refresh (first cycle after release; then annual).
S01	M07	S	SIM-RUN-AP	Arabic AP training corpus generated from the same pipeline.
S02	S01	S	SIM-RUN-REG	Regulations monitoring corpus generated from the same pipeline.
S03	M07	S	SIM-CONSUMER-DOC	Consumer integration runbook published (how Trial Balance, AP, and Regulations call Simula).
C01	S02 + v1.3 trigger (§16.7)	C	SIM-V13-EXPLORE	v1.3 scope kick-off (multi-tenant isolation, additional taxonomies).
C02	M07	C	SIM-TUNE	Quarterly prompt and classifier retraining from accept/edit telemetry.
W01	—	W	SIM-V12-EXTERNAL	External (non-SAP-OSS) customer access is *not* in v1.2 scope.
W02	—	W	SIM-V12-REALTIME	Real-time incremental generation is *not* in v1.2 scope; batch only.

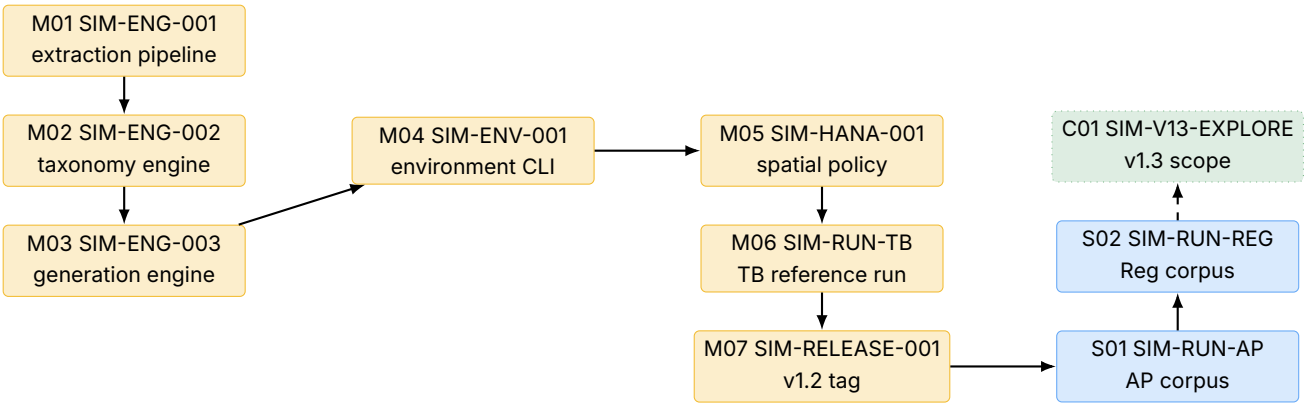

Figure 16.1: Simula dependency graph. Amber nodes are Must (critical path); blue are Should (consumer-corpus allowance); dotted green are Could (discretionary). Solid arrows are hard prerequisites; dashed arrows indicate conditional starts (the target is unlocked when the upstream completes *and* the v1.3 trigger in Section 16.7 fires).

16.4 Launch tiers and exit gates

The project walks the four-tier ladder shared across SAP OSS Finance projects. A tier is exited only when the named exit gate is satisfied; slippage requires sponsor sign-off.

16.5 Kill-switch conditions

The kill-switch is a pre-committed rulebook: a tripped condition enters the action column without further debate. GFAIC and sponsor are notified within 24 hours. The conditions are evaluated continuously (on every cycle) and do not carry dates.

Table 16.3: Simula launch tiers. The exit gate defines the MoSCoW items that must be closed to exit the tier.

Tier	Scope	Exit gate (references MoSCoW IDs)
Alpha	Engineering only, DEV + UAT, fixture corpora only.	M01, M02, M03 closed; schema-registry validation passes on the generated examples; reproducibility confirmed on two consecutive runs.
Beta	One consumer project (Trial Balance) drawing from Simula-generated corpora.	M04, M05, M06 closed; environment CLI reproduces the TB run end-to-end; HANA spatial strict mode green; no open SEV1/SEV2.
GA (limited)	Two consumer projects (TB plus Arabic AP).	M07, M08 closed; S01 closed; consumer-integration runbook (S03) exercised end-to-end; kill-switch green.
Full rollout	All three current consumers (TB, AP, Regulations).	S02 closed; two quarterly reviews sustained; business-case refresh complete; multi-consumer regression suite green.

Table 16.4: Simula kill-switch.

Condition	Evaluation	Action
Reference run non-reproducible across two consecutive seeds (output diff > 0)	per release	**Halt release**; investigate non-determinism; do not ship v1.2 until reproducibility is restored.
Schema-registry validation fails on $> 0.1\,\%$ of generated examples	per run	**Halt consumer delivery**; fix generation engine; re-run.
HANA spatial strict mode detects unknown SRID in production path	continuous	**Quarantine run**; incident-review before releasing corpus to consumer project.
Any consumer project (TB, AP, Regulations) declares red on a Simula-supplied corpus	continuous	**Pause Should items** S01–S03; remediate before next consumer wave.
Shared-platform dependency in Table 16.5 declared red by its owner project	continuous	**Pause downstream Must**; re-plan at next steering.

16.6 Shared-platform dependencies

This project depends on three shared-platform components. Each is tracked here only as a dependency, not a commitment this project can unilaterally adjust. Sequence is expressed in terms of the Must item that first consumes the dependency, not in calendar time.

Table 16.5: Shared-platform dependencies.

Component	Owner	First consumed by	Impact if late
Unified schema registry (`docs/schema/`)	Shared infra	M03	Already `migration.status=complete`; no remaining risk.
vLLM TurboQuant hosting (`vllm-only` policy)	Shared platform	M03	Graceful-degradation ladder defined in the Regulations spec provides fallback; batch-generation slip tolerable.
HANA Cloud spatial extensions	Shared infra	M05	No substitute; slip blocks spatial-policy certification and downstream consumer runs.

Shared-platform dependencies are surfaced at monthly steering. If any upstream owner declares **amber** or **red** on a component in Table 16.5, the affected downstream Must milestone is re-planned at the following steering.

16.7 v1.3 expansion trigger

The Could item C01 (SIM-V13-EXPLORE) is unlocked only when all of the following are observed after two consecutive consumer corpus generations under v1.2:

- aggregate consumer request volume for corpora outside the current taxonomy exceeds 20 % of total Simula runs;
- reproducibility and schema-validation rates sustained at or above release thresholds;
- business-case refresh (M08) explicitly budgets the v1.3 programme.

If any of the three conditions fails, C01 remains in the Could column and is re-evaluated at the next business-case refresh.

16.8 Governance and business-case cadence

Project governance.

The Simula project reports monthly to a project steering (Studio AI Lead, Training Data PM, Pipeline Engineer) and quarterly to the cross-project GFAIC. Monthly steering accepts or rejects amber/red signals on Must items; GFAIC reviews compliance posture and any cross-project dependencies (especially consumer-project risk).

Business-case refresh (M08).

The business case that authorises this project is refreshed on the following rhythm:

- **Mandatory refresh**: at the first cycle after release, and annually thereafter (ticket SIM-BC-REFRESH, priority M in Table 16.2). Refresh packet includes consumer demand assumptions, unit economics, GPU cost realised vs. forecast, and a deviation statement against the baseline.
- **Trigger-based refresh**: any of

  - consumer request volume changes by ≥ 20 % from forecast (either direction);
  - per-run generation cost changes by ≥ 30 % (GPU price, vLLM licensing, HANA I/O);
  - a consumer project's regulatory obligation expands Simula's in-scope taxonomy;
  - any kill-switch in Table 16.4 is tripped.

- **Owner**: Studio AI Lead. Approval routed through the project steering and filed at `docs/simula/business-case/refresh-history.yaml`.

Chapter 17

Implementation Roadmap

Integration with existing chapters

This chapter extends the project roadmap in Chapter 16 with enterprise deployment phases:

- **Chapter 16** (Project Roadmap): The v1.2 milestones (M01–M07) defined there are *prerequisites* for this enterprise roadmap. Phase 1 here begins only after M07 (v1.2 tag) is complete.
- **Chapter 13** (Agent Instructions): Tasks S1–S6 in that chapter deliver the technical foundation; this roadmap adds the business governance wrapper (BR-1.x through BR-5.x).

- **Chapter 3** (Business Requirements): Each phase in this roadmap addresses specific BR requirements, as shown in Table 17.5.
- **Figure 16.1** (Dependency Graph): This roadmap extends the dependency graph with enterprise deployment nodes that follow the Could (C01) milestone.

Like Chapter 16, this roadmap expresses *sequence* and *dependency*, not calendar dates. The phases are ordered by prerequisite relationships and risk, not by time estimates.

Scope of this chapter

This chapter defines the phased implementation roadmap for enterprise deployment of the Simula synthetic data framework. The roadmap is organized into three sequential phases, each building upon the capabilities established in the previous phase. Phases are sequenced by dependency and risk, not by time estimates.

Relationship to Chapter 16 Milestones

This enterprise roadmap extends the technical milestones defined in Chapter 16. The mapping below clarifies which project milestones are addressed by each enterprise phase:

Key Relationships:

- **Phase 1** corresponds to completing the "Should" milestones (S01–S03) plus establishing enterprise governance (BR-3.x, BR-4.x requirements).
- **Phase 2** corresponds to completing the "Could" milestones (C01–C02) plus achieving production scale (BR-2.x requirements).

Table 17.1: Mapping of Project Milestones (Ch. 16) to Enterprise Phases (Ch. 17)

Milestone	MoSCoW	Phase	Relationship
M01–M07	Must	*Prerequisite*	v1.2 technical foundation; completes before Phase 1
S01–S03	Should	Phase 1	Addressed during capability building
C01–C02	Could	Phase 2	Addressed during scaled deployment
M08 (Business Case Refresh)	Must	Phase 1 Exit	Gate check before Phase 2

- **Phase 3** represents continuous operations beyond the initial roadmap, addressing long-term optimization and BR-5.x requirements.
- **M08 (Business Case Refresh)** from Chapter 16 serves as the Phase 1 exit gate, requiring documented ROI before Phase 2 investment.

Typical Duration Estimates

While this roadmap does not commit to calendar dates, the following are typical duration ranges assuming dedicated resources:

Table 17.2: Typical Phase Durations (Resource-Dependent)

Phase	Typical Duration	Assumptions
Phase 1	2–4 months	3 FTE engineers, infrastructure ready
Phase 2	3–6 months	Pilot success, production resources allocated
Phase 3	Ongoing	Operational team in place

These estimates are indicative only. Actual duration depends on organizational readiness, infrastructure availability, and competing priorities.

17.1 Roadmap Overview

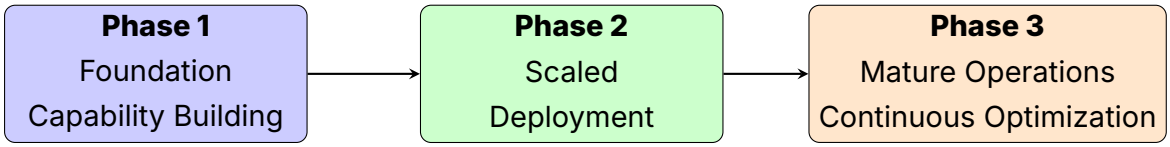

Figure 17.1: Simula implementation phases.

17.1.1 Phase Dependencies

Each phase has explicit entry criteria that must be satisfied before proceeding:

Table 17.3: Phase entry criteria.

Phase	Entry Criteria
Phase 1	Executive sponsorship, initial use case identified
Phase 2	Phase 1 complete, pilot success criteria met, production infrastructure available
Phase 3	Phase 2 complete, multiple use cases in production, operational team trained

17.2 Phase 1: Foundation Capability Building

Phase 1 establishes the foundational capabilities required for Simula deployment, focusing on a single high-value use case to validate the approach before scaling.

17.2.1 Phase 1 Objectives

1. Establish organizational readiness and governance integration
2. Deploy core Simula pipeline for pilot use case
3. Validate quality and utility through downstream evaluation
4. Build internal expertise and operational procedures

17.2.2 Phase 1 Deliverables

1.1 Use Case Definition and Value-Risk Assessment

1.1 Use Case and Risk Assessment

Requirements Addressed: BR-1.1, BR-1.2

Activities:

- Identify and document pilot use case with clear business objective
- Conduct structured value-risk assessment
- Define success criteria and measurement approach
- Obtain executive sponsor approval

Deliverables:

- Use case specification document
- Value-risk assessment matrix
- Success criteria definition
- Executive sign-off

Exit Criteria: Documented use case with approved risk assessment

1.2 Cross-Functional Team Assembly

1.2 Team Organization

Requirements Addressed: BR-1.4
Activities:

- Identify and assign team members across functions
- Establish communication cadence and escalation paths
- Define decision-making authority matrix
- Conduct initial training on Simula concepts

Deliverables:

- Team charter with named individuals
- RACI matrix for Simula activities
- Communication plan
- Training completion records

Exit Criteria: Functional team with clear roles and responsibilities

1.3 Governance Framework Integration

1.3 Governance Integration

Requirements Addressed: BR-1.3
Activities:

- Map synthetic data activities to existing governance frameworks
- Draft synthetic data policies and procedures
- Configure approval workflows in existing tools
- Update stakeholder training materials

Deliverables:

- Synthetic data policy document
- Updated governance framework documentation
- Configured approval workflows
- Training materials

SAP Integration: Configure SAP GRC for policy management
Exit Criteria: Governance policies documented and integrated

1.4 Domain Taxonomy Construction

1.4 Taxonomy Development

Requirements Addressed: BR-2.1
Activities:

- Extract domain factors from pilot use case schema
- Generate initial taxonomies using Simula taxonomy builder
- Conduct domain expert review and validation
- Establish taxonomy version control

Deliverables:

- Domain taxonomy artifacts (JSON/YAML)
- Expert validation sign-off
- Coverage gap analysis report
- Taxonomy versioning repository

SAP Integration: Store taxonomies in SAP HANA Cloud
Exit Criteria: Validated taxonomy with expert approval

1.5 Metadata and Provenance Infrastructure

1.5 Provenance System

Requirements Addressed: BR-2.4
Activities:

- Design metadata schema for generation provenance
- Implement automatic metadata capture in pipeline
- Configure tamper-evident storage
- Build metadata query interface

Deliverables:

- Metadata schema definition
- Provenance capture implementation
- Storage configuration
- Query interface

SAP Integration: Utilize SAP HANA Cloud for metadata storage with hash-chained integrity
Exit Criteria: Automated provenance capture operational

1.6 Quality Assessment Pipeline

1.6 Quality Evaluation

Requirements Addressed: BR-3.1, BR-3.2
Activities:

- Implement intrinsic quality metrics (coverage, diversity, complexity)
- Build automated quality report generation
- Configure downstream validation pipeline (TSTR)
- Establish quality thresholds and gates

Deliverables:

- Quality evaluation module
- Automated report templates
- TSTR pipeline configuration
- Quality gate definitions

SAP Integration: Execute validation via SAP AI Core
Exit Criteria: Quality pipeline generating actionable reports

17.2.3 Phase 1 Exit Criteria

Phase 1 is complete when:

✓ Pilot use case successfully generates synthetic dataset

✓ Downstream model trained on synthetic data meets quality gate

✓ Governance policies documented and approved

✓ Provenance and quality systems operational

✓ Team trained and operational procedures documented

17.3 Phase 2: Scaled Deployment

Phase 2 scales Simula capabilities to production usage with multiple use cases, implementing user-facing interfaces and comprehensive compliance mechanisms.

17.3.1 Phase 2 Objectives

1. Enable self-service generation for business users
2. Implement comprehensive privacy and bias controls
3. Deploy production access control and audit systems
4. Integrate with enterprise ML pipelines

17.3.2 Phase 2 Deliverables

2.1 Parameter Configuration Interface

2.1 User Configuration Interface

Requirements Addressed: BR-2.2
Activities:

- Design visual configuration interface
- Implement parameter validation and constraint checking
- Build preview mode for distribution estimation
- Create configuration templates for common scenarios

Deliverables:

- Web-based configuration UI
- Parameter validation logic
- Preview visualization
- Template library

SAP Integration: Build using SAP Build Apps; store configurations in SAP HANA Cloud
Exit Criteria: Business users can configure and launch generation tasks

2.2 Privacy Risk Assessment System

2.2 Privacy Controls

Requirements Addressed: BR-3.3
Activities:

- Implement membership inference attack testing
- Build similarity scanning against reference datasets
- Configure privacy risk scoring and thresholds
- Design remediation workflow for high-risk samples

Deliverables:

- Privacy risk evaluation module
- Similarity scanner
- Risk scoring dashboard
- Remediation workflow

SAP Integration: Integrate with SAP Data Privacy Integration
Exit Criteria: Automated privacy screening on all generated datasets

2.3 Bias Detection and Mitigation

2.3 Bias Controls

Requirements Addressed: BR-3.4
Activities:

- Implement distribution analysis on sensitive attributes
- Build fairness metric computation
- Design bias mitigation recommendations
- Create bias audit trail

Deliverables:

- Bias detection module
- Fairness metrics dashboard
- Mitigation recommendation engine
- Audit trail storage

Exit Criteria: Bias screening integrated into generation pipeline

2.4 Access Control and Permissions

2.4 RBAC Implementation

Requirements Addressed: BR-4.1
Activities:

- Configure role-based access control
- Implement dataset access approval workflows
- Define domain-specific access policies
- Enable comprehensive access logging

Deliverables:

- RBAC configuration
- Approval workflow
- Access policies
- Audit log configuration

SAP Integration: Leverage SAP Cloud Identity Services and XSUAA
Exit Criteria: Production access controls operational

2.5 ML Pipeline Integration

2.5 MLOps Integration

Requirements Addressed: BR-5.1
Activities:

- Implement REST API for generation task management
- Build SAP AI Core integration
- Configure webhook notifications
- Enable artifact registration in ML pipelines

Deliverables:

- Generation task API
- SAP AI Core connector
- Webhook configuration
- Pipeline integration documentation

SAP Integration: Primary integration with SAP AI Core MLOps
Exit Criteria: Simula integrated into production ML pipelines

2.6 Audit Logging System

2.6 Audit Trail

Requirements Addressed: BR-4.3
Activities:

- Deploy tamper-evident audit log storage
- Implement comprehensive event logging
- Build query interface for investigations
- Configure retention policies

Deliverables:

- Audit log storage
- Event logging implementation
- Query interface
- Retention configuration

SAP Integration: Utilize SAP Audit Log Service
Exit Criteria: Complete audit trail for all Simula activities

17.3.3 Phase 2 Exit Criteria

Phase 2 is complete when:

- ✓ Multiple use cases deployed in production
- ✓ Business users can self-service generation tasks
- ✓ Privacy and bias controls operational
- ✓ Access control and audit systems production-ready
- ✓ Integration with enterprise ML pipelines complete

17.4 Phase 3: Mature Operations and Continuous Optimization

Phase 3 establishes mature operational capabilities with continuous monitoring, optimization, and advanced features for long-term sustainability.

17.4.1 Phase 3 Objectives

1. Implement continuous quality monitoring and drift detection
2. Complete regulatory compliance and transparency requirements
3. Establish incident response capabilities
4. Optimize costs and enable advanced features

17.4.2 Phase 3 Deliverables

3.1 Continuous Monitoring and Drift Detection

3.1 Quality Monitoring

Requirements Addressed: BR-3.5
Activities:

- Implement scheduled quality sampling jobs
- Build drift detection with statistical testing
- Configure alerts for threshold breaches
- Enable automated remediation triggers

Deliverables:

- Monitoring job scheduler
- Drift detection module
- Alert configuration
- Remediation automation

SAP Integration: Configure via SAP AI Core; alerts via SAP Alert Notification Service
Exit Criteria: Continuous monitoring operational with alert coverage

3.2 Regulatory Compliance Completion

3.2 Compliance Framework

Requirements Addressed: BR-4.2
Activities:

- Document compliance position for each applicable regulation
- Obtain legal review sign-off
- Establish regulatory change monitoring
- Implement jurisdiction-specific variations

Deliverables:

- Compliance documentation
- Legal sign-off records
- Regulatory monitoring process
- Jurisdiction policy matrix

SAP Integration: Integrate with SAP GRC for regulatory tracking
Exit Criteria: Documented compliance for all applicable regulations

3.3 Transparency and Explainability

3.3 External Transparency

Requirements Addressed: BR-4.4
Activities:

- Create synthetic data documentation templates
- Develop stakeholder communication guidelines
- Build model card integration
- Prepare FAQ and training materials

Deliverables:

- Documentation templates
- Communication guidelines
- Model card integration
- FAQ materials

Exit Criteria: Transparency materials available for all use cases

3.4 Incident Response Procedures

3.4 Incident Management

Requirements Addressed: BR-4.5
Activities:

- Define incident classification and severity matrix
- Identify responsible parties for each scenario
- Establish escalation and notification paths
- Create remediation playbooks

Deliverables:

- Incident classification matrix
- Responsibility assignment
- Escalation procedures
- Playbooks

SAP Integration: Integrate with existing incident management
Exit Criteria: Incident response procedures documented and tested

3.5 Cost and Budget Optimization

3.5 Cost Management

Requirements Addressed: BR-2.3
Activities:

- Implement budget allocation by cost center
- Build real-time cost tracking
- Configure budget alerts
- Enable cost-per-example reporting

Deliverables:

- Budget allocation interface
- Cost tracking dashboard
- Alert configuration
- Cost reporting

SAP Integration: Integrate with SAP Analytics Cloud
Exit Criteria: Cost visibility and control for all generation activities

3.6 Incremental Generation Capability

3.6 Dataset Evolution

Requirements Addressed: BR-2.5
Activities:

- Implement taxonomy diff computation
- Build targeted generation for changed nodes
- Enable dataset merge with deduplication
- Create efficiency metrics reporting

Deliverables:

- Taxonomy diff module
- Targeted generation logic
- Merge/deduplication logic
- Efficiency reporting

Exit Criteria: Incremental generation available for all datasets

17.4.3 Phase 3 Exit Criteria

Phase 3 is complete when:

✓ Continuous monitoring detecting and alerting on quality issues

✓ Full regulatory compliance documented and approved

✓ Transparency materials available for external stakeholders

✓ Incident response procedures tested and operational

✓ Cost optimization and incremental generation operational

17.5 Key Success Factors

17.5.1 Critical Success Factors

1. **Domain Expert Engagement**: Simula's reasoning-driven approach achieves seedless generation, but taxonomy quality determines coverage and validity. Domain expert investment in taxonomy validation is essential.
2. **Governance-First Approach**: Synthetic data governance should be designed into the project from inception, not retrofitted after large-scale generation. Integrating with existing frameworks is more practical than building new ones.
3. **High-Value Use Case Focus**: The synthetic data market is growing rapidly (projected CAGR 43.9% from 2025–2032), but organizations should avoid broad deployment before validating ROI on focused use cases.
4. **Continuous Validation Loop**: Synthetic data is not a one-time generation task. Establishing "generate → validate → feedback → optimize" loops ensures sustained value.

17.5.2 Risk Mitigation Summary

Table 17.4: Key risks and mitigations by phase.

Phase	Risk	Mitigation
Phase 1	Taxonomy quality insufficient	Expert validation gates, iterative refinement
Phase 1	Governance friction	Early integration with existing frameworks
Phase 2	Privacy incidents	Automated screening, remediation workflows
Phase 2	User adoption low	Self-service interfaces, templates, training
Phase 3	Quality drift undetected	Continuous monitoring, statistical alerts
Phase 3	Cost overruns	Budget controls, cost visibility dashboards

17.6 Requirements Coverage by Phase

Table 17.5: Business requirements addressed by phase.

Category	Phase 1	Phase 2	Phase 3
BR-1.x Strategic	1.1, 1.2, 1.3, 1.4	—	1.5
BR-2.x Generation	2.1, 2.4	2.2	2.3, 2.5
BR-3.x Quality	3.1, 3.2	3.3, 3.4	3.5
BR-4.x Governance	—	4.1, 4.3	4.2, 4.4, 4.5
BR-5.x Integration	—	5.1	5.2, 5.3, 5.4

17.7 SAP Technology Stack Summary

Table 17.6: SAP technology integration by capability.

Capability	SAP Technology
Schema Extraction	SAP HANA Cloud
Taxonomy Storage	SAP HANA Cloud
ML Training Pipeline	SAP AI Core
Configuration Interface	SAP Build Apps
Cost Analytics	SAP Analytics Cloud
Access Control	SAP Cloud Identity Services, XSUAA
Audit Logging	SAP Audit Log Service
Alert Notifications	SAP Alert Notification Service
Governance	SAP GRC
Data Catalog	SAP Data Intelligence
Object Storage	SAP Object Store
Workflow Orchestration	SAP BTP Workflow

Chapter 18

Future Directions (v1.3+)

Future Work

Important Notice
All features described in this chapter are **explicitly out of scope for v1.2**. They represent planned enhancements for v1.3 and beyond. Do not include these features in v1.2 implementation planning or resource allocation.

This chapter is provided for:

- Strategic planning and roadmap visibility
- Stakeholder expectation setting
- Architecture decisions that preserve future extensibility

Integration with existing chapters

This chapter describes future enhancements that extend the foundational Simula specification established in earlier chapters:

- **Chapter 5** (Extraction Pipeline): Enhancement 1 (Multi-Modal) extends schema extraction to document images and time-series.
- **Chapter 6** (Taxonomy Engine): Enhancement 2 (Domain Calibration) adds knowledge anchors to the Algorithm 1 taxonomy builder.
- **Chapter 7** (Generation Engine): Enhancement 3 (Closed-Loop) adds feedback integration to Algorithm 2 data synthesis.
- **Chapter 16** (Project Roadmap): The business requirements in Chapter 3 and implementation roadmap in Chapter 17 are sequenced *after* the v1.2 milestones (M01–M07) complete.
- **Chapter 13** (Agent Instructions): Task S1–S6 remain the technical implementation tasks; the business requirements here provide the governance and compliance wrapper around those tasks.

The enhancements defined here are categorized as **Should** or **Could** items in the MoSCoW classification from Table 16.1, meaning they may be deferred to v1.3 if the Must milestones consume available capacity.

Scope of this chapter

This chapter defines enterprise-grade enhancements to the Simula framework, extending the core "mechanism design" approach to synthetic data generation for production deployment in regulated industries. These enhancements address the practical limitations identified in the original research while preserving Simula's key differentiators: dataset-level design, seedless generation, and independent control of coverage, complexity, and quality.

18.1 Enhancement Overview

The Simula framework, as described in "Reasoning-Driven Synthetic Data Generation and Evaluation" (OpenReview NALsdGEPhB, 2026) and the Google Research blog "Designing Synthetic Datasets for the Real World," introduces four foundational steps:

1. **Global Diversification** — Taxonomic sampling scaffolds for coverage
2. **Local Diversification** — Meta-prompt variation to prevent mode collapse
3. **Complexification** — Controlled difficulty progression
4. **Quality Assurance** — Reasoning-driven verification

This chapter extends each step with enterprise-ready capabilities:

Table 18.1: Simula enhancement dimensions.

Enhancement	Simula Step	Enterprise Value
Multi-Modal Extension	Global + Local	Vision, time-series, tabular data support
Domain Adaptation Calibration	Global	Vertical-specific accuracy (finance, healthcare)
Closed-Loop Validation	Quality Assurance	Measurable ROI via downstream correlation
Programmable Workflow Engine	All Steps	Declarative specifications, audit compliance

18.2 Enhancement 1: Multi-Modal Extension

18.2.1 Current Limitation

The Simula research primarily demonstrates text-based synthetic data generation. Enterprise use cases frequently require multi-modal data:

- **Document processing**: Scanned invoices with embedded tables (Arabic AP domain)
- **Time-series analytics**: Financial metrics with temporal patterns (Trial Balance domain)
- **Structured data**: Tabular records with categorical and numerical fields

18.2.2 Cross-Modal Reasoning Scaffolds

Extend the taxonomic scaffolding approach to multi-modal scenarios by constructing *cross-modal reasoning graphs* that capture relationships between modalities.

Listing 18.1: Cross-modal taxonomy extension.

```
@dataclass
class MultiModalTaxonomyNode(TaxonomyNode):
    """Extended taxonomy node with cross-modal relationships."""

    # Primary modality for this node
    primary_modality: str = "text"  # text, image, tabular, time_series

    # Cross-modal mappings
    cross_modal_mappings: dict[str, str] = field(default_factory=dict)
    # Example: {"image": "invoice_header_region", "text":
    "vendor_name_field"}

    # Modality-specific constraints
    modality_constraints: dict[str, Any] = field(default_factory=dict)
```

```
    # Example: {"image": {"min_resolution": [800, 600], "format": "PNG"}}

class CrossModalTaxonomyBuilder(SimulaTaxonomyBuilder):
    """Build taxonomies with cross-modal consistency."""

    async def build_cross_modal_taxonomy(
        self,
        factor: Factor,
        modalities: list[str],
    ) -> MultiModalTaxonomy:
        """
        Build taxonomy ensuring cross-modal consistency.

        For example, in medical imaging:
        disease_type -> imaging_presentation -> text_description
        ensures generated image-text pairs are semantically aligned.
        """
        # Build primary modality taxonomy
        primary_taxonomy = await self.build_taxonomy(factor)

        # Extend with cross-modal mappings
        for node in primary_taxonomy.get_leaf_nodes():
            mappings = await self._infer_cross_modal_mappings(
                node, modalities
            )
            node.cross_modal_mappings = mappings

        return MultiModalTaxonomy(
            primary_taxonomy=primary_taxonomy,
            modalities=modalities,
        )
```

18.2.3 Modality-Aware Complexification

Define modality-specific complexity metrics:

Table 18.2: Modality-specific complexity dimensions.

Modality	Complexity Dimensions
Text	Vocabulary diversity, syntactic complexity, reasoning depth
Image	Background clutter, lighting variation, occlusion, scale variance
Tabular	Cardinality, sparsity, categorical diversity, numerical range
Time-Series	Trend strength, seasonality, anomaly density, noise ratio

18.2.4 SAP Integration: Document Information Extraction

For SAP Document Information Extraction (DOX) scenarios:

Listing 18.2: SAP DOX multi-modal integration.

```
class SAPDOXModalityAdapter:
    """Adapter for SAP Document Information Extraction."""

    MODALITY_MAPPING = {
        "invoice_header": ["image_region", "text_field", "table_cell"],
        "line_items": ["table_structure", "text_sequence"],
        "totals": ["numerical_field", "currency_format"],
    }

    def generate_cross_modal_example(
        self,
        taxonomy_path: list[str],
        complexity_level: str,
    ) -> MultiModalTrainingExample:
        """
        Generate aligned multi-modal training example.

        Output includes:
        - Synthetic document image (PNG/PDF)
        - Extracted text annotations
        - Structured field mappings
        - Ground truth bounding boxes
        """
        pass
```

18.3 Enhancement 2: Domain Adaptation Calibration

18.3.1 Current Limitation

Simula's seedless, reasoning-driven approach relies on the LLM's inherent knowledge to construct taxonomies. In specialized domains (finance, healthcare, manufacturing), this knowledge may be incomplete or biased toward common patterns, missing critical long-tail scenarios.

18.3.2 Human-in-Loop Knowledge Anchors

Introduce lightweight "knowledge anchors" that guide taxonomy construction without violating the seedless principle:

Listing 18.3: Domain knowledge anchor injection.

```
@dataclass
class DomainKnowledgeAnchor:
    """Lightweight knowledge injection for domain calibration."""

    # Domain identifier
    domain: str  # "trial_balance", "arabic_ap", "regulations"

    # Required coverage nodes (MUST appear in taxonomy)
    required_nodes: list[str] = field(default_factory=list)

    # Prohibited nodes (MUST NOT appear - regulatory/compliance)
    prohibited_nodes: list[str] = field(default_factory=list)

    # Expert-provided complexity anchors
    complexity_anchors: dict[str, float] = field(default_factory=dict)
    # Example: {"multi_currency_reconciliation": 0.85,
    "simple_variance": 0.2}

    # Domain-specific vocabulary
    domain_vocabulary: list[str] = field(default_factory=list)
```

```python

class CalibratedTaxonomyBuilder(SimulaTaxonomyBuilder):
    """Taxonomy builder with domain calibration."""

    async def build_calibrated_taxonomy(
        self,
        factor: Factor,
        anchors: DomainKnowledgeAnchor,
    ) -> Taxonomy:
        """
        Build taxonomy with domain expert calibration.

        Process:
        1. Generate initial taxonomy via LLM reasoning
        2. Validate required nodes are present
        3. Remove prohibited nodes
        4. Calibrate complexity scores against anchors
        5. Present to expert for final review (optional)
        """
        # Step 1: Initial generation
        taxonomy = await self.build_taxonomy(factor)

        # Step 2: Validate required nodes
        existing_nodes = {n.name.lower() for n in taxonomy.get_leaf_nodes()}
        missing = set(anchors.required_nodes) - existing_nodes

        if missing:
            logger.warning(f"Missing required nodes: {missing}")
            taxonomy = await self._inject_missing_nodes(taxonomy, missing)

        # Step 3: Remove prohibited nodes
        taxonomy = self._prune_prohibited(taxonomy, anchors.prohibited_nodes)

        # Step 4: Calibrate complexity
        taxonomy = self._calibrate_complexity(taxonomy, anchors.complexity_anchors)

        return taxonomy
```

18.3.3 Vertical-Specific Complexity Metrics

Define industry-standard complexity dimensions:

Table 18.3: Vertical-specific complexity definitions.

Vertical	Metric	Definition
Finance	Regulatory complexity	Number of compliance rules applicable
	Temporal complexity	Multi-period analysis requirements
	Entity complexity	Cross-entity consolidation depth
Healthcare	Diagnostic complexity	Differential diagnosis breadth
	Protocol complexity	Multi-step treatment pathways
	Comorbidity complexity	Co-occurring condition interactions
Manufacturing	Process complexity	Operation sequence length
	Tolerance complexity	Precision requirements (6-sigma)
	Supply chain complexity	Multi-tier supplier dependencies

18.3.4 SAP Integration: S/4HANA Domain Schemas

Listing 18.4: SAP S/4HANA domain calibration.

```
# Domain anchors for SAP verticals
SAP_DOMAIN_ANCHORS = {
    "trial_balance": DomainKnowledgeAnchor(
        domain="trial_balance",
        required_nodes=[
            "intercompany_elimination",
            "currency_translation",
            "statutory_adjustment",
            "management_overlay",
        ],
        prohibited_nodes=[
            "personal_account_detail",  # Privacy
            "salary_individual",        # Compliance
```

```
        ],
        complexity_anchors={
            "simple_variance_analysis": 0.2,
            "multi_currency_consolidation": 0.75,
            "complex_elimination_entry": 0.9,
        },
    ),
    "arabic_ap": DomainKnowledgeAnchor(
        domain="arabic_ap",
        required_nodes=[
            "vat_compliance",
            "withholding_tax",
            "payment_terms_discount",
            "multi_currency_invoice",
        ],
        complexity_anchors={
            "standard_invoice": 0.15,
            "credit_note_reconciliation": 0.5,
            "multi_vendor_payment_run": 0.8,
        },
    ),
}
```

18.4 Enhancement 3: Closed-Loop Validation

18.4.1 Current Limitation

Simula's quality assurance relies on reasoning-model self-evaluation. This lacks direct feedback from downstream model performance, making it difficult to:

- Quantify ROI of synthetic data generation
- Identify systematic blind spots in coverage
- Adapt generation parameters based on model learning curves

18.4.2 Training-Validation Feedback Loop

Integrate Simula with the model training pipeline to create a closed feedback loop:

Listing 18.5: Closed-loop validation architecture.

```python
@dataclass
class DownstreamPerformanceSignal:
    """Performance signal from downstream model training."""

    # Model identifier
    model_id: str

    # Performance by complexity band
    accuracy_by_complexity: dict[str, float]  # {"easy": 0.95, "hard": 0.72}

    # Performance by taxonomy path
    accuracy_by_taxonomy: dict[str, float]  # {"sql/aggregation/sum": 0.98}

    # Identified weak spots
    underperforming_nodes: list[str]

    # Timestamp
    evaluated_at: str

class ClosedLoopGenerator(SimulaDataGenerator):
    """Generator with downstream feedback integration."""

    async def generate_with_feedback(
        self,
        initial_count: int,
        downstream_evaluator: Callable,
        max_iterations: int = 3,
    ) -> list[TrainingExample]:
        """
        Generate data with iterative refinement based on downstream performance.

        Algorithm:
        1. Generate initial dataset
        2. Train downstream model
        3. Evaluate performance by taxonomy node
```

```
        4. Identify underperforming regions
        5. Adjust sampling weights to oversample weak regions
        6. Generate additional targeted examples
        7. Repeat until convergence or max iterations
        """
        examples = await self.generate_examples(initial_count)

        for iteration in range(max_iterations):
            # Train and evaluate downstream model
            signal = await downstream_evaluator(examples)

            # Check convergence
            if self._is_converged(signal):
                break

            # Adjust sampling weights
            weight_adjustments = self._compute_weight_adjustments(signal)
            self._apply_weight_adjustments(weight_adjustments)

            # Generate targeted examples for weak regions
            additional = await self._generate_targeted(
                signal.underperforming_nodes,
                count=initial_count // 4,
            )
            examples.extend(additional)

        return examples

    def _compute_weight_adjustments(
        self,
        signal: DownstreamPerformanceSignal,
    ) -> dict[str, float]:
        """
        Compute sampling weight adjustments based on performance gaps.

        Regions with lower accuracy receive higher sampling weights
        in subsequent generations.
        """
        adjustments = {}
        mean_accuracy = sum(signal.accuracy_by_taxonomy.values()) /
    len(signal.accuracy_by_taxonomy)

        for node, accuracy in signal.accuracy_by_taxonomy.items():
            if accuracy < mean_accuracy:
```

```
                # Boost underperforming regions
                gap = mean_accuracy - accuracy
                adjustments[node] = 1.0 + (gap * 2)  # Up to 2x boost
            else:
                adjustments[node] = 1.0

        return adjustments
```

18.4.3 Adversarial Verification Layer

Add an adversarial detector to identify systematic biases:

Listing 18.6: Adversarial verification for distribution drift.

```
class AdversarialVerifier:
    """Detect systematic biases in synthetic data."""

    async def detect_distribution_drift(
        self,
        synthetic_examples: list[TrainingExample],
        reference_distribution: dict[str, float],
    ) -> DriftReport:
        """
        Compare synthetic data distribution against reference.

        Detects:
        - Mode collapse (over-concentration in specific patterns)
        - Coverage gaps (missing taxonomy regions)
        - Complexity skew (imbalanced difficulty distribution)
        """
        pass

    async def adversarial_probe(
        self,
        synthetic_examples: list[TrainingExample],
        probe_queries: list[str],
    ) -> ProbeReport:
        """
        Use adversarial probes to test for memorization.

        Checks if generated data inadvertently reproduces
        patterns from the reasoning model's training data.
```

```
        """
        pass
```

18.4.4 SAP Integration: AI Core MLOps

Listing 18.7: SAP AI Core integration for closed-loop training.

```
class SAPAICoreIntegration:
    """Integration with SAP AI Core for MLOps feedback loop."""

    async def submit_training_job(
        self,
        training_data_path: str,
        model_config: dict,
    ) -> str:
        """Submit training job to SAP AI Core."""
        # Use SAP AI Core SDK
        pass

    async def get_evaluation_metrics(
        self,
        execution_id: str,
    ) -> DownstreamPerformanceSignal:
        """Retrieve evaluation metrics from completed training."""
        # Parse AI Core metrics
        pass

    async def trigger_retraining(
        self,
        signal: DownstreamPerformanceSignal,
        threshold: float = 0.85,
    ) -> Optional[str]:
        """Trigger retraining if performance below threshold."""
        if signal.overall_accuracy < threshold:
            # Generate additional targeted examples
            # Submit new training job
            pass
```

18.5 Enhancement 4: Programmable Workflow Engine

18.5.1 Current Limitation

Simula provides conceptual controllability but lacks:

- User-friendly interfaces for business users
- Declarative specification of dataset requirements
- Version control and reproducibility guarantees
- Audit trails for compliance

18.5.2 Declarative Dataset Specification Language

Design a YAML-based DSL for dataset specification:

Listing 18.8: Simula dataset specification DSL.

```yaml
# simula-dataset-spec.yaml
apiVersion: simula/v1
kind: DatasetSpecification
metadata:
  name: trial-balance-text2sql-v3
  domain: trial_balance
  owner: finance-ai-team
  created: 2026-04-21

spec:
  # Target dataset size
  target_count: 10000

  # Coverage requirements
  coverage:
```

```yaml
    taxonomy_paths:
      - path: "query_type/aggregation/*"
        min_coverage: 0.15
      - path: "query_type/filtering/*"
        min_coverage: 0.20
      - path: "complexity/multi_table_join"
        min_coverage: 0.10

  # Complexity distribution
  complexity:
    distribution:
      easy: 0.30
      medium: 0.45
      hard: 0.25
    calibration_anchors:
      - example_id: "expert_example_001"
        target_score: 0.85

  # Quality thresholds
  quality:
    critic_pass_rate: 0.90
    diversity_threshold: 0.70

  # Generation parameters
  generation:
    model: "gpt-4-turbo"
    temperature: 0.7
    best_of_n: 5
    enable_complexification: true
    complexification_ratio: 0.30

  # Validation
  validation:
    schema: "docs/schema/simula/training-example.schema.json"
    downstream_eval: true

  # Output
  output:
    format: jsonl
    path: "data/hana_generated/tb_text2sql_v3.jsonl"
    include_metadata: true
```

18.5.3 Version Control and Reproducibility

Listing 18.9: Git-like versioning for datasets.

```python
@dataclass
class DatasetVersion:
    """Version metadata for reproducibility."""

    version_id: str  # Semantic version or hash
    parent_version: Optional[str]

    # Generation configuration snapshot
    spec_snapshot: dict
    taxonomy_snapshot: dict

    # Reproducibility parameters
    random_seed: int
    model_version: str

    # Statistics
    example_count: int
    coverage_report: dict
    quality_report: dict

    # Audit trail
    created_by: str
    created_at: str
    approval_status: str  # draft, reviewed, approved
    approver: Optional[str]

class DatasetVersionControl:
    """Git-like version control for synthetic datasets."""

    def commit(
        self,
        examples: list[TrainingExample],
        spec: DatasetSpecification,
        message: str,
    ) -> DatasetVersion:
        """Commit a new dataset version."""
        pass

    def diff(
        self,
```

```
        version_a: str,
        version_b: str,
    ) -> DatasetDiff:
        """Compare two dataset versions."""
        pass

    def rollback(
        self,
        version_id: str,
    ) -> list[TrainingExample]:
        """Rollback to a previous version."""
        pass

    def audit_log(
        self,
        version_id: str,
    ) -> list[AuditEntry]:
        """Get complete audit trail for a version."""
        pass
```

18.5.4 SAP Integration: BTP and HANA Cloud

Listing 18.10: SAP BTP integration for workflow orchestration.

```
class SAPBTPWorkflowIntegration:
    """Integration with SAP BTP for workflow orchestration."""

    def __init__(self, btp_config: BTPConfig):
        self.object_store = btp_config.object_store  # SAP Object Store
        self.hana_client = HanaClient(btp_config.hana)
        self.workflow_service = btp_config.workflow_service

    async def execute_workflow(
        self,
        spec: DatasetSpecification,
    ) -> WorkflowExecution:
        """
        Execute Simula workflow on SAP BTP.

        Steps:
        1. Parse specification
        2. Extract schema from HANA Cloud
```

```
        3. Build taxonomies
        4. Generate examples
        5. Validate against schema
        6. Upload to HANA Cloud
        7. Trigger downstream training (optional)
        """
        pass

    async def store_dataset_version(
        self,
        version: DatasetVersion,
        examples: list[TrainingExample],
    ) -> str:
        """Store dataset version in SAP Object Store with HANA
    metadata."""
        # Store examples in Object Store
        object_path = f"simula/{version.version_id}/examples.jsonl"

        # Store metadata in HANA Cloud
        self.hana_client.insert_dataset_version(version)

        return object_path
```

18.6 Regulatory Framework Alignment

These enhancements are designed to align with emerging regulatory guidance on synthetic data:

18.6.1 UK FCA Synthetic Data Expert Group

The Financial Conduct Authority (FCA) Synthetic Data Expert Group identifies nine governance principles[1]:

[1] FCA, "Synthetic Data Expert Group Report," 2024

1. **Transparency** — Enhancement 4 (audit trails, version control)
2. **Explainability** — Enhancement 2 (knowledge anchors, calibration)
3. **Purpose limitation** — Enhancement 4 (declarative specifications)
4. **Data minimization** — Enhancement 1 (targeted generation)
5. **Accuracy** — Enhancement 3 (closed-loop validation)
6. **Security** — SAP BTP integration (access control)
7. **Accountability** — Enhancement 4 (ownership, approval workflows)
8. **Fairness** — Enhancement 3 (adversarial bias detection)
9. **Continuous monitoring** — Enhancement 3 (drift detection)

18.6.2 IEEE Synthetic Data White Paper

The IEEE white paper on synthetic data governance emphasizes[2]:

- **Privacy risk quantification** — Enhancement 3 (adversarial probing)
- **Fidelity-utility tradeoff** — Enhancement 3 (downstream validation)
- **Provenance tracking** — Enhancement 4 (version control)

18.6.3 World Economic Forum Guidelines

The WEF framework for synthetic data in financial services notes[3]:

- Cross-functional governance teams are essential (addressed in Chapter 3)
- Traceability and provenance are foundational (Enhancement 4)

[2] IEEE, "Synthetic Data: Generation, Validation, and Governance," 2025
[3] WEF, "Synthetic Data for Financial Services," 2024

- Continuous monitoring prevents quality degradation (Enhancement 3)

18.7 Implementation Checklist

[ ] Multi-modal taxonomy builder with cross-modal mappings

[ ] Modality-specific complexification metrics

[ ] SAP DOX integration for document extraction

[ ] Domain knowledge anchor injection mechanism

[ ] Vertical-specific complexity calibration

[ ] Closed-loop training feedback integration

[ ] Adversarial verification layer

[ ] SAP AI Core MLOps integration

[ ] Declarative DSL parser

[ ] Dataset version control system

[ ] SAP BTP workflow orchestration

[ ] HANA Cloud metadata storage

- Future Directions (v1.3+)

Chapter A

Implementation Gap Review and Agent Rule Pack

Warning

Point-in-Time Assessment
This appendix is a snapshot of implementation status, not a normative specification.
Last Updated: 2026-04-21

> **Next Review Due**: 2026-05-21 (or after completion of any v1.2 milestone)
>
> If this document is more than 4 weeks old without updates, re-verify findings before relying on them.

A.1 Review scope and constraints

This chapter records a repository review performed against the *Simula Training Data Framework* specification PDF and the current code under `/Users/user/Documents/sap-oss/src` as of 2026-04-21. The review covered four targets:

1. the normative Simula implementation path in `src/training/`;
2. the implementation-agent and maintenance rule location under `src/training/`;
3. the operational frontend and API wrapper in `src/generativeUI/training-webcomponents-ngx/`;
4. runtime monitoring and alerting support for Simula runs.

Two delivery constraints shaped the outcome:

- This workstream was instructed to change only files under `docs/`. No source changes were made under `src/`.
- The intended agent location is `src/training/`. Because source-tree edits were out of scope for this task, the implementation-agent and runtime-monitor definitions are captured here as canonical proposed artefacts to be materialised later under `src/training/.clinerules` and a companion runtime monitor rule file in the same subtree.

A.2 Executive assessment

Table A.1: Current implementation status against the Simula specification.

Area	Status	Assessment
Default CLI and pipeline entrypoints	**Red**	The repository default still executes the legacy CSV/template/Spider pipeline, not the HANA-direct Simula flow described in chapters 5–7.
Simula core modules	**Amber**	HANA extraction, taxonomy building, generation, and evaluator modules exist, but they are only partially integrated and several evaluation components fail at runtime because their data contracts drifted.
Evaluation framework	**Red**	Complexity calibration, coverage, diversity, taxonomy quality, and critic calibration are present as isolated modules, but they are not orchestrated by the runtime path that produces datasets.
Frontend training console	**Amber/Red**	The shell, governance surfaces, and HANA explorer exist, but the pipeline UX still models the older workflow and does not expose Simula-specific artefacts or evaluator outputs.
Runtime monitoring and alerting	**Amber/Red**	Generic health and notification primitives exist, but there is no Simula-specific runtime monitor that watches spec-critical outputs, thresholds, and evaluator failures.
`src/training` agent rule materialisation	**Amber**	The canonical agent location is `src/training`, but the rule files are documented here rather than created in source because this task was constrained to `docs/`.

A.3 Validation performed

The review did not rely only on static reading. The following checks were run against the current repository state:

- `pytest -q src/training/pipeline/tests/test_pipeline.py` passed (**20/20**); this confirms that the *legacy* pipeline remains internally healthy.

- `pytest -q packages/api-server/tests/test_main.py -k 'hana_stats_falls_back_to_preview_when_unconfigured or hana_query_returns_preview_rows_when_unconfigured'` passed (**2/2**); this confirms that preview fallback behaviour is functioning as designed in the current API wrapper.
- Runtime reproductions demonstrated that `SimulaCriticValidator` and `SimulaCoverageEvaluator` are not execution-safe against the current `TrainingExample` and `Taxonomy` classes: the observed failures were `TypeError: TrainingExample.__init__() got an unexpected keyword argument 'schema_context'` and `AttributeError: 'Taxonomy' object has no attribute 'id'`.

A.4 Gap-to-Milestone Mapping Summary

The following table provides a direct mapping from each P0 gap identified in this appendix to the specific project milestone(s) it blocks. This enables prioritized remediation planning aligned with the project roadmap in Chapter 16.

Critical Path Impact: The four P0 gaps collectively block milestones M01 through M07, meaning no v1.2 release is possible until all P0 gaps are resolved. The recommended remediation order in §A.9 sequences these fixes to unblock the critical path as early as possible.

A.5 Findings in `src/training`

Table A.2: P0 Gap to Milestone Blocking Relationship

Gap ID	Severity	Blocks Milestone	Impact Summary
SIM-GAP-001	P0	M01, M04	Default CLI runs legacy pipeline; blocks extraction pipeline certification and environment CLI
SIM-GAP-002	P0	M04, M06	API/job wiring to legacy generators; blocks environment CLI and TB reference run
SIM-GAP-003	P0	M02, M03	Evaluator contract drift; blocks taxonomy engine and generation engine certification
UI-GAP-001	P0	M06, M07	Pipeline UI models old workflow; blocks TB reference run sign-off and v1.2 release
P1 gaps (for reference):			
SIM-GAP-004	P1	M03	Evaluation framework not in runtime path
SIM-GAP-005	P1	M03	Adaptive controls not applied in generator
SIM-GAP-006	P1	M04, M06	Output contract mismatch across runner/API/UI
SIM-GAP-007	P2	M07	Tests only prove legacy path
UI-GAP-002	P1	M06	Simulated telemetry (not real metrics)
UI-GAP-003	P1	M06	Preview/production distinction unclear
UI-GAP-004	P1	M07	No Simula-specific UI for evaluation artefacts
UI-GAP-005	P2	M07	Notifications generic, not Simula-specific

SIM-GAP-001 (P0): the default execution path is still the legacy pipeline

Blocks Milestone: M01 (Extraction pipeline), M04 (Environment CLI)

Observed state. `src/training/pipeline/main.py` still exposes the legacy command set (`extract-schema`, `parse-templates`, `expand`, `format-spider`, `upload`), and the default Makefile target still executes the same CSV/template/Spider flow.

Evidence.

- `src/training/pipeline/main.py:3--13, 30--39, 136--141`
- `src/training/pipeline/Makefile:4--11, 29--30, 35--82`
- By contrast, the Simula orchestration path lives in `src/training/scripts/run_hana_pipeline.py:3--10, 51--196` and is not the repository default.

Impact. An operator following the current default entrypoint runs the old pipeline, not the spec-defined HANA extraction, taxonomy generation, and reasoning-driven synthesis flow. This is the single largest end-to-end compliance gap.

Required remediation. Promote the Simula path to the authoritative CLI/runtime entrypoint, and demote or quarantine the legacy template pipeline behind an explicit compatibility switch.

SIM-GAP-002 (P0): API and job orchestration are wired to the legacy generators

Blocks Milestone: M04 (Environment CLI), M06 (TB reference run)

Observed state. The training console backend does not call `run_hana_pipeline.py`. Instead:

- the pipeline websocket worker runs `make all` in `src/training/pipeline/`;

- the training-generation job runner executes `schema_pipeline/data_generator.py` followed by `prepare_training_data.py`.

Evidence.

- `packages/api-server/src/main.py:2904--3019`
- `packages/api-server/src/main.py:2222--2357`

Impact. Even if the Simula modules were complete, the deployed frontend/API stack would still launch the older data path and publish the older artefact contract. The Simula implementation therefore remains present in code but absent from the operational surface.

Required remediation. Replace the API worker commands with the Simula runner and its report artefacts; keep the current legacy commands only for rollback or migration support.

SIM-GAP-003 (P0): evaluation modules have drifted away from the current data contracts

Blocks Milestone: M02 (Taxonomy engine), M03 (Generation engine)

Observed state. The current `TrainingExample` dataclass contains only `id`, `question`, `sql`, `domain`, `table`, `difficulty`, `taxonomy_mix`, `meta_prompt`, `is_complexified`, and `critic_verdict`. However:

- `SimulaCriticValidator` constructs `TrainingExample` objects using nonexistent fields such as `schema_context`, `taxonomy_path`, `complexity_score`, `critic_passed`, and `generation_metadata`;
- `SimulaCoverageEvaluator` expects each `Taxonomy` to expose `id` and `factor`, and each `TaxonomyNode` to expose `id`; the current taxonomy classes expose `factor_name` and `name/description/children` instead.

Evidence.

- `src/training/pipeline/simula_data_generator.py:32--44`
- `src/training/pipeline/simula_critic_validator.py:212--238, 268--272`
- `src/training/pipeline/simula_coverage_evaluator.py:157--170, 203--215`
- Runtime reproductions raised the concrete exceptions recorded in the validation section above.

Impact. The spec's evaluation chapter is not merely unintegrated; two of its core evaluators are currently unusable against the live Simula data model. This blocks critic calibration, coverage measurement, and any automatic gap-filling workflow built on top of them.

Required remediation. Unify the `TrainingExample`, `Taxonomy`, and evaluator contracts before further orchestration work. Until those contracts are repaired, downstream evaluation integration should be treated as blocked.

SIM-GAP-004 (P1): the Simula orchestration script stops before the full evaluation framework

Observed state. `run_hana_pipeline.py` performs three phases:

1. HANA schema extraction;
2. taxonomy construction;
3. example generation and file write-out.

It writes `extracted_schema.json`, `taxonomies.json`, a timestamped training JSONL, and `pipeline_stats.json`. It does *not* call the evaluator modules for complexity calibration, coverage, diversity, taxonomy quality, critic calibration, or adaptive gap filling.

Evidence.

- `src/training/scripts/run_hana_pipeline.py:51--196`

- Search results show that `SimulaComplexityCalibrator`, `SimulaCoverageEvaluator`, `SimulaCriticValidator`, `SimulaDiversityAnalyzer`, `SimulaDiversityOptimizer`, and `SimulaTaxonomyEvaluator` are defined and re-exported, but not invoked from any orchestration path.

Impact. The implementation currently reaches chapter 5 (*Generation Engine*) but not chapter 7 (*Evaluation Framework*) in an operational sense. The spec's quality gates therefore exist largely as library code and prose, not as emitted run artefacts.

Required remediation. Extend the Simula runner so that every completed generation run emits all evaluation reports required by the spec and records their threshold outcomes into governance and monitoring surfaces.

SIM-GAP-005 (P1): adaptive controls are specified in configuration but not applied in the generator

Observed state. `GenerationConfig` defines `get_adaptive_complexity_ratio()` and `get_adaptive_critic_threshold()`, but the main generation loop continues to use the raw `complexity_ratio` and a plain double-critic pass/fail call without consulting those helpers.

Evidence.

- `src/training/pipeline/simula_config.py:269--326`
- `src/training/pipeline/simula_data_generator.py:494--500, 599--614`
- Repository-wide search found no call sites for `get_adaptive_complexity_ratio()` or `get_adaptive_critic_threshold()` outside their definitions.

Impact. Adaptive features introduced in the v1.2 spec are represented as configuration helpers only. The runtime path does not yet change its behaviour in response to teacher strength or measured complexity.

Required remediation. Thread the adaptive controls through generation, critic acceptance, and report emission so the pipeline can prove that the v1.2 adaptive logic is active.

SIM-GAP-006 (P1): output contracts do not line up across the Simula runner, API, and UI

Observed state. The Simula runner writes timestamped files such as `simula_training_{timestamp}.jsonl` under `data/hana_generated/`. The API and frontend continue to expect the older path `training/pipeline/output/train.jsonl`.

Evidence.

- `src/training/pipeline/simula_data_generator.py:672--716`
- `packages/api-server/src/main.py:2960--2963, 3007--3012, 3354--3360`
- `apps/angular-shell/src/app/pages/pipeline/pipeline.component.ts:549--554`

Impact. Even after the Simula runner is wired in, the frontend and API will not discover the artefacts they are looking for unless the output contract is normalised.

Required remediation. Choose one authoritative artefact layout for Simula runs and apply it consistently across runner, backend, governance records, preview endpoints, and UI components.

SIM-GAP-007 (P2): automated tests still prove only the legacy pipeline path

Observed state. The repository contains passing unit tests for the legacy pipeline modules, but there are no corresponding tests for `run_hana_pipeline.py` or the end-to-end Simula evaluator stack.

Evidence.

- src/training/pipeline/tests/test_pipeline.py
- The passing test run recorded in the validation section above exercised only those legacy tests.

Impact. The project currently has high confidence in the deprecated path and low confidence in the normative one.

Required remediation. Add unit and integration coverage for extraction, taxonomy, critic filtering, evaluator compatibility, output contracts, and the API worker that launches the Simula flow.

A.6 Findings in src/generativeUI/training

UI-GAP-001 (P0): the pipeline UI still models the old six-stage workflow

Blocks Milestone: M06 (TB reference run), M07 (v1.2 release tag)

Observed state. The pipeline page shows the following stages: Preconvert, Extract Schema, Parse Templates, Expand, Build SQL, and Format. Stage transitions are inferred from keywords such as converting, parsing, and formatting.

Evidence.

- apps/angular-shell/src/app/pages/pipeline/pipeline.component.ts:333--339
- apps/angular-shell/src/app/pages/pipeline/pipeline.component.ts:488--497

Impact. The operational UI tells users that the product pipeline is still the preconvert/template/Spider flow. It does not surface the Simula phase model from this spec: HANA extraction, taxonomy building, data generation, and evaluation.

Required remediation. Replace the stage model with the Simula phase graph and emit stage-specific artefacts and metrics from the backend.

UI-GAP-002 (P1): runtime telemetry in the pipeline page is simulated rather than sourced from the backend

Observed state. The pipeline page's worker grid and throughput readout are generated from random numbers whenever the pipeline state is `running`.

Evidence.

- `apps/angular-shell/src/app/pages/pipeline/pipeline.component.ts:369--379`

Impact. The page provides motion, but not observability. It cannot support a runtime-monitoring role because the visual telemetry is not derived from actual execution state.

Required remediation. Bind the concurrency, throughput, and stage timing displays to real worker metrics, log events, and evaluator outcomes.

UI-GAP-003 (P1): the frontend/API operational surface is still wired to preview or compatibility paths

Observed state. The HANA explorer and related backend endpoints deliberately fall back to preview data when credentials or connectivity are missing. The API wrapper exposes this cleanly, but the Simula pipeline has no explicit distinction between a live production run and a preview-backed demonstration surface.

Evidence.

- `packages/api-server/src/main.py:3030--3038, 3149--3197, 3349--3395`
- `apps/angular-shell/src/app/pages/hana-explorer/hana-explorer.component.ts:17--25, 73--82, 112--118`

- API preview tests passed, confirming that preview mode is an intended and working behaviour in the current stack.

Impact. Preview fallback is appropriate for resilience, but the current product surface does not yet enforce the sharper contract the Simula spec needs: production corpus generation must be visibly live, never silently satisfied by preview artefacts.

Required remediation. Add explicit live-vs-preview run state, block governed production launches when required dependencies are only in preview mode, and surface that policy directly in the UI.

UI-GAP-004 (P1): there is no Simula-specific UI for taxonomy and evaluation artefacts

Observed state. The frontend contains generic dashboard, analytics, registry, governance, HANA explorer, and data-product screens, but there is no route or component dedicated to:

- extracted schema inspection as a Simula run artefact;
- taxonomy visualisation and taxonomy quality review;
- coverage, diversity, critic calibration, or Elo complexity reports;
- student-teacher gap tracking or adaptive-complexity decisions.

Evidence.

- `apps/angular-shell/src/app/app.routes.ts:1--92`
- Repository-wide searches for `SimulaCoverageEvaluator`, `SimulaCriticValidator`, `SimulaComplexityCalibrator`, `SimulaDiversityAnalyzer`, and `SimulaTaxonomyEvaluator` found no corresponding Angular surfaces.

Impact. The spec's evaluator outputs cannot yet be reviewed in the operational console, which means human sign-off, explainability, and runtime triage still depend on raw files or logs rather than first-class product screens.

Required remediation. Add explicit Simula result views or extend existing governance and registry pages so that each spec-defined artefact becomes inspectable.

UI-GAP-005 (P2): notification and alerting are generic, not Simula-specific

Observed state. Notification CRUD exists in the API, but the shell notification tray is currently composed from generic health, pipeline state, and GPU pressure held in the store. No Simula-specific alert rules are registered for evaluator failures, missing artefacts, contract drift, or unexpected preview mode.

Evidence.

- `packages/api-server/src/notification_api.py`
- `packages/api-server/src/main.py:590`
- `apps/angular-shell/src/app/components/shell/shell.component.ts:645--691`

Impact. The plumbing for notifications exists, but the actual runtime monitor requested for Simula has not yet been encoded as alerting logic.

Required remediation. Back the notification surface with rules derived from Simula run artefacts and thresholds, not only generic platform health.

A.7 Proposed implementation-agent rule file

Because this review was restricted to `docs/`, the following content is the canonical proposed implementation-agent rule file. It should be materialised at `src/training/.clinerules` once source-tree edits are allowed.

```
# ==============================================================================
# .clinerules - Simula implementation and maintenance agent
# ==============================================================================

Mission:
- Implement and maintain the Simula training data framework in src/training.
- Treat the Simula PDF spec as normative over legacy pipeline behaviour.

Repository reality check:
- The canonical agent location is src/training.
- The default runtime path is still the legacy CSV/template pipeline.
- Simula modules already exist in src/training/pipeline and must become the
  authoritative execution path.

Read these files first on every Simula task:
- docs/pdf/specs/simula-training-spec.pdf
- docs/latex/specs/simula/chapters/12-implementation-instructions.tex
- src/training/scripts/run_hana_pipeline.py
- src/training/pipeline/main.py
- src/training/pipeline/Makefile
- src/training/pipeline/simula_data_generator.py
- src/training/pipeline/simula_taxonomy_builder.py
- src/training/pipeline/hana_schema_extractor.py
- src/training/pipeline/simula_critic_validator.py
- src/training/pipeline/simula_coverage_evaluator.py

Non-negotiable implementation rules:
1. Promote the Simula HANA-direct path over the legacy template path.
2. Keep a single canonical data contract across TrainingExample, Taxonomy,
   evaluation reports, API payloads, and UI artefacts.
3. Do not add new evaluator fields in one module without updating every
   producer and consumer.
4. Wire evaluation into the runtime path:
   - complexity calibration
   - critic calibration
   - taxonomy quality
   - coverage
   - diversity
   - student-teacher gap reporting
5. Emit stable artefacts for every run:
   - extracted_schema.json
   - taxonomies.json
   - training_examples.jsonl (or one agreed canonical JSONL name)
```

```
   - pipeline_stats.json
   - evaluation reports
6. Make API and frontend consume the same artefact contract as the runner.
7. Preserve legacy pipeline code only as an explicit compatibility path.

Required tests before considering a change complete:
- unit tests for any changed Simula module
- contract test for TrainingExample and Taxonomy serialization
- integration test for the runner entrypoint used by the API
- regression test proving frontend/API artefact paths still resolve

Definition of done:
- the default pipeline command launches the Simula flow
- evaluator modules execute without contract errors
- governance captures run status and evaluation outputs
- frontend surfaces the real Simula artefacts instead of legacy placeholders
```

A.8 Proposed runtime-monitor agent rule file

The runtime-monitor agent below is the canonical proposed companion to the implementation agent. It should be materialised later at a path such as `src/training/.clinerules.runtime-monitor` or another agreed operational-agent location.

```
# ==============================================================================
# .clinerules - Simula runtime monitor and alerting agent
# ==============================================================================

Mission:
- Monitor live Simula runs and raise actionable alerts.
- Distinguish live production runs from preview/demo fallbacks.

Primary evidence sources:
- governance training runs and gate checks
- /pipeline/status and websocket pipeline logs
- /jobs and job websocket updates
- HANA health and mode (live vs preview)
- vLLM / TurboQuant health
- emitted run artefacts in the agreed output directory

Treat these as critical alert conditions:
1. Pipeline state = error or job status = failed.
2. HANA extraction completed with zero tables.
3. Taxonomy build completed with zero taxonomies.
4. Example generation completed with zero examples.
```

```
5. Expected artefact missing:
   - extracted_schema.json
   - taxonomies.json
   - training JSONL
   - pipeline_stats.json
6. Any evaluator raises a contract error or runtime exception.
7. Production-governed run is operating in preview mode.
8. Coverage or taxonomy quality falls below agreed thresholds.
9. Critic calibration degrades:
   - p(accept | correct) ≤ 0.80
   - p(accept | corrupted) ≥ 0.30
10. Diversity falls outside target band:
   - global diversity materially below target
   - local diversity materially below target

Recommended severities:
- Critical:
  failed run, preview mode on a production run, missing core artefact,
  contract drift in evaluator modules
- High:
  zero extracted tables, zero taxonomies, zero examples,
  critic calibration failure
- Medium:
  low coverage, low diversity, repeated governance blocks

Required alert payload:
- run_id
- job_id
- team or tenant
- stage
- severity
- exact failing condition
- last relevant log lines
- affected artefact path
- recommended next action

Response rules:
1. Open a notification and attach it to the owning user/team.
2. Append a governance note for Critical and High alerts.
3. Never report "healthy" if the run is only in preview mode.
4. Prefer precise evidence over generic summaries.
5. Collapse duplicate alerts for the same run and condition.
```

A.9 Sequenced remediation order

The findings above should be addressed in the following order:

1. Replace the default execution path and API worker wiring so the Simula runner becomes authoritative.
2. Repair the `TrainingExample`/`Taxonomy`/evaluator contract drift.
3. Normalise run artefact naming and paths across runner, API, governance, and frontend.
4. Integrate the evaluation framework into the operational run.
5. Replace simulated UI telemetry with real runtime metrics and evaluator outputs.
6. Materialise the two proposed agent rule files once source-tree edits outside `docs/` are permitted.

Chapter B

SAP Integration Guide

This appendix provides detailed technical guidance for integrating Simula with SAP technology stack components. It complements the high-level integration points described in Chapter 3 and Chapter 17.

B.1 SAP HANA Cloud Integration

B.1.1 Connection Configuration

Simula requires read-only access to SAP HANA Cloud for schema extraction.

Listing B.1: HANA Cloud connection configuration

```
# Environment variables for HANA Cloud connection
SIMULA_HANA_HOST=your-hana-instance.hana.ondemand.com
SIMULA_HANA_PORT=443
SIMULA_HANA_USER=SIMULA_SERVICE_USER
SIMULA_HANA_PASSWORD_SECRET=hana-password  # Reference to BTP Credential
    Store
SIMULA_HANA_ENCRYPT=true
SIMULA_HANA_SCHEMA=TRAINING_SOURCE

# Python connection using hdbcli
from hdbcli import dbapi

connection = dbapi.connect(
    address=os.environ["SIMULA_HANA_HOST"],
    port=int(os.environ["SIMULA_HANA_PORT"]),
    user=os.environ["SIMULA_HANA_USER"],
    password=get_secret(os.environ["SIMULA_HANA_PASSWORD_SECRET"]),
    encrypt=True,
    sslValidateCertificate=True
)
```

B.1.2 Required Database Permissions

The service user requires the following minimum privileges:

Table B.1: HANA Cloud permissions for Simula service user

Permission	Purpose
`SELECT` on `SYS.TABLES`	List available tables
`SELECT` on `SYS.COLUMNS`	Extract column metadata
`SELECT` on `SYS.CONSTRAINTS`	Extract foreign key relationships
`SELECT` on `SYS.INDEXES`	Identify indexed columns
`SELECT` on target schemas	Read table/view definitions
`CATALOG READ`	Access system views

Listing B.2: Granting minimum permissions

```sql
-- Create service user
CREATE USER SIMULA_SERVICE_USER PASSWORD "***" NO
    FORCE_FIRST_PASSWORD_CHANGE;

-- Grant catalog read
GRANT CATALOG READ TO SIMULA_SERVICE_USER;

-- Grant schema access (repeat for each source schema)
GRANT SELECT ON SCHEMA "TRAINING_SOURCE" TO SIMULA_SERVICE_USER;

-- For metadata extraction
GRANT SELECT ON SYS.TABLES TO SIMULA_SERVICE_USER;
GRANT SELECT ON SYS.COLUMNS TO SIMULA_SERVICE_USER;
GRANT SELECT ON SYS.CONSTRAINTS TO SIMULA_SERVICE_USER;
GRANT SELECT ON SYS.INDEXES TO SIMULA_SERVICE_USER;
```

B.1.3 HANA-Specific Type Mapping

Table B.2: HANA type to Simula schema type mapping

HANA Type	Simula Type	Notes
`TINYINT`, `SMALLINT`, `INTEGER`, `BIGINT`	`integer`	—
`DECIMAL`, `SMALLDECIMAL`	`decimal`	Preserve precision/scale
`REAL`, `DOUBLE`	`float`	—
`VARCHAR`, `NVARCHAR`	`string`	Note max length
`DATE`	`date`	Format: YYYY-MM-DD
`TIME`	`time`	Format: HH:MM:SS
`TIMESTAMP`, `SECONDDATE`	`timestamp`	ISO 8601 format
`BOOLEAN`	`boolean`	—
`VARBINARY`, `BLOB`	`binary`	Excluded from SQL generation
`CLOB`, `NCLOB`	`text`	Large text handling
`ST_GEOMETRY`, `ST_POINT`	`geometry`	See §B.1.4

B.1.4 Spatial Type Handling

HANA spatial types require special handling during schema extraction and SQL generation:

Listing B.3: Spatial type detection and handling

```
SPATIAL_TYPES = {
    "ST_GEOMETRY", "ST_POINT", "ST_LINESTRING", "ST_POLYGON",
    "ST_MULTIPOINT", "ST_MULTILINESTRING", "ST_MULTIPOLYGON",
    "ST_GEOMETRYCOLLECTION"
}

def extract_column_metadata(cursor, table_name: str) ->
    list[ColumnMetadata]:
    """Extract column metadata with spatial type handling."""
    query = """
        SELECT COLUMN_NAME, DATA_TYPE_NAME, LENGTH, SCALE, IS_NULLABLE
        FROM SYS.COLUMNS
        WHERE TABLE_NAME = ?
        ORDER BY POSITION
    """
    cursor.execute(query, (table_name,))

    columns = []
    for row in cursor.fetchall():
        col_type = row[1]
        is_spatial = col_type in SPATIAL_TYPES

        columns.append(ColumnMetadata(
            name=row[0],
            data_type=col_type,
            simula_type="geometry" if is_spatial else
    map_hana_type(col_type),
            length=row[2],
            scale=row[3],
            nullable=row[4] == "TRUE",
            is_spatial=is_spatial,
            spatial_srid=get_srid(cursor, table_name, row[0]) if
    is_spatial else None
        ))

    return columns

```

```python
def get_srid(cursor, table_name: str, column_name: str) -> int | None:
    """Get SRID for spatial column."""
    query = """
        SELECT SRS_ID FROM SYS.ST_GEOMETRY_COLUMNS
        WHERE TABLE_NAME = ? AND COLUMN_NAME = ?
    """
    cursor.execute(query, (table_name, column_name))
    row = cursor.fetchone()
    return row[0] if row else 4326  # Default to WGS84
```

SQL Generation Guidelines for Spatial Columns:

- Use `ST_AsText()` for human-readable output
- Use `ST_Distance()` for proximity queries
- Use `ST_Within()` for containment predicates
- Avoid including raw geometry in `SELECT *` queries

B.2 SAP AI Core Integration

B.2.1 Deployment Configuration

Simula can be deployed as an SAP AI Core serving endpoint:

Listing B.4: SAP AI Core serving configuration

```yaml
# ai-core-serving.yaml
apiVersion: ai.sap.com/v1alpha1
kind: ServingTemplate
metadata:
  name: simula-generator
  labels:
    scenarios.ai.sap.com/id: "simula-training"
    ai.sap.com/version: "1.2.0"
spec:
  template:
    apiVersion: serving.kubeflow.org/v1alpha2
```

```yaml
    spec:
      predictor:
        imagePullSecrets:
          - name: docker-registry-secret
        containers:
          - name: simula
            image: sap-oss/simula:1.2.0
            ports:
              - containerPort: 8080
                protocol: TCP
            env:
              - name: SIMULA_VLLM_ENDPOINT
                valueFrom:
                  secretKeyRef:
                    name: simula-secrets
                    key: vllm-endpoint
              - name: SIMULA_HANA_HOST
                valueFrom:
                  configMapKeyRef:
                    name: simula-config
                    key: hana-host
            resources:
              requests:
                memory: "4Gi"
                cpu: "2"
              limits:
                memory: "8Gi"
                cpu: "4"
```

B.2.2 Training Pipeline Integration

Register Simula outputs with SAP AI Core MLOps:

Listing B.5: Registering training data with AI Core

```python
from ai_core_sdk.models import Artifact, ArtifactScenario

def register_training_dataset(
    ai_core_client,
    dataset_path: str,
    run_id: str,
    metadata: dict
```

```
) -> str:
    """Register generated dataset as AI Core artifact."""

    artifact = Artifact(
        name=f"simula-training-data-{run_id}",
        kind="dataset",
        url=dataset_path,  # S3/Object Store URL
        scenario_id="simula-training",
        labels={
            "run_id": run_id,
            "examples_count": str(metadata["total_examples"]),
            "coverage_score": str(metadata["coverage"]["global"]),
            "diversity_score": str(metadata["diversity"]["global"]),
            "generation_date": metadata["timestamp"]
        }
    )

    result = ai_core_client.artifact.create(artifact)
    return result.id
```

B.3 SAP BTP Services Integration

B.3.1 Credential Store

Store sensitive credentials in SAP BTP Credential Store:

Listing B.6: Retrieving secrets from BTP Credential Store

```
import requests
from cfenv import AppEnv

def get_secret(secret_name: str) -> str:
    """Retrieve secret from BTP Credential Store."""
    env = AppEnv()
    cred_store = env.get_service(name="credential-store")

    response = requests.get(

    f"{cred_store.credentials['url']}/api/v1/credentials/{secret_name}",
```

```
        headers={
            "Authorization": f"Bearer {get_token(cred_store)}",
            "Content-Type": "application/json"
        }
    )
    response.raise_for_status()
    return response.json()["value"]
```

B.3.2 Audit Log Service

Emit audit events to SAP Audit Log Service:

Listing B.7: Audit logging integration

```
from sap_audit_logging import AuditLogger, SecurityMessage

audit_logger = AuditLogger()

def log_generation_event(
    run_id: str,
    user_id: str,
    action: str,
    details: dict
):
    """Log generation event to SAP Audit Log Service."""

    message = SecurityMessage(
        user=user_id,
        action=action,
        object_type="simula.training.dataset",
        object_id=run_id,
        attributes=[
            {"name": "examples_count", "old": None, "new": str(details.get("count"))},
            {"name": "source_schema", "old": None, "new": details.get("schema")},
            {"name": "quality_passed", "old": None, "new": str(details.get("passed"))}
        ]
    )

    audit_logger.log_security_message(message)
```

B.3.3 Alert Notification Service

Configure alerts for quality gate failures:

Listing B.8: Alert notification for quality failures

```python
import requests

def send_quality_alert(
    ans_endpoint: str,
    run_id: str,
    failures: list[dict]
):
    """Send alert via SAP Alert Notification Service."""

    alert = {
        "eventType": "simula.quality.failure",
        "severity": "WARNING",
        "category": "NOTIFICATION",
        "subject": f"Simula quality gate failure: {run_id}",
        "body": format_failure_report(failures),
        "resource": {
            "resourceName": "simula-generator",
            "resourceType": "application"
        },
        "tags": {
            "run_id": run_id,
            "failure_count": str(len(failures))
        }
    }

    response = requests.post(
        f"{ans_endpoint}/cf/producer/v1/resource-events",
        json=alert,
        headers={"Content-Type": "application/json"}
    )
    response.raise_for_status()
```

B.4 SAP Analytics Cloud Integration

B.4.1 Cost Tracking Dashboard

Export generation metrics to SAP Analytics Cloud:

Listing B.9: Exporting metrics to Analytics Cloud

```
def export_metrics_to_sac(
    sac_client,
    run_metrics: dict,
    cost_data: dict
):
    """Export run metrics to SAP Analytics Cloud model."""

    data_row = {
        "run_id": run_metrics["run_id"],
        "date": run_metrics["timestamp"][:10],
        "examples_generated": run_metrics["total_examples"],
        "llm_calls": cost_data["total_llm_calls"],
        "llm_tokens_input": cost_data["tokens"]["input"],
        "llm_tokens_output": cost_data["tokens"]["output"],
        "cost_usd": cost_data["estimated_cost_usd"],
        "cost_per_example": cost_data["estimated_cost_usd"] / run_metrics["total_examples"],
        "coverage_score": run_metrics["coverage"]["global"],
        "diversity_score": run_metrics["diversity"]["global"],
        "critic_acceptance_rate": run_metrics["critic"]["acceptance_rate"]
    }

    sac_client.import_data(
        model_id="simula_generation_metrics",
        data=[data_row]
    )
```

B.4.2 Recommended Dashboard Widgets

Table B.3: SAC dashboard widget recommendations

Widget	Type	Metrics
Generation Volume	Time Series	Examples/day, cumulative total
Cost Analysis	Stacked Bar	Cost by component (LLM, compute, storage)
Quality Trends	Line Chart	Coverage, diversity, critic rate over time
Use Case Distribution	Pie Chart	Examples by taxonomy factor
Cost Efficiency	KPI Card	Cost per example, trend arrow
Quality Gate Status	Traffic Light	Current run pass/fail status

B.5 SAP GRC Integration

B.5.1 Policy Management

Register synthetic data policies in SAP GRC:

Listing B.10: GRC policy registration

```yaml
# grc-policy-registration.yaml
policy:
  id: "SYNTH-DATA-001"
  name: "Synthetic Training Data Generation Policy"
  category: "Data Governance"
  owner: "AI Platform Team"

  controls:
    - id: "CTRL-001"
      name: "Quality Gate Verification"
      description: "All generated datasets must pass quality gates"
      test_procedure: "Verify coverage ≥ 80%, diversity ≥ 65%"
      frequency: "Per generation run"

    - id: "CTRL-002"
      name: "Provenance Logging"
```

```yaml
      description: "All generation runs must have complete provenance"
      test_procedure: "Verify audit trail completeness"
      frequency: "Per generation run"

    - id: "CTRL-003"
      name: "Privacy Screening"
      description: "Generated data must pass privacy checks"
      test_procedure: "Run membership inference test"
      frequency: "Before production use"
```

B.5.2 Compliance Evidence Collection

Automatically collect compliance evidence:

Listing B.11: GRC evidence collection

```python
def collect_compliance_evidence(run_id: str, metrics: dict) -> dict:
    """Collect compliance evidence for GRC reporting."""

    return {
        "run_id": run_id,
        "timestamp": datetime.utcnow().isoformat(),
        "controls": {
            "CTRL-001": {
                "status": "PASS" if metrics["coverage"]["global"] ≥ 0.80
                          and metrics["diversity"]["global"] ≥ 0.65
    else "FAIL",
                "evidence": {
                    "coverage": metrics["coverage"]["global"],
                    "diversity": metrics["diversity"]["global"]
                }
            },
            "CTRL-002": {
                "status": "PASS" if metrics.get("provenance_complete")
    else "FAIL",
                "evidence": {
                    "audit_events": metrics.get("audit_event_count", 0),
                    "provenance_hash": metrics.get("provenance_hash")
                }
            }
        }
    }
```

B.6 SAP Build Apps Integration

For self-service generation interfaces, Simula exposes a REST API compatible with SAP Build Apps:

B.6.1 API Endpoints

Table B.4: Simula REST API endpoints for Build Apps

Endpoint	Method	Purpose
`/api/v1/schemas`	GET	List available schemas
`/api/v1/taxonomies`	GET, POST	List/create taxonomies
`/api/v1/runs`	GET, POST	List/create generation runs
`/api/v1/runs/{id}`	GET	Get run status and results
`/api/v1/runs/{id}/cancel`	POST	Cancel running generation
`/api/v1/artifacts/{id}`	GET	Download generated artifacts
`/api/v1/metrics`	GET	Get aggregated metrics

B.6.2 Build Apps Data Entity

Listing B.12: Build Apps data entity for generation runs

```
{
  "name": "SimulaRun",
  "fields": [
    {"name": "id", "type": "text", "key": true},
    {"name": "name", "type": "text"},
    {"name": "status", "type": "enum", "values": ["pending", "running",
    "completed", "failed"]},
    {"name": "target_count", "type": "number"},
    {"name": "generated_count", "type": "number"},
    {"name": "coverage_score", "type": "number"},
    {"name": "diversity_score", "type": "number"},
    {"name": "created_at", "type": "datetime"},
```

```
12      {"name": "completed_at", "type": "datetime"},
13      {"name": "error_message", "type": "text"}
14    ],
15    "dataSource": {
16      "type": "REST",
17      "baseUrl": "https://simula-api.your-domain.com",
18      "authentication": "OAuth2",
19      "endpoints": {
20        "getAll": "/api/v1/runs",
21        "get": "/api/v1/runs/{id}",
22        "create": "/api/v1/runs"
23      }
24    }
25  }
```

B.7 Quick Reference: Environment Variables

Table B.5: Complete environment variable reference

Variable	Default	Description
HANA Cloud Connection		
`SIMULA_HANA_HOST`	(required)	HANA Cloud host
`SIMULA_HANA_PORT`	443	HANA Cloud port
`SIMULA_HANA_USER`	(required)	Service user
`SIMULA_HANA_PASSWORD_SECRET`	(required)	Credential Store reference
`SIMULA_HANA_SCHEMA`	(required)	Default schema
LLM Configuration		
`SIMULA_VLLM_ENDPOINT`	(required)	vLLM TurboQuant endpoint
`SIMULA_TEACHER_MODEL`	deepseek-coder-33b	Teacher model ID
`SIMULA_STUDENT_MODEL`	codellama-7b	Student model ID
`SIMULA_LLM_TIMEOUT`	60	Request timeout (seconds)
Taxonomy Configuration		
`SIMULA_TAXONOMY_DEPTH`	3	Default taxonomy depth
`SIMULA_BEST_OF_N`	5	Proposals per node
`SIMULA_TAXONOMY_TEMPERATURE`	0.8	Expansion temperature
Generation Configuration		
`SIMULA_TARGET_EXAMPLES`	10000	Default target count
`SIMULA_COMPLEXIFICATION_RATIO`	0.5	Complexification fraction
`SIMULA_SAMPLING_STRATEGY`	uniform	Node sampling strategy
`SIMULA_SCHEMA_MAX_TOKENS`	4096	Schema context budget
`SIMULA_MAX_COLUMNS_PER_TABLE`	30	Column limit per table
Quality Thresholds		
`SIMULA_COVERAGE_THRESHOLD_L1`	1.0	Level 1 coverage
`SIMULA_COVERAGE_THRESHOLD_L2`	0.95	Level 2 coverage
`SIMULA_COVERAGE_THRESHOLD_L3`	0.80	Level 3 coverage
`SIMULA_DIVERSITY_THRESHOLD`	0.65	Global diversity
`SIMULA_CRITIC_ACCEPT_MIN`	0.80	Min P(accept|correct)
`SIMULA_CRITIC_REJECT_MAX`	0.30	Max P(accept|corrupted)
SAP Integration		
`SIMULA_AI_CORE_URL`	—	AI Core endpoint
`SIMULA_AUDIT_LOG_URL`	—	Audit Log Service
`SIMULA_ANS_URL`	—	Alert Notification Service

Chapter C

References

C.1 Source Documents

Table C.1: Source Documents

Document	Path
Simula Research Paper	`docs/6171_Reasoning_Driven_Syntheti.pdf`
Trial Balance Spec	`docs/technical-reports/04-trial-balance-ai-specification.tex`
Arabic AP Spec	`docs/arabic/structured/` (external)

C.2 Implementation Files

Table C.2: Implementation Files

File	Lines
`src/training/pipeline/simula_config.py`	Configuration dataclasses
`src/training/pipeline/simula_llm_client.py`	LLM client with Best-of-N, double-critic
`src/training/pipeline/hana_schema_extractor.py`	HANA MCP extraction
`src/training/pipeline/simula_taxonomy_builder.py`	Algorithm 1 implementation
`src/training/pipeline/simula_data_generator.py`	Algorithm 2 implementation
`src/training/pipeline/simula_complexity_calibrator.py`	Algorithm 3 (Elo scoring)
`src/training/pipeline/simula_coverage_evaluator.py`	Level Ratio Coverage
`src/training/pipeline/simula_diversity_analyzer.py`	Global + local diversity
`src/training/pipeline/simula_critic_validator.py`	Critic calibration
`src/training/scripts/run_hana_pipeline.py`	CLI entry point
`src/training/scripts/split_by_complexity.py`	Complexity split (Figure 7)

C.3 Tools and Dependencies

Table C.3: Tools and Dependencies

Tool	Version	Purpose
vLLM TurboQuant	≥0.8.x	LLM inference backend
AI Core PAL	1.x	HANA MCP server
Python	3.11+	Runtime
asyncio	stdlib	Async orchestration
httpx	0.27+	HTTP client

C.4 Academic References

1. Elo, A. E., & Sloan, S. (1978). *The Rating of Chess Players, Past and Present*. Arco.
2. Yu, P., et al. (2023). "Large Language Models as Simulated Economic Agents." *arXiv preprint arXiv:2305.09857*.
3. Liu, Y., et al. (2024). "Simula: Reasoning-Driven Synthetic Data Generation for Domain-Specific Knowledge Distillation." *OpenReview NALsdGEPhB*.
4. Brown, T., et al. (2020). "Language Models are Few-Shot Learners." *Advances in Neural Information Processing Systems*, 33.
5. Wei, J., et al. (2022). "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models." *NeurIPS 2022*.

C.5 Related Specifications

Arabic AP Specification
`docs/latex/specs/arabic/arabic-ap-spec.tex`

Trial Balance Specification
`docs/latex/specs/tb/tb-review-spec.tex`

Regulations Framework Specification
`docs/latex/specs/regulations/regulations-spec.tex`

Agent Rule Pack Specification
`docs/latex/specs/clinerules-agents/clinerules-agents-spec.tex`

C.6 Schema Registry

All schemas referenced in this specification are registered in the unified schema registry:

Registry Index `docs/schema/registry.json`

Simula Schemas `docs/schema/simula/`

- `training-example.schema.json`
- `taxonomy.schema.json`
- `hana-schema-entry.schema.json`
- `config.schema.json`
- `meta-prompt.schema.json`

Common Types `docs/schema/common/base-types.schema.json`

Shared Enums `docs/schema/common/enums.yaml`

C.7 External Standards

JSON Schema Draft-07 (https://json-schema.org/draft-07/schema)

ODPS 4.1 SAP ODPS Data Product Specification v4.1

MAS MGF Monetary Authority of Singapore Model Governance Framework for AI

AI Verify 2.0 Singapore AI governance testing framework

Moonshot CI/CD 1.1.0 Red-teaming and benchmark toolkit

Chapter D

Global Wave-Dependency Manifest

D.1 Critical Path Visualization (Waves 1–3)

The following chart defines the normative sequencing for implementation agents. A task in a later wave cannot be marked complete until its prerequisites in an earlier wave are validated.

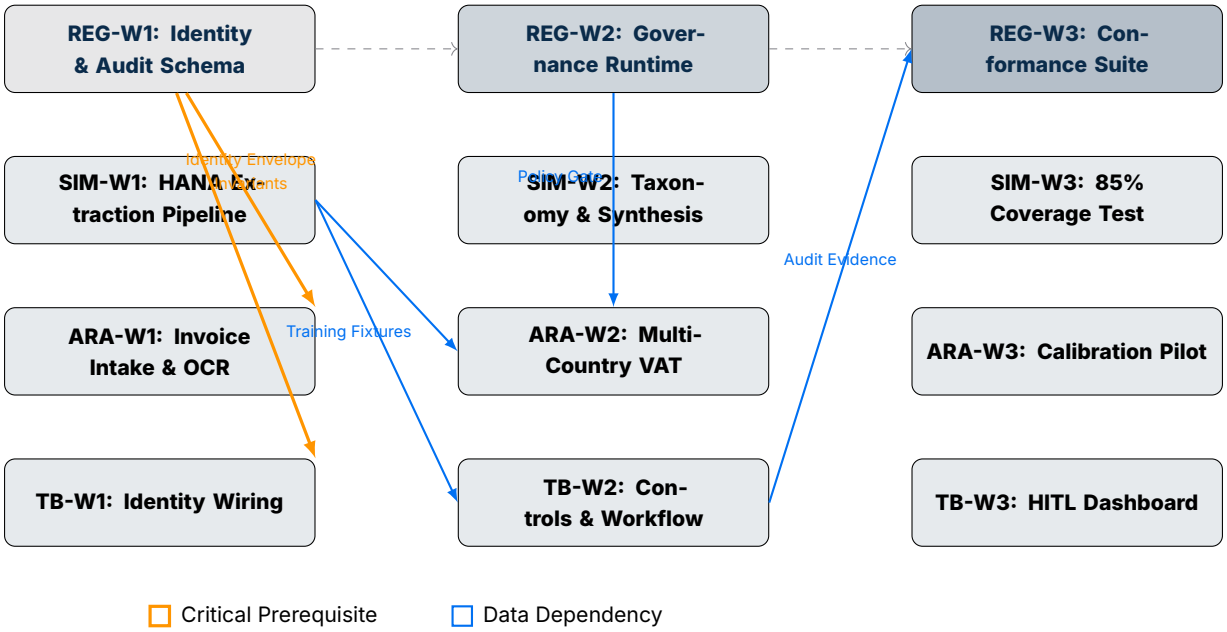

Figure D.1: Cross-specification dependency graph: The blueprint for platform rollout.

D.2 Inter-Domain Prerequisite Matrix

Dependency Enforcement

Implementation agents are strictly forbidden from "Fast-Forwarding" into Wave 2 if Wave 1 critical prerequisites (marked in **Gold**) are in a gap or partial state.

Table D.1: Normative Cross-Domain Prerequisites.

Target Task	Prerequisite	Ref ID	Rationale
ARA-W1	REG-W1 (Identity)	`REG-IDENTITY-01`	Invoice intake must be audit-traced from Day 1.
TB-W2	SIM-W1 (Extraction)	`SIM-EXTRACT-01`	Anomaly detection requires HANA schema metadata.
REG-W2	SIM-W1 (Registry)	`SIM-REGISTRY-01`	Governance monitoring relies on the schema registry.
ARA-W2	REG-W2 (Pillar 1)	`REG-MGF-P1`	VAT release is an HITL gate governed by MAS Pillar 1.
TB-W3	ARA-W2 (Posted Invoices)	`ARA-AP-POST`	Review dashboard depends on posted AP data from Arabic domain.

This specification is part of the SAP-OSS AI Platform documentation suite.
For questions or feedback, contact the Controls Architecture Team.

Typeset with using the sap-oss-apple document class.
Typography: Inter (body), JetBrains Mono (code).

Version 1.2.0 • April 2026